Fedora 24

SELinux User's and Administrator's Guide

Basic and advanced configuration of Security-Enhanced Linux (SELinux)



Mirek Jahoda

Barbora Ančincová

Murray McAllister

Scott Radvan

Daniel Walsh

Dominick Grift

Eric Paris

_		
James	\mathbf{n}	Orric
Jailles	IVI	UHIO

Fedora 24 SELinux User's and Administrator's Guide Basic and advanced configuration of Security-Enhanced Linux (SELinux) Edition 1

Author	Mirek Jahoda	<pre>mirek.jahoda@redhat.com</pre>
Author	Barbora Ančincová	bancinco@redhat.com
Author	Murray McAllister	mmcallis@redhat.com
Author	Scott Radvan	sradvan@redhat.com
Author	Daniel Walsh	dwalsh@redhat.com
Author	Dominick Grift	domg472@gmail.com
Author	Eric Paris	eparis@parisplace.org
Author	James Morris	jmorris@redhat.com

Copyright © 2016 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution—Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at http://creativecommons.org/licenses/by-sa/3.0/. The original authors of this document, and Red Hat, designate the Fedora Project as the "Attribution Party" for purposes of CC-BY-SA. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, JBoss, MetaMatrix, Fedora, the Infinity Logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

For guidelines on the permitted uses of the Fedora trademarks, refer to https://fedoraproject.org/wiki/Legal:Trademark guidelines.

Linux® is the registered trademark of Linus Torvalds in the United States and other countries.

Java® is a registered trademark of Oracle and/or its affiliates.

XFS® is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL® is a registered trademark of MySQL AB in the United States, the European Union and other countries.

All other trademarks are the property of their respective owners.

This book consists of two parts: *SELinux* and *Managing Confined Services*. The former describes the basics and principles upon which SELinux functions, the latter is more focused on practical tasks to set up and configure various services.



Preface	ix
1. Document Conventions	ix
1.1. Typographic Conventions	ix
1.2. Pull-quote Conventions	X
1.3. Notes and Warnings	
2. We Need Feedback!	xi
I. SELinux	1
1. Introduction	3
1.1. Benefits of running SELinux	4
1.2. Examples	5
1.3. SELinux Architecture	
1.4. SELinux States and Modes	6
2. SELinux Contexts	7
2.1. Domain Transitions	8
2.2. SELinux Contexts for Processes	
2.3. SELinux Contexts for Users	. 10
3. Targeted Policy	11
3.1. Confined Processes	11
3.2. Unconfined Processes	
3.3. Confined and Unconfined Users	
3.3.1. The sudo Transition and SELinux Roles	. 18
4. Working with SELinux	21
4.1. SELinux Packages	. 21
4.2. Which Log File is Used	. 22
4.3. Main Configuration File	. 23
4.4. Permanent Changes in SELinux States and Modes	
4.4.1. Enabling SELinux	
4.4.2. Permissive Mode	
4.5. Disabling SELinux	
4.6. Booleans	
4.6.1. Listing Booleans	
4.6.2. Configuring Booleans	
4.6.3. Shell Auto-Completion	
4.7.1 Tomperent Changes: about	
4.7.1. Temporary Changes: chcon	
4.8. The file_t and default_t Types	
4.9. Mounting File Systems	
4.9.1. Context Mounts	
4.9.2. Changing the Default Context	
4.9.3. Mounting an NFS Volume	
4.9.4. Multiple NFS Mounts	
4.9.5. Making Context Mounts Persistent	
4.10. Maintaining SELinux Labels	
4.10.1. Copying Files and Directories	
4.10.2. Moving Files and Directories	. 43
4.10.3. Checking the Default SELinux Context	
4.10.4. Archiving Files with tar	. 45
4.10.5. Archiving Files with star	. 46
4.11. Information Gathering Tools	. 48

	4.12. Multi-Level Security (MLS)	50
	4.12.1. MLS and System Privileges	51
	4.12.2. Enabling MLS in SELinux	51
	4.12.3. Creating a User With a Specific MLS Range	53
	4.12.4. Setting Up Polyinstantiated Directories	54
	4.13. File Name Transition	55
	4.14. Disable ptrace()	56
	4.15. Thumbnail Protection	57
	5. The sepolicy Suite	59
	5.1. The sepolicy Python Bindings	59
	5.2. Generating SELinux Policy Modules: sepolicy generate	59
	5.3. Prioritizing SELinux Policy Modules	
	5.4. Understanding Domain Transitions: sepolicy transition	
	5.5. Generating Manual Pages: sepolicy manpage	61
	6. Confining Users	63
	6.1. Linux and SELinux User Mappings	63
	6.2. Confining New Linux Users: useradd	63
	6.3. Confining Existing Linux Users: semanage login	
	6.4. Changing the Default Mapping	
	6.5. xguest: Kiosk Mode	
	6.6. Booleans for Users Executing Applications	67
	7. sVirt	69
	7.1. Security and Virtualization	70
	7.2. sVirt Labeling	70
	8. Secure Linux Containers	73
	9. SELinux systemd Access Control	75
	9. SELinux systemd Access Control 9.1. SELinux Access Permissions for Services	
		75
	9.1. SELinux Access Permissions for Services	75
	9.1. SELinux Access Permissions for Services	75 78 79
	9.1. SELinux Access Permissions for Services	75 78 79 79
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald	75 78 79 79 80
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald 10. Troubleshooting 10.1. What Happens when Access is Denied 10.2. Top Three Causes of Problems	75 78 79 79 80 80
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald 10. Troubleshooting 10.1. What Happens when Access is Denied 10.2. Top Three Causes of Problems 10.2.1. Labeling Problems	75 78 79 79 80 80 81
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald	75 78 79 79 80 80 81 83
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald 10. Troubleshooting 10.1. What Happens when Access is Denied 10.2. Top Three Causes of Problems 10.2.1. Labeling Problems 10.2.2. How are Confined Services Running? 10.2.3. Evolving Rules and Broken Applications 10.3. Fixing Problems 10.3.1. Linux Permissions	75 78 79 79 80 80 81 83 83
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald 10. Troubleshooting 10.1. What Happens when Access is Denied 10.2. Top Three Causes of Problems 10.2.1. Labeling Problems 10.2.2. How are Confined Services Running? 10.2.3. Evolving Rules and Broken Applications 10.3. Fixing Problems 10.3.1. Linux Permissions 10.3.2. Possible Causes of Silent Denials	75 78 79 80 80 81 83 83 83 84
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald 10. Troubleshooting 10.1. What Happens when Access is Denied 10.2. Top Three Causes of Problems 10.2.1. Labeling Problems 10.2.2. How are Confined Services Running? 10.2.3. Evolving Rules and Broken Applications 10.3. Fixing Problems 10.3.1. Linux Permissions 10.3.2. Possible Causes of Silent Denials 10.3.3. Manual Pages for Services	75 78 79 79 80 81 83 83 84 84
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald 10. Troubleshooting 10.1. What Happens when Access is Denied 10.2. Top Three Causes of Problems 10.2.1. Labeling Problems 10.2.2. How are Confined Services Running? 10.2.3. Evolving Rules and Broken Applications 10.3. Fixing Problems 10.3.1. Linux Permissions 10.3.2. Possible Causes of Silent Denials 10.3.3. Manual Pages for Services 10.3.4. Permissive Domains	75 78 79 79 80 81 83 83 84 84 85
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald	75 78 79 80 81 83 83 84 84 85 86
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald 10. Troubleshooting 10.1. What Happens when Access is Denied 10.2. Top Three Causes of Problems 10.2.1. Labeling Problems 10.2.2. How are Confined Services Running? 10.2.3. Evolving Rules and Broken Applications 10.3. Fixing Problems 10.3.1. Linux Permissions 10.3.2. Possible Causes of Silent Denials 10.3.3. Manual Pages for Services 10.3.4. Permissive Domains 10.3.5. Searching For and Viewing Denials 10.3.6. Raw Audit Messages	75 78 79 80 81 83 83 84 84 85 86 89
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald	75 78 79 79 80 81 83 83 84 84 85 86 89 90
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald 10. Troubleshooting 10.1. What Happens when Access is Denied 10.2. Top Three Causes of Problems 10.2.1. Labeling Problems 10.2.2. How are Confined Services Running? 10.2.3. Evolving Rules and Broken Applications 10.3. Fixing Problems 10.3.1. Linux Permissions 10.3.2. Possible Causes of Silent Denials 10.3.3. Manual Pages for Services 10.3.4. Permissive Domains 10.3.5. Searching For and Viewing Denials 10.3.6. Raw Audit Messages 10.3.7. sealert Messages 10.3.8. Allowing Access: audit2allow	75 78 79 80 81 83 83 84 84 85 86 89 90
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald	75 78 79 80 81 83 83 84 85 86 89 90 92
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald	75 78 79 80 81 83 83 84 85 86 89 90 92
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald	75 78 79 80 81 83 83 84 85 86 89 90 92
	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald 10. Troubleshooting 10.1. What Happens when Access is Denied 10.2. Top Three Causes of Problems 10.2.1. Labeling Problems 10.2.2. How are Confined Services Running? 10.2.3. Evolving Rules and Broken Applications 10.3. Fixing Problems 10.3.1. Linux Permissions 10.3.2. Possible Causes of Silent Denials 10.3.3. Manual Pages for Services 10.3.4. Permissive Domains 10.3.5. Searching For and Viewing Denials 10.3.6. Raw Audit Messages 10.3.7. sealert Messages 10.3.8. Allowing Access: audit2allow 11. Further Information 11.1. Contributors 11.2. Other Resources	75 78 79 80 81 83 83 84 85 86 89 90 92
П. 1	9.1. SELinux Access Permissions for Services 9.2. SELinux and journald	75 78 79 80 81 83 83 84 85 86 89 90 92

13.	The Apache HTTP Server	101
	13.1. The Apache HTTP Server and SELinux	. 101
	13.2. Types	. 103
	13.3. Booleans	. 106
	13.4. Configuration examples	. 109
	13.4.1. Running a static site	
	13.4.2. Sharing NFS and CIFS volumes	
	13.4.3. Sharing files between services	
	13.4.4. Changing port numbers	
	10. In a changing port numbers imminiminiminiminiminiminiminiminiminim	
14.	Samba	117
	14.1. Samba and SELinux	. 117
	14.2. Types	. 118
	14.3. Booleans	. 118
	14.4. Configuration examples	. 119
	14.4.1. Sharing directories you create	119
	14.4.2. Sharing a website	121
15.	File Transfer Protocol	125
	15.1. FTP and SELinux	
	15.2. Types	
	15.3. Booleans	
	15.4. Configuration Examples	
	15.4.1. Uploading to an FTP site	. 128
16	Network File System	133
10.	16.1. NFS and SELinux	
	16.2. Types	
	16.3. Booleans	
	16.4. Configuration Examples	
	16.4.1. Enabling SELinux Labeled NFS Support	. I35
17.	Berkeley Internet Name Domain	137
	17.1. BIND and SELinux	. 137
	17.2. Types	. 137
	17.3. Booleans	
	17.4. Configuration Examples	
	17.4.1. Dynamic DNS	
	•	
18.	Concurrent Versioning System	141
	18.1. CVS and SELinux	
	18.2. Types	. 141
	18.3. Booleans	. 141
	18.4. Configuration Examples	142
	18.4.1. Setting up CVS	. 142
10	Cauld Cooking Provis	145
19.	Squid Caching Proxy	
	19.1. Squid Caching Proxy and SELinux	
	19.2. Types	
	19.3. Booleans	
	19.4. Configuration Examples	
	19.4.1. Squid Connecting to Non-Standard Ports	. 149
20	MariaDB (a replacement for MySQL)	151
_5.	20.1. MariaDB and SELinux	
	20.2. Types	
	20.3. Booleans	
	-0.0. D00104110	

SELinux User's and Administrator's Guide

20.4. Configuration Examples	153
20.4.1. MariaDB Changing Database Location	153
21. PostgreSQL	157
21.1. PostgreSQL and SELinux	157
21.2. Types	158
21.3. Booleans	159
21.4. Configuration Examples	159
21.4.1. PostgreSQL Changing Database Location	159
22. rsync	163
22.1. rsync and SELinux	163
22.2. Types	163
22.3. Booleans	164
22.4. Configuration Examples	164
22.4.1. Rsync as a daemon	164
23. Postfix	169
23.1. Postfix and SELinux	169
23.2. Types	170
23.3. Booleans	170
23.4. Configuration Examples	171
23.4.1. SpamAssassin and Postfix	171
24. DHCP	173
24.1. DHCP and SELinux	173
24.2. Types	174
25. References	175
A. Revision History	177

Preface

1. Document Conventions

This manual uses several conventions to highlight certain words and phrases and draw attention to specific pieces of information.

In PDF and paper editions, this manual uses typefaces drawn from the *Liberation Fonts*¹ set. The Liberation Fonts set is also used in HTML editions if the set is installed on your system. If not, alternative but equivalent typefaces are displayed. Note: Red Hat Enterprise Linux 5 and later includes the Liberation Fonts set by default.

1.1. Typographic Conventions

Four typographic conventions are used to call attention to specific words and phrases. These conventions, and the circumstances they apply to, are as follows.

Mono-spaced Bold

Used to highlight system input, including shell commands, file names and paths. Also used to highlight keycaps and key combinations. For example:

To see the contents of the file my_next_bestselling_novel in your current working directory, enter the cat my_next_bestselling_novel command at the shell prompt and press Enter to execute the command.

The above includes a file name, a shell command and a keycap, all presented in mono-spaced bold and all distinguishable thanks to context.

Key combinations can be distinguished from keycaps by the hyphen connecting each part of a key combination. For example:

Press **Enter** to execute the command.

Press **Ctrl+Alt+F2** to switch to the first virtual terminal. Press **Ctrl+Alt+F1** to return to your X-Windows session.

The first paragraph highlights the particular keycap to press. The second highlights two key combinations (each a set of three keycaps with each set pressed simultaneously).

If source code is discussed, class names, methods, functions, variable names and returned values mentioned within a paragraph will be presented as above, in **mono-spaced bold**. For example:

File-related classes include **filesystem** for file systems, **file** for files, and **dir** for directories. Each class has its own associated set of permissions.

Proportional Bold

This denotes words or phrases encountered on a system, including application names; dialog box text; labeled buttons; check-box and radio button labels; menu titles and sub-menu titles. For example:

Choose $System \rightarrow Preferences \rightarrow Mouse$ from the main menu bar to launch Mouse Preferences. In the Buttons tab, click the Left-handed mouse check box and click

¹ https://fedorahosted.org/liberation-fonts/

Close to switch the primary mouse button from the left to the right (making the mouse suitable for use in the left hand).

To insert a special character into a **gedit** file, choose **Applications** \rightarrow **Accessories** \rightarrow **Character Map** from the main menu bar. Next, choose **Search** \rightarrow **Find...** from the **Character Map** menu bar, type the name of the character in the **Search** field and click **Next**. The character you sought will be highlighted in the **Character Table**. Double-click this highlighted character to place it in the **Text to copy** field and then click the **Copy** button. Now switch back to your document and choose **Edit** \rightarrow **Paste** from the **gedit** menu bar.

The above text includes application names; system-wide menu names and items; application-specific menu names; and buttons and text found within a GUI interface, all presented in proportional bold and all distinguishable by context.

Mono-spaced Bold Italic or Proportional Bold Italic

Whether mono-spaced bold or proportional bold, the addition of italics indicates replaceable or variable text. Italics denotes text you do not input literally or displayed text that changes depending on circumstance. For example:

To connect to a remote machine using ssh, type **ssh** *username@domain.name* at a shell prompt. If the remote machine is **example.com** and your username on that machine is john, type **ssh john@example.com**.

The **mount** -o **remount file**-**system** command remounts the named file system. For example, to remount the **/home** file system, the command is **mount** -o **remount /home**.

To see the version of a currently installed package, use the rpm -q package command. It will return a result as follows: package-version-release.

Note the words in bold italics above — username, domain.name, file-system, package, version and release. Each word is a placeholder, either for text you enter when issuing a command or for text displayed by the system.

Aside from standard usage for presenting the title of a work, italics denotes the first use of a new and important term. For example:

Publican is a *DocBook* publishing system.

1.2. Pull-quote Conventions

Terminal output and source code listings are set off visually from the surrounding text.

Output sent to a terminal is set in **mono-spaced roman** and presented thus:

```
books Desktop documentation drafts mss photos stuff svn
books_tests Desktop1 downloads images notes scripts svgs
```

Source-code listings are also set in mono-spaced roman but add syntax highlighting as follows:

```
package org.jboss.book.jca.ex1;
import javax.naming.InitialContext;
```

```
public class ExClient
  public static void main(String args[])
       throws Exception
     InitialContext iniCtx = new InitialContext();
                           = iniCtx.lookup("EchoBean");
     Object
                    ref
                    home = (EchoHome) ref;
     EchoHome
     Echo
                          = home.create();
                    echo
     System.out.println("Created Echo");
     System.out.println("Echo.echo('Hello') = " + echo.echo("Hello"));
  }
}
```

1.3. Notes and Warnings

Finally, we use three visual styles to draw attention to information that might otherwise be overlooked.



Note

Notes are tips, shortcuts or alternative approaches to the task at hand. Ignoring a note should have no negative consequences, but you might miss out on a trick that makes your life easier.



Important

Important boxes detail things that are easily missed: configuration changes that only apply to the current session, or services that need restarting before an update will apply. Ignoring a box labeled 'Important' will not cause data loss but may cause irritation and frustration.



Warning

Warnings should not be ignored. Ignoring warnings will most likely cause data loss.

2. We Need Feedback!

If you find a typographical error in this manual, or if you have thought of a way to make this manual better, we would love to hear from you! Please submit a report in Bugzilla: http://bugzilla.redhat.com/bugzilla/ against the product **Fedora Documentation.**

When submitting a bug report, be sure to mention the manual's identifier: selinux-guide

If you have a suggestion for improving the documentation, try to be as specific as possible when describing it. If you have found an error, please include the section number and some of the surrounding text so we can find it easily.

Part I. SELinux



Introduction

Security-Enhanced Linux (SELinux) is an implementation of a *mandatory access control* mechanism in the Linux kernel, checking for allowed operations after standard *discretionary access controls* are checked. SELinux can enforce rules on files and processes in a Linux system, and on their actions, based on defined policies.

When using SELinux, files, including directories and devices, are referred to as objects. Processes, such as a user running a command or the Mozilla Firefox application, are referred to as subjects. Most operating systems use a Discretionary Access Control (DAC) system that controls how subjects interact with objects, and how subjects interact with each other. On operating systems using DAC, users control the permissions of files (objects) that they own. For example, on Linux operating systems, users could make their home directories world-readable, giving users and processes (subjects) access to potentially sensitive information, with no further protection over this unwanted action.

Relying on DAC mechanisms alone is fundamentally inadequate for strong system security. DAC access decisions are only based on user identity and ownership, ignoring other security-relevant information such as the role of the user, the function and trustworthiness of the program, and the sensitivity and integrity of the data. Each user typically has complete discretion over their files, making it difficult to enforce a system-wide security policy. Furthermore, every program run by a user inherits all of the permissions granted to the user and is free to change access to the user's files, so minimal protection is provided against malicious software. Many system services and privileged programs run with coarse-grained privileges that far exceed their requirements, so that a flaw in any one of these programs could be exploited to obtain further system access.¹

The following is an example of permissions used on Linux operating systems that do not run Security-Enhanced Linux (SELinux). The permissions and output in these examples may differ slightly from your system. Use the following command to view file permissions:

```
~]$ ls -l file1
-rwxrw-r-- 1 user1 group1 0 2009-08-30 11:03 file1
```

In this example, the first three permission bits, **rwx**, control the access the Linux **user1** user (in this case, the owner) has to **file1**. The next three permission bits, **rw-**, control the access the Linux **group1** group has to **file1**. The last three permission bits, **r--**, control the access everyone else has to **file1**, which includes all users and processes.

Security-Enhanced Linux (SELinux) adds Mandatory Access Control (MAC) to the Linux kernel, and is enabled by default in Fedora. A general purpose MAC architecture needs the ability to enforce an administratively-set security policy over all processes and files in the system, basing decisions on labels containing a variety of security-relevant information. When properly implemented, it enables a system to adequately defend itself and offers critical support for application security by protecting against the tampering with, and bypassing of, secured applications. MAC provides strong separation of applications that permits the safe execution of untrustworthy applications. Its ability to limit the privileges associated with executing processes limits the scope of potential damage that can result from the exploitation of vulnerabilities in applications and system services. MAC enables information

¹ "Integrating Flexible Support for Security Policies into the Linux Operating System", by Peter Loscocco and Stephen Smalley. This paper was originally prepared for the National Security Agency and is, consequently, in the public domain. See the *original paper* [http://www.nsa.gov/research/_files/selinux/papers/freenix01/index.shtml] for details and the document as it was first released. Any edits and changes were done by Murray McAllister.

to be protected from legitimate users with limited authorization as well as from authorized users who have unwittingly executed malicious applications.²

The following is an example of the labels containing security-relevant information that are used on processes, Linux users, and files, on Linux operating systems that run SELinux. This information is called the SELinux *context*, and is viewed using the following command:

```
~]$ ls -Z file1
-rwxrw-r-- user1 group1 unconfined_u:object_r:user_home_t:s0 file1
```

In this example, SELinux provides a user (unconfined_u), a role (object_r), a type (user_home_t), and a level (s0). This information is used to make access control decisions. With DAC, access is controlled based only on Linux user and group IDs. It is important to remember that SELinux policy rules are checked *after* DAC rules. SELinux policy rules are not used if DAC rules deny access first.



Linux and SELinux Users

On Linux operating systems that run SELinux, there are Linux users as well as SELinux users. SELinux users are part of SELinux policy. Linux users are mapped to SELinux users. To avoid confusion, this guide uses *Linux user* and *SELinux user* to differentiate between the two.

1.1. Benefits of running SELinux

- All processes and files are labeled with a type. A type defines a domain for processes, and a type
 for files. Processes are separated from each other by running in their own domains, and SELinux
 policy rules define how processes interact with files, as well as how processes interact with each
 other. Access is only allowed if an SELinux policy rule exists that specifically allows it.
- Fine-grained access control. Stepping beyond traditional UNIX permissions that are controlled at user discretion and based on Linux user and group IDs, SELinux access decisions are based on all available information, such as an SELinux user, role, type, and, optionally, a level.
- SELinux policy is administratively-defined, enforced system-wide, and is not set at user discretion.
- Reduced vulnerability to privilege escalation attacks. Processes run in domains, and are therefore separated from each other. SELinux policy rules define how processes access files and other processes. If a process is compromised, the attacker only has access to the normal functions of that process, and to files the process has been configured to have access to. For example, if the Apache HTTP Server is compromised, an attacker cannot use that process to read files in user home directories, unless a specific SELinux policy rule was added or configured to allow such access.
- SELinux can be used to enforce data confidentiality and integrity, as well as protecting processes from untrusted inputs.

However, SELinux is not:

² "Meeting Critical Security Objectives with Security-Enhanced Linux", by Peter Loscocco and Stephen Smalley. This paper was originally prepared for the National Security Agency and is, consequently, in the public domain. See the *original paper* [http://www.nsa.gov/research/_files/selinux/papers/ottawa01/index.shtml] for details and the document as it was first released. Any edits and changes were done by Murray McAllister.

- · antivirus software,
- · a replacement for passwords, firewalls, or other security systems,
- · an all-in-one security solution.

SELinux is designed to enhance existing security solutions, not replace them. Even when running SELinux, it is important to continue to follow good security practices, such as keeping software up-to-date, using hard-to-guess passwords, firewalls, and so on.

1.2. Examples

The following examples demonstrate how SELinux increases security:

- The default action is deny. If an SELinux policy rule does not exist to allow access, such as for a process opening a file, access is denied.
- SELinux can confine Linux users. A number of confined SELinux users exist in SELinux policy.
 Linux users can be mapped to confined SELinux users to take advantage of the security rules and
 mechanisms applied to them. For example, mapping a Linux user to the SELinux user_u user,
 results in a Linux user that is not able to run (unless configured otherwise) set user ID (setuid)
 applications, such as sudo and su, as well as preventing them from executing files and applications
 in their home directory. If configured, this prevents users from executing malicious files from their
 home directories.
- Process separation is used. Processes run in their own domains, preventing processes from
 accessing files used by other processes, as well as preventing processes from accessing other
 processes. For example, when running SELinux, unless otherwise configured, an attacker cannot
 compromise a Samba server, and then use that Samba server as an attack vector to read and write
 to files used by other processes, such as databases used by MariaDB.
- SELinux helps limit the damage made by configuration mistakes. Domain Name System (DNS) servers often replicate information between each other in what is known as a zone transfer.
 Attackers can use zone transfers to update DNS servers with false information. When running the Berkeley Internet Name Domain (BIND) as a DNS server in Fedora, even if an administrator forgets to limit which servers can perform a zone transfer, the default SELinux policy prevents zone files from being updated via zone transfers, by the BIND named daemon itself, and by other processes.
- See the NetworkWorld.com⁴ article, A seatbelt for server software: SELinux blocks real-world exploits⁵⁶, for background information about SELinux, and information about various exploits that SELinux has prevented.

1.3. SELinux Architecture

SELinux is a Linux security module that is built into the Linux kernel. SELinux is driven by loadable policy rules. When security-relevant access is taking place, such as when a process attempts to open a file, the operation is intercepted in the kernel by SELinux. If an SELinux policy rule allows the operation, it continues, otherwise, the operation is blocked and the process receives an error.

³ Text files that include information, such as host name to IP address mappings, that are used by DNS servers.

⁴ http://www.networkworld.com

 $^{^{5}}$ http://www.networkworld.com/news/2008/022408-selinux.html

⁶ Marti, Don. "A seatbelt for server software: SELinux blocks real-world exploits". Published 24 February 2008. Accessed 27 August 2009: http://www.networkworld.com/news/2008/022408-selinux.html.

SELinux decisions, such as allowing or disallowing access, are cached. This cache is known as the Access Vector Cache (AVC). When using these cached decisions, SELinux policy rules need to be checked less, which increases performance. Remember that SELinux policy rules have no effect if DAC rules deny access first.

1.4. SELinux States and Modes

SELinux can be either in the enabled or disabled state. When disabled, only DAC rules are used. When enabled, SELinux can run in one of the following modes:

- Enforcing: SELinux policy is enforced. SELinux denies access based on SELinux policy rules.
- Permissive: SELinux policy is not enforced. SELinux does not deny access, but denials are logged for actions that would have been denied if running in enforcing mode.

Use the **setenforce** utility to change between enforcing and permissive mode. Changes made with **setenforce** do not persist across reboots. To change to enforcing mode, as the Linux root user, run the **setenforce** 1 command. To change to permissive mode, run the **setenforce** 0 command. Use the **getenforce** utility to view the current SELinux mode:



Persistent states and modes changes are covered in Section 4.4, "Permanent Changes in SELinux States and Modes".

SELinux Contexts

Processes and files are labeled with an SELinux context that contains additional information, such as an SELinux user, role, type, and, optionally, a level. When running SELinux, all of this information is used to make access control decisions. In Fedora, SELinux provides a combination of Role-Based Access Control (RBAC), Type Enforcement (TE), and, optionally, Multi-Level Security (MLS).

The following is an example showing SELinux context. SELinux contexts are used on processes, Linux users, and files, on Linux operating systems that run SELinux. Use the following command to view the SELinux context of files and directories:

```
~]$ ls -Z file1
-rwxrw-r-- user1 group1 unconfined_u:object_r:user_home_t:s0 file1
```

SELinux contexts follow the SELinux user:role:type:level syntax. The fields are as follows:

SELinux user

The SELinux user identity is an identity known to the policy that is authorized for a specific set of roles, and for a specific MLS/MCS range. Each Linux user is mapped to an SELinux user via SELinux policy. This allows Linux users to inherit the restrictions placed on SELinux users. The mapped SELinux user identity is used in the SELinux context for processes in that session, in order to define what roles and levels they can enter. Run the following command as root to view a list of mappings between SELinux and Linux user accounts (you need to have the *policycoreutils-python* package installed):

```
~]# semanage login -l
Login Name SELinux User MLS/MCS Range Service

__default__ unconfined_u s0-s0:c0.c1023 *
root unconfined_u s0-s0:c0.c1023 *
system_u system_u s0-s0:c0.c1023 *
```

Output may differ slightly from system to system:

- The Login Name column lists Linux users.
- The SELinux User column lists which SELinux user the Linux user is mapped to. For processes, the SELinux user limits which roles and levels are accessible.
- The MLS/MCS Range column, is the level used by Multi-Level Security (MLS) and Multi-Category Security (MCS).
- The **Service** column determines the correct SELinux context, in which the Linux user is supposed to be logged in to the system. By default, the asterisk (*) character is used, which stands for any service.

role

Part of SELinux is the Role-Based Access Control (RBAC) security model. The role is an attribute of RBAC. SELinux users are authorized for roles, and roles are authorized for domains. The role serves as an intermediary between domains and SELinux users. The roles that can be entered determine which domains can be entered; ultimately, this controls which object types can be accessed. This helps reduce vulnerability to privilege escalation attacks.

type

The type is an attribute of Type Enforcement. The type defines a domain for processes, and a type for files. SELinux policy rules define how types can access each other, whether it be a domain accessing a type, or a domain accessing another domain. Access is only allowed if a specific SELinux policy rule exists that allows it.

level

The level is an attribute of MLS and MCS. An MLS range is a pair of levels, written as *lowlevel-highlevel* if the levels differ, or *lowlevel* if the levels are identical (s0-s0 is the same as s0). Each level is a sensitivity-category pair, with categories being optional. If there are categories, the level is written as *sensitivity:category-set*. If there are no categories, it is written as *sensitivity*.

If the category set is a contiguous series, it can be abbreviated. For example, c0.c3 is the same as c0, c1, c2, c3. The /etc/selinux/targeted/setrans.conf file maps levels (s0:c0) to human-readable form (that is CompanyConfidential). In Fedora, targeted policy enforces MCS, and in MCS, there is just one sensitivity, s0. MCS in Fedora supports 1024 different categories: c0 through to c1023. s0-s0:c0.c1023 is sensitivity s0 and authorized for all categories.

MLS enforces the Bell-La Padula Mandatory Access Model, and is used in Labeled Security Protection Profile (LSPP) environments. To use MLS restrictions, install the *selinux-policy-mls* package, and configure MLS to be the default SELinux policy. The MLS policy shipped with Fedora omits many program domains that were not part of the evaluated configuration, and therefore, MLS on a desktop workstation is unusable (no support for the X Window System); however, an MLS policy from the *upstream SELinux Reference Policy*¹ can be built that includes all program domains. For more information on MLS configuration, see *Section 4.12*, "Multi-Level Security (MLS)".

2.1. Domain Transitions

A process in one domain transitions to another domain by executing an application that has the entrypoint type for the new domain. The entrypoint permission is used in SELinux policy and controls which applications can be used to enter a domain. The following example demonstrates a domain transition:

Procedure 2.1. An Example of a Domain Transition

 A user wants to change their password. To do this, they run the passwd utility. The /usr/bin/ passwd executable is labeled with the passwd_exec_t type:

```
~]$ ls -Z /usr/bin/passwd
-rwsr-xr-x root root system_u:object_r:passwd_exec_t:s0 /usr/bin/passwd
```

The passwd utility accesses /etc/shadow, which is labeled with the shadow_t type:

```
~]$ ls -Z /etc/shadow
-r----. root root system_u:object_r:shadow_t:s0 /etc/shadow
```

2. An SELinux policy rule states that processes running in the passwd_t domain are allowed to read and write to files labeled with the shadow_t type. The shadow_t type is only applied to files

http://oss.tresys.com/projects/refpolicy

that are required for a password change. This includes /etc/gshadow, /etc/shadow, and their backup files.

- 3. An SELinux policy rule states that the passwd_t domain has entrypoint permission to the passwd_exec_t type.
- 4. When a user runs the passwd utility, the user's shell process transitions to the passwd_t domain. With SELinux, since the default action is to deny, and a rule exists that allows (among other things) applications running in the passwd_t domain to access files labeled with the shadow_t type, the passwd application is allowed to access /etc/shadow, and update the user's password.

This example is not exhaustive, and is used as a basic example to explain domain transition. Although there is an actual rule that allows subjects running in the passwd_t domain to access objects labeled with the shadow_t file type, other SELinux policy rules must be met before the subject can transition to a new domain. In this example, Type Enforcement ensures:

- The passwd_t domain can only be entered by executing an application labeled with the passwd_exec_t type; can only execute from authorized shared libraries, such as the lib_t type; and cannot execute any other applications.
- Only authorized domains, such as passwd_t, can write to files labeled with the shadow_t type.
 Even if other processes are running with superuser privileges, those processes cannot write to files labeled with the shadow_t type, as they are not running in the passwd_t domain.
- Only authorized domains can transition to the passwd_t domain. For example, the sendmail process running in the sendmail_t domain does not have a legitimate reason to execute passwd; therefore, it can never transition to the passwd_t domain.
- Processes running in the passwd_t domain can only read and write to authorized types, such as
 files labeled with the etc_t or shadow_t types. This prevents the passwd application from being
 tricked into reading or writing arbitrary files.

2.2. SELinux Contexts for Processes

Use the **ps** -eZ command to view the SELinux context for processes. For example:

Procedure 2.2. View the SELinux Context for the passwd Utility

- 1. Open a terminal, such as **Applications** \rightarrow **System Tools** \rightarrow **Terminal**.
- 2. Run the passwd utility. Do not enter a new password:

```
~]$ passwd
Changing password for user user_name.
Changing password for user_name.
(current) UNIX password:
```

3. Open a new tab, or another terminal, and run the following command. The output is similar to the following:

```
~]$ ps -eZ | grep passwd
unconfined_u:unconfined_r:passwd_t:s0-s0:c0.c1023 13212 pts/1 00:00:00 passwd
```

4. In the first tab/terminal, press Ctrl+C to cancel the passwd utility.

In this example, when the passwd utility (labeled with the passwd_exec_t type) is executed, the user's shell process transitions to the passwd_t domain. Remember that the type defines a domain for processes, and a type for files.

To view the SELinux contexts for all running processes, run the ps utility again. Note that below is a truncated example of the output, and may differ on your system:

The system_r role is used for system processes, such as daemons. Type Enforcement then separates each domain.

2.3. SELinux Contexts for Users

Use the following command to view the SELinux context associated with your Linux user:

```
~]$ id -Z
unconfined_u:unconfined_t:s0-s0:c0.c1023
```

In Fedora, Linux users run unconfined by default. This SELinux context shows that the Linux user is mapped to the SELinux unconfined_u user, running as the unconfined_r role, and is running in the unconfined_t domain. s0-s0 is an MLS range, which in this case, is the same as just s0. The categories the user has access to is defined by c0.c1023, which is all categories (c0 through to c1023).

Targeted Policy

Targeted policy is the default SELinux policy used in Fedora. When using targeted policy, processes that are targeted run in a confined domain, and processes that are not targeted run in an unconfined domain. For example, by default, logged-in users run in the unconfined_t domain, and system processes started by init run in the unconfined_service_t domain; both of these domains are unconfined.

Unconfined domains (as well as confined domains) are subject to executable and writeable memory checks. By default, subjects running in an unconfined domain cannot allocate writeable memory and execute it. This reduces vulnerability to buffer overflow attacks. These memory checks are disabled by setting Booleans, which allow the SELinux policy to be modified at runtime. Boolean configuration is discussed later.

3.1. Confined Processes

Almost every service that listens on a network, such as sshd or httpd, is confined in Fedora. Also, most processes that run as the root user and perform tasks for users, such as the passwd utility, are confined. When a process is confined, it runs in its own domain, such as the httpd process running in the httpd_t domain. If a confined process is compromised by an attacker, depending on SELinux policy configuration, an attacker's access to resources and the possible damage they can do is limited.

Complete this procedure to ensure that SELinux is enabled and the system is prepared to perform the following example:

Procedure 3.1. How to Verify SELinux Status

1. Confirm that SELinux is enabled, is running in enforcing mode, and that targeted policy is being used. The correct output should look similar to the output below:

```
~]$ sestatus

SELinux status: enabled

SELinuxfs mount: /selinux

Current mode: enforcing

Mode from config file: enforcing

Policy version: 24

Policy from config file: targeted
```

See Section 4.4, "Permanent Changes in SELinux States and Modes" for detailed information about changing SELinux modes.

2. As root, create a file in the /var/www/html/ directory:

```
~]# touch /var/www/html/testfile
```

3. Run the following command to view the SELinux context of the newly created file:

```
~]$ ls -Z /var/www/html/testfile
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 /var/www/html/
testfile
```

By default, Linux users run unconfined in Fedora, which is why the **testfile** file is labeled with the SELinux unconfined_u user. RBAC is used for processes, not files. Roles do not have a meaning for files; the object_r role is a generic role used for files (on persistent storage

and network file systems). Under the **/proc/** directory, files related to processes may use the system_r role. The httpd_sys_content_t type allows the httpd process to access this file.

The following example demonstrates how SELinux prevents the Apache HTTP Server (httpd) from reading files that are not correctly labeled, such as files intended for use by Samba. This is an example, and should not be used in production. It assumes that the *httpd* and *wget* packages are installed, the SELinux targeted policy is used, and that SELinux is running in enforcing mode.

Procedure 3.2. An Example of Confined Process

1. As root, start the httpd daemon:

```
~]# systemctl start httpd.service
```

Confirm that the service is running. The output should include the information below (only the time stamp will differ):

```
~]$ systemctl status httpd.service
httpd.service - The Apache HTTP Server
Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled)
Active: active (running) since Mon 2013-08-05 14:00:55 CEST; 8s ago
```

2. Change into a directory where your Linux user has write access to, and run the following command. Unless there are changes to the default configuration, this command succeeds:

3. The **chcon** command relabels files; however, such label changes do not survive when the file system is relabeled. For permanent changes that survive a file system relabel, use the semanage utility, which is discussed later. As root, run the following command to change the type to a type used by Samba:

```
~]# chcon -t samba_share_t /var/www/html/testfile
```

Run the following command to view the changes:

```
~]$ ls -Z /var/www/html/testfile
-rw-r--r-- root root unconfined_u:object_r:samba_share_t:s0 /var/www/html/testfile
```

4. Note that the current DAC permissions allow the httpd process access to **testfile**. Change into a directory where your user has write access to, and run the following command. Unless there are changes to the default configuration, this command fails:

```
~]$ wget http://localhost/testfile
--2009-11-06 14:11:23-- http://localhost/testfile
Resolving localhost... 127.0.0.1
Connecting to localhost|127.0.0.1|:80... connected.
HTTP request sent, awaiting response... 403 Forbidden
2009-11-06 14:11:23 ERROR 403: Forbidden.
```

5. As root, remove **testfile**:

```
~]# rm -i /var/www/html/testfile
```

6. If you do not require httpd to be running, as root, run the following command to stop it:

```
~]# systemctl stop httpd.service
```

This example demonstrates the additional security added by SELinux. Although DAC rules allowed the httpd process access to **testfile** in step 2, because the file was labeled with a type that the httpd process does not have access to, SELinux denied access.

If the auditd daemon is running, an error similar to the following is logged to **/var/log/audit/audit.log**:

```
type=AVC msg=audit(1220706212.937:70): avc: denied { getattr } for pid=1904 comm="httpd"
path="/var/www/html/testfile" dev=sda5 ino=247576 scontext=unconfined_u:system_r:httpd_t:s0
tcontext=unconfined_u:object_r:samba_share_t:s0 tclass=file

type=SYSCALL msg=audit(1220706212.937:70): arch=40000003 syscall=196 success=no exit=-13
a0=b9e21da0 a1=bf9581dc a2=555ff4 a3=2008171 items=0 ppid=1902 pid=1904 auid=500 uid=48
gid=48 euid=48 suid=48 fsuid=48 egid=48 sgid=48 fsgid=48 tty=(none) ses=1 comm="httpd"
exe="/usr/sbin/httpd" subj=unconfined_u:system_r:httpd_t:s0 key=(null)
```

Also, an error similar to the following is logged to /var/log/httpd/error_log:

```
[Wed May 06 23:00:54 2009] [error] [client 127.0.0.1] (13)Permission denied: access to / testfile denied
```

3.2. Unconfined Processes

Unconfined processes run in unconfined domains, for example, unconfined services executed by init end up running in the unconfined_service_t domain, unconfined services executed by kernel end up running in the kernel_t domain, and unconfined services executed by unconfined Linux users end up running in the unconfined_t domain. For unconfined processes, SELinux policy rules are applied, but policy rules exist that allow processes running in unconfined domains almost all access. Processes running in unconfined domains fall back to using DAC rules exclusively. If an unconfined process is compromised, SELinux does not prevent an attacker from gaining access to system resources and data, but of course, DAC rules are still used. SELinux is a security enhancement on top of DAC rules — it does not replace them.

To ensure that SELinux is enabled and the system is prepared to perform the following example, complete the *Procedure 3.1*, "How to Verify SELinux Status" described in Section 3.1, "Confined Processes".

The following example demonstrates how the Apache HTTP Server (httpd) can access data intended for use by Samba, when running unconfined. Note that in Fedora, the httpd process runs in

the confined httpd_t domain by default. This is an example, and should not be used in production. It assumes that the *httpd*, *wget*, *dbus* and *audit* packages are installed, that the SELinux targeted policy is used, and that SELinux is running in enforcing mode.

Procedure 3.3. An Example of Unconfined Process

 The **chcon** command relabels files; however, such label changes do not survive when the file system is relabeled. For permanent changes that survive a file system relabel, use the semanage utility, which is discussed later. As the root user, run the following command to change the type to a type used by Samba:

```
~]# chcon -t samba_share_t /var/www/html/testfile
```

View the changes:

```
~]$ ls -Z /var/www/html/testfile
-rw-r--r-- root root unconfined_u:object_r:samba_share_t:s0 /var/www/html/testfile
```

Run the following command to confirm that the httpd process is not running:

```
~]$ systemctl status httpd.service
httpd.service - The Apache HTTP Server
Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled)
Active: inactive (dead)
```

If the output differs, run the following command as root to stop the httpd process:

```
~]# systemctl stop httpd.service
```

3. To make the httpd process run unconfined, run the following command as root to change the type of the /usr/sbin/httpd file, to a type that does not transition to a confined domain:

```
~]# chcon -t bin_t /usr/sbin/httpd
```

4. Confirm that /usr/sbin/httpd is labeled with the bin_t type:

```
~]$ ls -Z /usr/sbin/httpd
-rwxr-xr-x. root root system_u:object_r:bin_t:s0 /usr/sbin/httpd
```

5. As root, start the httpd process and confirm, that it started successfully:

```
~]# systemctl start httpd.service
```

```
~]# systemctl status httpd.service
httpd.service - The Apache HTTP Server
Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled)
Active: active (running) since Thu 2013-08-15 11:17:01 CEST; 5s ago
```

6. Run the following command to view httpd running in the unconfined_service_t domain:

```
~]$ ps -eZ | grep httpd
```

```
system_u:system_r:unconfined_service_t:s0 11884 ? 00:00:00 httpd
system_u:system_r:unconfined_service_t:s0 11885 ? 00:00:00 httpd
system_u:system_r:unconfined_service_t:s0 11886 ? 00:00:00 httpd
system_u:system_r:unconfined_service_t:s0 11887 ? 00:00:00 httpd
system_u:system_r:unconfined_service_t:s0 11888 ? 00:00:00 httpd
system_u:system_r:unconfined_service_t:s0 11889 ? 00:00:00 httpd
```

7. Change into a directory where your Linux user has write access to, and run the following command. Unless there are changes to the default configuration, this command succeeds:

Although the httpd process does not have access to files labeled with the samba_share_t type, httpd is running in the unconfined unconfined_service_t domain, and falls back to using DAC rules, and as such, the wget command succeeds. Had httpd been running in the confined httpd_t domain, the wget command would have failed.

8. The restorecon utility restores the default SELinux context for files. As root, run the following command to restore the default SELinux context for /usr/sbin/httpd:

```
~]# restorecon -v /usr/sbin/httpd
restorecon reset /usr/sbin/httpd context system_u:object_r:unconfined_exec_t:s0-
>system_u:object_r:httpd_exec_t:s0
```

Confirm that /usr/sbin/httpd is labeled with the httpd_exec_t type:

```
~]$ ls -Z /usr/sbin/httpd
-rwxr-xr-x root root system_u:object_r:httpd_exec_t:s0 /usr/sbin/httpd
```

9. As root, run the following command to restart httpd. After restarting, confirm that httpd is running in the confined httpd_t domain:

```
~]# systemctl restart httpd.service
```

```
~]$ ps -eZ | grep httpd
system_u:system_r:httpd_t:s0
                               8883 ?
                                             00:00:00 httpd
system_u:system_r:httpd_t:s0
                               8884 ?
                                             00:00:00 httpd
                               8885 ?
                                             00:00:00 httpd
system_u:system_r:httpd_t:s0
                               8886 ?
system_u:system_r:httpd_t:s0
                                             00:00:00 httpd
                               8887 ?
system_u:system_r:httpd_t:s0
                                             00:00:00 httpd
system_u:system_r:httpd_t:s0
                               8888 ?
                                             00:00:00 httpd
system_u:system_r:httpd_t:s0
                               8889 ?
                                             00:00:00 httpd
```

10. As root, remove **testfile**:

```
~]# rm -i /var/www/html/testfile
rm: remove regular empty file `/var/www/html/testfile'? y
```

11. If you do not require httpd to be running, as root, run the following command to stop httpd:

```
~]# systemctl stop httpd.service
```

The examples in these sections demonstrate how data can be protected from a compromised confined-process (protected by SELinux), as well as how data is more accessible to an attacker from a compromised unconfined-process (not protected by SELinux).

3.3. Confined and Unconfined Users

Each Linux user is mapped to an SELinux user using SELinux policy. This allows Linux users to inherit the restrictions on SELinux users. This Linux user mapping is seen by running the **semanage login**-1 command as root:

```
~]# semanage login -l

Login Name SELinux User MLS/MCS Range Service

__default__ unconfined_u s0-s0:c0.c1023 *
root unconfined_u s0-s0:c0.c1023 *
system_u system_u s0-s0:c0.c1023 *
```

In Fedora, Linux users are mapped to the SELinux <u>__default__</u> login by default, which is mapped to the SELinux unconfined_u user. The following line defines the default mapping:

```
__default__ unconfined_u s0-s0:c0.c1023
```

The following procedure demonstrates how to add a new Linux user to the system and how to map that user to the SELinux unconfined_u user. It assumes that the root user is running unconfined, as it does by default in Fedora:

Procedure 3.4. Mapping a New Linux User to the SELinux unconfined_u User

1. As root, run the following command to create a new Linux user named newuser:

```
~]# useradd newuser
```

2. To assign a password to the Linux newuser user. Run the following command as root:

```
~]# passwd newuser
Changing password for user newuser.
New UNIX password: Enter a password
Retype new UNIX password: Enter the same password again
passwd: all authentication tokens updated successfully.
```

3. Log out of your current session, and log in as the Linux newuser user. When you log in, the pam_selinux PAM module automatically maps the Linux user to an SELinux user (in this case, unconfined_u), and sets up the resulting SELinux context. The Linux user's shell is then launched with this context. Run the following command to view the context of a Linux user:

[newuser@localhost ~]\$ id -Z
unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023



Note

If you no longer need the **newuser** user on your system, log out of the Linux **newuser**'s session, log in with your account, and run the **userdel -r newuser** command as root. It will remove **newuser** along with their home directory.

Confined and unconfined Linux users are subject to executable and writeable memory checks, and are also restricted by MCS or MLS.

If an unconfined Linux user executes an application that SELinux policy defines as one that can transition from the unconfined_t domain to its own confined domain, the unconfined Linux user is still subject to the restrictions of that confined domain. The security benefit of this is that, even though a Linux user is running unconfined, the application remains confined. Therefore, the exploitation of a flaw in the application can be limited by the policy.

Similarly, we can apply these checks to confined users. However, each confined Linux user is restricted by a confined user domain against the unconfined_t domain. The SELinux policy can also define a transition from a confined user domain to its own target confined domain. In such a case, confined Linux users are subject to the restrictions of that target confined domain. The main point is that special privileges are associated with the confined users according to their role. In the table below, you can see examples of basic confined domains for Linux users in Fedora:

T-1-1-1	0.4	\circ		to the least of	On the second second	
Table	3.1.	SEL	ınux	User	Capabilities	

User	Role	Domain	X Window System	su or sudo	Execute in home directory and /tmp/ (default)	Networking
sysadm_u	sysadm_r	sysadm_t	yes	su and sudo	yes	yes
staff_u	staff_r	staff_t	yes	only sudo	yes	yes
user_u	user_r	user_t	yes	no	yes	yes
guest_u	guest_r	guest_t	no	no	no	no
xguest_u	xguest_r	xguest_t	yes	no	no	Firefox only

- Linux users in the user_t, guest_t, and xguest_t domains can only run set user ID (setuid) applications if SELinux policy permits it (for example, passwd). These users cannot run the **su** and **sudo** setuid applications, and therefore cannot use these applications to become root.
- Linux users in the sysadm_t, staff_t, user_t, and xguest_t domains can log in via the X Window System and a terminal.
- By default, Linux users in the guest_t and xguest_t domains cannot execute applications in their home directories or the /tmp/ directory, preventing them from executing applications, which inherit

users' permissions, in directories they have write access to. This helps prevent flawed or malicious applications from modifying users' files.

- By default, Linux users in the staff_t and user_t domains can execute applications in their home directories and /tmp/. See Section 6.6, "Booleans for Users Executing Applications" for information about allowing and preventing users from executing applications in their home directories and /tmp/.
- The only network access Linux users in the xguest_t domain have is Firefox connecting to web pages.

Alongside with the already mentioned SELinux users, there are special roles, that can be mapped to those users. These roles determine what SELinux allows the user to do:

- webadm_r can only administrate SELinux types related to the Apache HTTP Server. See Section 13.2, "Types" for further information.
- dbadm_r can only administrate SELinux types related to the MariaDB database and the PostgreSQL database management system. See Section 20.2, "Types" and Section 21.2, "Types" for further information.
- logadm_r can only administrate SELinux types related to the syslog and auditlog processes.
- secadm_r can only administrate SELinux.
- auditadm_r can only administrate processes related to the audit subsystem.

To list all available roles, run the following command:

```
~]$ seinfo -r
```

Note that the **seinfo** command is provided by the *setools-console* package, which is not installed by default.

3.3.1. The sudo Transition and SELinux Roles

In certain cases, confined users need to perform an administrative task that require root privileges. To do so, such a confined user has to gain a *confined administrator* SELinux role using the **sudo** command. The **sudo** command is used to give trusted users administrative access. When users precede an administrative command with **sudo**, they are prompted for their *own* password. Then, when they have been authenticated and assuming that the command is permitted, the administrative command is executed as if they were the root user.

As shown in *Table 3.1, "SELinux User Capabilities"*, only the staff_u and sysadm_u SELinux confined users are permitted to use **sudo** by default. When such users execute a command with **sudo**, their role is changed based on the rules specified in the /etc/sudoers configuration file or in a respective file in the /etc/sudoers.d/ directory if such a file exists.

For more information about **sudo**, see *Fedora System Administrator's Guide* available at *http://docs.fedoraproject.org/*.

Procedure 3.5. Configuring the sudo Transition

This procedure shows how to set up **sudo** to transition a newly-created *SELinux_user_u* confined user to a *administrator_r* confined administrator. To configure a confined administrator role for an already existing SELinux user, skip the first two steps. Also note that the following commands must be run as the root user.

1. Create a new SELinux user and specify the default SELinux role and a supplementary confined administrator role for this user:

```
~]# semanage user -a -r s0-s0:c0.c1023 -R
"default_role_r administrator_r" SELinux_user_u
```

In the example below, the default role of the newly-created confined_u SELinux user is staff_r and the confined administrator role is webadm_r:

```
~]# semanage user -a -r s0-s0:c0.c1023 -R "staff_r webadm_r" confined_u
```

Set up the default SElinux policy context file. For example, to have the same SELinux rules as the staff_u SELinux user, copy the staff_u context file:

```
\sim]# cp /etc/selinux/targeted/contexts/users/staff_u /etc/selinux/targeted/contexts/users/SELinux_user
```

3. Map the newly-created SELinux user to an existing Linux user:

```
~]# semanage login -a -s SELinux_user_u -rs0:c0.c1023 linux_user
```

4. Create a new configuration file with the same name as your Linux user in the /etc/sudoers.d/ directory and add the following string to it:

```
~]# echo "linux_user ALL=(ALL) TYPE=administaror_t ROLE=administrator_r /bin/sh " > / etc/sudoers.d/linux_user
```

For example:

```
~]# echo "linux_user ALL=(ALL) TYPE=webadm_t ROLE=webadm_r /bin/sh " > /etc/sudoers.d/linux_user
```

5. Use the restorecon utility to relabel the *linux_user* home directory:

```
~]# restorecon -R -v /home/linux_user
```

6. Reboot the system:

```
~]# systemctl reboot
```

7. When you log in to the system as the newly-created Linux user, the user is labeled with the default SELinux role:

```
~]$ id -Z
SELinux_user_u:default_role_r:default_role_t:s0:c0.c1023
```

After running **sudo**, the user's SELinux context changes to the supplementary SELinux role as specified in **/etc/sudoers.d/linux_user**. The **-i** option used with **sudo** caused that an interactive shell is executed:

```
~]$ sudo -i
```

```
~]# id -Z
SELinux_user_u:administrator_r:administrator_t:s0-s0:c0.c1023
```

For the confined_u SELinux user from the example specified in the first step the output looks like below:

```
~]$ id -Z
confined_u:staff_r:staff_t:s0:c0.c1023
~]$ sudo -i
~]# id -Z
confined_u:webadm_r:webadm_t:s0:c0.c1023
```

Working with SELinux

The following sections give a brief overview of the main SELinux packages in Fedora; installing and updating packages; which log files are used; the main SELinux configuration file; enabling and disabling SELinux; SELinux modes; configuring Booleans; temporarily and persistently changing file and directory labels; overriding file system labels with the **mount** command; mounting NFS volumes; and how to preserve SELinux contexts when copying and archiving files and directories.

4.1. SELinux Packages

In Fedora full installation, the SELinux packages are installed by default unless they are manually excluded during installation. If performing a minimal installation in text mode, the *policycoreutils-python* and the *policycoreutils-gui* package are not installed by default. Also, by default, SELinux runs in enforcing mode and the SELinux targeted policy is used. The following SELinux packages are installed on your system by default:

- policycoreutils provides utilities such as restorecon, secon, setfiles, semodule, load_policy, and setsebool, for operating and managing SELinux.
- selinux-policy provides configuration for the SELinux Reference policy. The SELinux Reference
 Policy is a complete SELinux policy, and is used as a basis for other policies, such as the SELinux
 targeted policy; see the Tresys Technology SELinux Reference Policy¹ page for further information.
 This package contains the selinux-policy.conf file and RPM macros.
- · selinux-policy-targeted provides the SELinux targeted policy.
- libselinux provides an API for SELinux applications.
- *libselinux-utils* provides the avcstat, getenforce, getsebool, matchpathcon, selinuxconlist, selinuxdefcon, selinuxenabled, and setenforce utilities.
- libselinux-python provides Python bindings for developing SELinux applications.

The following packages are not installed by default but can be optionally installed by running the **dnf install** <**package-name>** command:

- selinux-policy-devel provides utilities for creating a custom SELinux policy and policy modules. It also contains manual pages that describe how to configure SELinux altogether with various services.
- selinux-policy-mls provides the MLS (Multi-Level Security) SELinux policy.
- setroubleshoot-server translates denial messages, produced when access is denied by SELinux, into detailed descriptions that can be viewed with the **sealert** utility, also provided in this package.
- setools-console provides the *Tresys Technology SETools distribution*², a number of utilities and libraries for analyzing and querying policy, audit log monitoring and reporting, and file context management. The setools package is a meta-package for SETools. The setools-gui package provides the apol and seaudit utilities. The setools-console package provides the sechecker, sediff, seinfo, sesearch, and findcon command-line utilities. See the *Tresys Technology SETools*³ page for information about these utilities.

¹ http://oss.tresys.com/projects/refpolicy

² http://oss.tresys.com/projects/setools

³ http://oss.tresys.com/projects/setools

- mcstrans translates levels, such as s0-s0:c0.c1023, to a form that is easier to read, such as SystemLow-SystemHigh.
- policycoreutils-python provides utilities such as **semanage**, **audit2allow**, **audit2why**, and **chcat**, for operating and managing SELinux.
- policycoreutils-gui provides system-config-selinux, a graphical utility for managing SELinux.

4.2. Which Log File is Used

In Fedora, the *dbus* and *audit* packages are installed by default, unless they are removed from the default package selection. The *setroubleshoot-server* must be installed via DNF (use the **dnf install setroubleshoot** command).

If the auditd daemon is running, an SELinux denial message, such as the following, is written to / var/log/audit/audit.log by default:

```
type=AVC msg=audit(1223024155.684:49): avc: denied { getattr } for pid=2000 comm="httpd" path="/var/www/html/file1" dev=dm-0 ino=399185 scontext=unconfined_u:system_r:httpd_t:s0 tcontext=system_u:object_r:samba_share_t:s0 tclass=file
```

In addition, a message similar to the one below is written to the /var/log/message file:

```
May 7 18:55:56 localhost setroubleshoot: SELinux is preventing httpd (httpd_t) "getattr" to /var/www/html/file1 (samba_share_t). For complete SELinux messages. run sealert -l de7e30d6-5488-466d-a606-92c9f40d316d
```

In Fedora 24, setroubleshootd no longer constantly runs as a service. However, it is still used to analyze the AVC messages. Two new programs act as a method to start setroubleshoot when needed:

- The sedispatch utility runs as a part of the audit subsystem. When an AVC denial message is returned, sedispatch sends a message using dbus. These messages go straight to setroubleshootd if it is already running. If it is not running, sedispatch starts it automatically.
- The seapplet utility runs in the system toolbar, waiting for dbus messages in setroubleshootd. It launches the notification bubble, allowing the user to review AVC messages.

Procedure 4.1. Starting Daemons Automatically

1. To configure the auditd and rsyslog daemons to automatically start at boot, run the following commands as the root user:

```
~]# systemctl enable auditd.service

~]# systemctl enable rsyslog.service
```

2. To ensure that the daemons are enabled, type the following commands at the shell prompt:

```
~]$ systemctl is-enabled auditd enabled
```

```
~]$ systemctl is-enabled rsyslog
enabled
```

Alternatively, use the **systemctl status service-name**.**service** command and search for the keyword **enabled** in the command output, for example:

```
~]$ systemctl status auditd.service | grep enabled
auditd.service - Security Auditing Service
Loaded: loaded (/usr/lib/systemd/system/auditd.service; enabled)
```

To learn more on how the systemd daemon manages system services, see the *Managing System Services*⁴ chapter in the *System Administrator's Guide*⁵.

4.3. Main Configuration File

The /etc/selinux/config file is the main SELinux configuration file. It controls whether SELinux is enabled or disabled and which SELinux mode and SELinux policy is used:

```
# This file controls the state of SELinux on the system.
# SELINUX= can take one of these three values:
# enforcing - SELinux security policy is enforced.
# permissive - SELinux prints warnings instead of enforcing.
# disabled - No SELinux policy is loaded.
SELINUX=enforcing
# SELINUXTYPE= can take one of these two values:
# targeted - Targeted processes are protected,
# mls - Multi Level Security protection.
SELINUXTYPE=targeted
```

SELINUX=

The **SELINUX** option sets whether SELinux is disabled or enabled and in which mode - enforcing or permissive - it is running:

- When using **SELINUX=enforcing**, SELinux policy is enforced, and SELinux denies access based on SELinux policy rules. Denial messages are logged.
- When using SELINUX=permissive, SELinux policy is not enforced. SELinux does not deny
 access, but denials are logged for actions that would have been denied if running SELinux in
 enforcing mode.
- When using **SELINUX=disabled**, SELinux is disabled, the SELinux module is not registered with the Linux kernel, and only DAC rules are used.

SELINUXTYPE=

The **SELINUXTYPE** option sets the SELinux policy to use. Targeted policy is the default policy. Only change this option if you want to use the MLS policy. For information on how to enable the MLS policy, see *Section 4.12.2, "Enabling MLS in SELinux"*.

⁴ https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/System_Administrators_Guide/sect-Managing_Services_with_systemd-Services.html

https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/System_Administrators_Guide/index.html



Important

When systems run with SELinux in permissive or disabled mode, users have permission to label fies incorrectly. Also, files created while SELinux is disabled are not labeled. This causes problems when changing to enforcing mode. To prevent incorrectly labeled and unlabeled files from causing problems, file systems are automatically relabeled when changing from disabled mode to permissive or enforcing mode.

4.4. Permanent Changes in SELinux States and Modes

As discussed in *Section 1.4, "SELinux States and Modes"*, SELinux can be enabled or disabled. When enabled, SELinux has two modes: enforcing and permissive.

Use the **getenforce** or **sestatus** commands to check in which mode SELinux is running. The **getenforce** command returns **Enforcing**, **Permissive**, or **Disabled**.

The **sestatus** command returns the SELinux status and the SELinux policy being used:

~]\$ sestatus

SELinux status: enabled

/sys/fs/selinux SELinuxfs mount: SELinux root directory: /etc/selinux Loaded policy name: targeted Current mode: enforcing Mode from config file: enforcing Policy MLS status: enabled Policy deny_unknown status: allowed Max kernel policy version: 30



Note

When systems run SELinux in permissive mode, users are able to label files incorrectly. Files created while SELinux is disabled are not labeled at all. This behavior causes problems when changing to enforcing mode because files are labeled incorrectly or are not labeled at all. To prevent incorrectly labeled and unlabeled files from causing problems, file systems are automatically relabeled when changing from the disabled state to permissive or enforcing mode.

4.4.1. Enabling SELinux

When enabled, SELinux can run in one of two modes: enforcing or permissive. The following sections show how to permanently change into these modes. On systems with SELinux disabled, the **SELINUX=disabled** option is configured in **/etc/selinux/config**:

```
# This file controls the state of SELinux on the system.

# SELINUX= can take one of these three values:

# enforcing - SELinux security policy is enforced.

# permissive - SELinux prints warnings instead of enforcing.

# disabled - No SELinux policy is loaded.

SELINUX=disabled
```

```
# SELINUXTYPE= can take one of these three values:

# targeted - Targeted processes are protected,

# minimum - Modification of targeted policy. Only selected processes are protected.

# mls - Multi Level Security protection.

SELINUXTYPE=targeted
```

Also, the **getenforce** command returns Disabled:

```
~]$ getenforce
Disabled
```

This guide assumes that the following packages are installed:

- · selinux-policy-targeted
- · selinux-policy
- libselinux
- · libselinux-python
- libselinux-utils
- · policycoreutils
- policycoreutils-python
- setroubleshoot
- setroubleshoot-server
- · setroubleshoot-plugins

To confirm that the aforementioned packages are installed, use the rpm utility:

```
~]$ rpm -qa | grep selinux
selinux-policy-3.12.1-136.el7.noarch
libselinux-2.2.2-4.el7.x86_64
selinux-policy-targeted-3.12.1-136.el7.noarch
libselinux-utils-2.2.2-4.el7.x86_64
libselinux-python-2.2.2-4.el7.x86_64
```

```
~]$ rpm -qa | grep policycoreutils
policycoreutils-2.2.5-6.el7.x86_64
policycoreutils-python-2.2.5-6.el7.x86_64
```

```
~]$ rpm -qa | grep setroubleshoot
setroubleshoot-server-3.2.17-2.el7.x86_64
setroubleshoot-3.2.17-2.el7.x86_64
setroubleshoot-plugins-3.0.58-2.el7.noarch
```

If they are not installed, use the DNF utility as root to install them:

```
~]# dnf install package_name
```

The following packages are optional:

- · policycoreutils-gui
- · setroubleshoot
- mcstrans

Following procedure shows how to enable SELinux:

Procedure 4.2. Enabling SELinux

1. Edit the /etc/selinux/config file as follows:

```
# This file controls the state of SELinux on the system.
# SELINUX= can take one of these three values:
# enforcing - SELinux security policy is enforced.
# permissive - SELinux prints warnings instead of enforcing.
# disabled - No SELinux policy is loaded.
SELINUX=enforcing
# SELINUXTYPE= can take one of these two values:
# targeted - Targeted processes are protected,
# mls - Multi Level Security protection.
SELINUXTYPE=targeted
```

2. Reboot the system:

```
~]# reboot
```

On the next boot, SELinux relabels all the files and directories within the system and adds SELinux context for files and directories that were created when SELinux was disabled.



Note

After changing to enforcing mode, SELinux may deny some actions because of incorrect or missing SELinux policy rules. To view what actions SELinux denies, run the following command as root:

```
~]# journalctl | grep "SELinux is preventing"
```

If SELinux denies some actions, see *Chapter 10, Troubleshooting* for information about troubleshooting.

Temporary changes in modes are covered in Section 1.4, "SELinux States and Modes".

4.4.2. Permissive Mode

When SELinux is running in permissive mode, SELinux policy is not enforced. The system remains operational and SELinux does not deny any operations but only logs AVC messages, which can be then used for troubleshooting, debugging, and SELinux policy improvements.

To permanently change mode to permissive, follow the procedure below:

Procedure 4.3. Changing to Permissive Mode

1. Edit the /etc/selinux/config file as follows:

```
# This file controls the state of SELinux on the system.

# SELINUX= can take one of these three values:

# enforcing - SELinux security policy is enforced.

# permissive - SELinux prints warnings instead of enforcing.

# disabled - No SELinux policy is loaded.

SELINUX=permissive

# SELINUXTYPE= can take one of these two values:

# targeted - Targeted processes are protected,

# mls - Multi Level Security protection.

SELINUXTYPE=targeted
```

2. Reboot the system:

```
~]# reboot
```

Temporary changes in modes are covered in Section 1.4, "SELinux States and Modes".

4.5. Disabling SELinux

When SELinux is disabled, SELinux policy is not loaded at all; it is not enforced and AVC messages are not logged. Therefore, all benefits of running SELinux listed in Section 1.1, "Benefits of running SELinux" are lost.



Important

It is strongly recommended to use permissive mode instead of permanently disabling SELinux. See *Section 4.4.2, "Permissive Mode"* for more information about permissive mode.

To permanently disable SELinux, follow the procedure below:

Procedure 4.4. Disabling SELinux

1. Configure **SELINUX=disabled** in the **/etc/selinux/config** file:

```
# This file controls the state of SELinux on the system.

# SELINUX= can take one of these three values:

# enforcing - SELinux security policy is enforced.

# permissive - SELinux prints warnings instead of enforcing.

# disabled - No SELinux policy is loaded.

SELINUX=disabled

# SELINUXTYPE= can take one of these two values:

# targeted - Targeted processes are protected,

# mls - Multi Level Security protection.

SELINUXTYPE=targeted
```

2. Reboot your system. After reboot, confirm that the **getenforce** command returns **Disabled**:

```
~]$ getenforce
Disabled
```

4.6. Booleans

Booleans allow parts of SELinux policy to be changed at runtime, without any knowledge of SELinux policy writing. This allows changes, such as allowing services access to NFS volumes, without reloading or recompiling SELinux policy.

4.6.1. Listing Booleans

For a list of Booleans, an explanation of what each one is, and whether they are on or off, run the **semanage boolean -1** command as the Linux root user. The following example does not list all Booleans and the output is shortened for brevity:

```
~]# semanage boolean -1
SELinux boolean State Default Description

ftp_home_dir (off , off) Determine whether ftpd can read...
smartmon_3ware (off , off) Determine whether smartmon can...
mpd_enable_homedirs (off , off) Determine whether mpd can traverse...
```

The **SELinux boolean** column lists Boolean names. The **Description** column lists whether the Booleans are on or off, and what they do.

In the following example, the ftp_home_dir Boolean is off, preventing the FTP daemon (vsftpd) from reading and writing to files in user home directories:

```
ftp_home_dir (off , off) Determine whether ftpd can read...
```

The **getsebool** -a command lists Booleans, whether they are on or off, but does not give a description of each one. The following example does not list all Booleans:

```
~]$ getsebool -a
cvs_read_shadow --> off
daemons_dump_core --> on
ftp_home_dir --> off
```

Run the **getsebool boolean-name** command to only list the status of the **boolean-name** Boolean:

```
~]$ getsebool cvs_read_shadow cvs_read_shadow --> off
```

Use a space-separated list to list multiple Booleans:

```
~]$ getsebool cvs_read_shadow daemons_dump_core ftp_home_dir
cvs_read_shadow --> off
daemons_dump_core --> on
ftp_home_dir --> off
```

4.6.2. Configuring Booleans

Run the setsebool utility in the **setsebool boolean_name on/off** form to enable or disable Booleans.

The following example demonstrates configuring the httpd_can_network_connect_db Boolean:

Procedure 4.5. Configuring Booleans

1. By default, the httpd_can_network_connect_db Boolean is off, preventing Apache HTTP Server scripts and modules from connecting to database servers:

```
~]$ getsebool httpd_can_network_connect_db
httpd_can_network_connect_db --> off
```

2. To temporarily enable Apache HTTP Server scripts and modules to connect to database servers, run the following command as root:

```
~]# setsebool httpd_can_network_connect_db on
```

3. Use the getsebool utility to verify the Boolean has been enabled:

```
~]$ getsebool httpd_can_network_connect_db
httpd_can_network_connect_db --> on
```

This allows Apache HTTP Server scripts and modules to connect to database servers.

4. This change is not persistent across reboots. To make changes persistent across reboots, run the **setsebool** -**P boolean-name on** command as root:

```
~]# setsebool -P httpd_can_network_connect_db on
```

4.6.3. Shell Auto-Completion

It is possible to use shell auto-completion with the getsebool, setsebool, and semanage utilities. Use the auto-completion with getsebool and setsebool to complete both command-line parameters and Booleans. To list only the command-line parameters, add the hyphen character ("-") after the command name and hit the **Tab** key:

```
~]# setsebool -[Tab]
-P
```

To complete a Boolean, start writing the Boolean name and then hit Tab:

```
~]$ getsebool samba_[Tab]
samba_create_home_dirs samba_export_all_ro samba_run_unconfined
samba_domain_controller samba_export_all_rw samba_share_fusefs
samba_enable_home_dirs samba_portmapper samba_share_nfs
```

```
~]# setsebool -P virt_use_[Tab]
virt_use_comm virt_use_nfs virt_use_sanlock
virt_use_execmem virt_use_rawip virt_use_usb
virt_use_fusefs virt_use_samba virt_use_xserver
```

⁶ To temporarily revert to the default behavior, as the Linux root user, run the **setsebool httpd_can_network_connect_db off** command. For changes that persist across reboots, run the **setsebool -P httpd_can_network_connect_db off** command.

The semanage utility is used with several command-line arguments that are completed one by one. The first argument of a **semanage** command is an option, which specifies what part of SELinux policy is managed:

```
~]# semanage [Tab]
boolean export import login node port
dontaudit fcontext interface module permissive user
```

Then, one or more command-line parameters follow:

```
~]# semanage fcontext -[Tab]
                      --equal
-a
            - D
                                     --help
                                                  – m
                                                               -0
            --delete
                        -f
                                                  --modify
                                                               -S
--add
                                     -1
-C
            --deleteall --ftype
                                     --list
                                                  -n
                                                               -t
-d
                         -h
                                     --locallist --noheading
                                                               --type
```

Finally, complete the name of a particular SELinux entry, such as a Boolean, SELinux user, domain, or another. Start typing the entry and hit **Tab**:

Command-line parameters can be chained in a command:

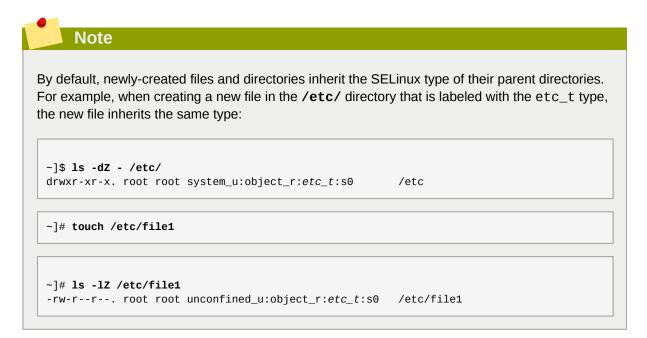
```
~]# semanage port -a -t http_port_t -p tcp 81
```

4.7. SELinux Contexts - Labeling Files

On systems running SELinux, all processes and files are labeled in a way that represents security-relevant information. This information is called the SELinux context. For files, this is viewed using the **1s** -**Z** command:

```
~]$ ls -Z file1
-rw-rw-r-- user1 group1 unconfined_u:object_r:user_home_t:s0 file1
```

In this example, SELinux provides a user (unconfined_u), a role (object_r), a type (user_home_t), and a level (s0). This information is used to make access control decisions. On DAC systems, access is controlled based on Linux user and group IDs. SELinux policy rules are checked after DAC rules. SELinux policy rules are not used if DAC rules deny access first.



There are multiple commands for managing the SELinux context for files, such as **chcon**, **semanage fcontext**, and **restorecon**.

4.7.1. Temporary Changes: chcon

The **chcon** command changes the SELinux context for files. However, changes made with the **chcon** command do not survive a file system relabel, or the execution of the **restorecon** command. SELinux policy controls whether users are able to modify the SELinux context for any given file. When using **chcon**, users provide all or part of the SELinux context to change. An incorrect file type is a common cause of SELinux denying access.

Quick Reference

• Run the **chcon** -t *type file-name* command to change the file type, where *type* is an SELinux type, such as httpd_sys_content_t, and *file-name* is a file or directory name:

```
~]$ chcon -t httpd_sys_content_t file-name
```

 Run the chcon -R -t type directory-name command to change the type of the directory and its contents, where type is an SELinux type, such as httpd_sys_content_t, and directory-name is a directory name:

```
~]$ chcon -R -t httpd_sys_content_t directory-name
```

Procedure 4.6. Changing a File's or Directory's Type

The following procedure demonstrates changing the type, and no other attributes of the SELinux context. The example in this section works the same for directories, for example, if **file1** was a directory.

- 1. Change into your home directory.
- 2. Create a new file and view its SELinux context:

```
~]$ touch file1
```

```
~]$ ls -Z file1
-rw-rw-r-- user1 group1 unconfined_u:object_r:user_home_t:s0 file1
```

In this example, the SELinux context for **file1** includes the SELinux unconfined_u user, object_r role, user_home_t type, and the s0 level. For a description of each part of the SELinux context, see *Chapter 2*, *SELinux Contexts*.

3. Run the following command to change the type to samba_share_t. The **-t** option only changes the type. Then view the change:

```
~]$ chcon -t samba_share_t file1

~]$ ls -Z file1
```

4. Use the following command to restore the SELinux context for the **file1** file. Use the **-v** option to view what changes:

-rw-rw-r-- user1 group1 unconfined_u:object_r:samba_share_t:s0 file1

```
~]$ restorecon -v file1
restorecon reset file1 context unconfined_u:object_r:samba_share_t:s0-
>system_u:object_r:user_home_t:s0
```

In this example, the previous type, samba_share_t, is restored to the correct, user_home_t type. When using targeted policy (the default SELinux policy in Fedora), the **restorecon** command reads the files in the **/etc/selinux/targeted/contexts/files/** directory, to see which SELinux context files should have.

Procedure 4.7. Changing a Directory and its Contents Types

The following example demonstrates creating a new directory, and changing the directory's file type along with its contents to a type used by the Apache HTTP Server. The configuration in this example is used if you want Apache HTTP Server to use a different document root (instead of /var/www/html/):

As the root user, create a new /web/ directory and then 3 empty files (file1, file2, and file3) within this directory. The /web/ directory and files in it are labeled with the default_t type:

```
~]# mkdir /web

~]# touch /web/file{1,2,3}

~]# ls -dz /web

drwxr-xr-x root root unconfined_u:object_r:default_t:s0 /web

~]# ls -lz /web
```

-rw-r--r- root root unconfined_u:object_r:default_t:s0 file1

```
-rw-r--r-- root root unconfined_u:object_r:default_t:s0 file2
-rw-r--r-- root root unconfined_u:object_r:default_t:s0 file3
```

2. As root, run the following command to change the type of the /web/ directory (and its contents) to httpd_sys_content_t:

```
~]# chcon -R -t httpd_sys_content_t /web/
```

```
~]# ls -dZ /web/
drwxr-xr-x root root unconfined_u:object_r:httpd_sys_content_t:s0 /web/
```

```
~]# ls -lz /web/
-rw-r--r- root root unconfined_u:object_r:httpd_sys_content_t:s0 file1
-rw-r--r- root root unconfined_u:object_r:httpd_sys_content_t:s0 file2
-rw-r--r- root root unconfined_u:object_r:httpd_sys_content_t:s0 file3
```

3. To restore the default SELinux contexts, use the restorecon utility as root:

```
~]# restorecon -R -v /web/
restorecon reset /web context unconfined_u:object_r:httpd_sys_content_t:s0-
>system_u:object_r:default_t:s0
restorecon reset /web/file2 context unconfined_u:object_r:httpd_sys_content_t:s0-
>system_u:object_r:default_t:s0
restorecon reset /web/file3 context unconfined_u:object_r:httpd_sys_content_t:s0-
>system_u:object_r:default_t:s0
restorecon reset /web/file1 context unconfined_u:object_r:httpd_sys_content_t:s0-
>system_u:object_r:default_t:s0
```

See the chcon(1) manual page for further information about **chcon**.



Note

Type Enforcement is the main permission control used in SELinux targeted policy. For the most part, SELinux users and roles can be ignored.

4.7.2. Persistent Changes: semanage fcontext

The **semanage fcontext** command is used to change the SELinux context of files. When using targeted policy, changes are written to files located in the **/etc/selinux/targeted/contexts/files/** directory:

- The **file_contexts** file specifies default contexts for many files, as well as contexts updated via **semanage fcontext**.
- The **file_contexts.local** file stores contexts to newly created files and directories not found in **file_contexts**.

Two utilities read these files. The setfiles utility is used when a file system is relabeled and the restorecon utility restores the default SELinux contexts. This means that changes made by **semanage fcontext** are persistent, even if the file system is relabeled. SELinux policy controls whether users are able to modify the SELinux context for any given file.

Ouick Reference

To make SELinux context changes that survive a file system relabel:

1. Run the following command, remembering to use the full path to the file or directory:

```
~]# semanage fcontext -a options file-name|directory-name
```

2. Use the restorecon utility to apply the context changes:

```
~]# restorecon -v file-name|directory-name
```

Procedure 4.8. Changing a File's or Directory 's Type

The following example demonstrates changing a file's type, and no other attributes of the SELinux context. This example works the same for directories, for instance if **file1** was a directory.

 As the root user, create a new file in the /etc/ directory. By default, newly-created files in /etc/ are labeled with the etc_t type:

```
~]# touch /etc/file1

~]$ ls -z /etc/file1
-rw-r--r-- root root unconfined_u:object_r:etc_t:s0 /etc/file1
```

To list information about a directory, use the following command:

```
~]$ ls -dZ directory_name
```

2. As root, run the following command to change the **file1** type to samba_share_t. The **-a** option adds a new record, and the **-t** option defines a type (samba_share_t). Note that running this command does not directly change the type; **file1** is still labeled with the etc_t type:

```
~]# semanage fcontext -a -t samba_share_t /etc/file1

~]# ls -z /etc/file1
-rw-r--r-- root root unconfined_u:object_r:etc_t:s0 /etc/file1
```

The semanage fcontext -a -t samba_share_t /etc/file1 command adds the following entry to /etc/selinux/targeted/contexts/files/file_contexts.local:

```
/etc/file1 unconfined_u:object_r:samba_share_t:s0
```

3. As root, use the restorecon utility to change the type. Because semanage added an entry to **file.contexts.local** for **/etc/file1**, restorecon changes the type to samba_share_t:

```
~]# restorecon -v /etc/file1
restorecon reset /etc/file1 context unconfined_u:object_r:etc_t:s0-
>system_u:object_r:samba_share_t:s0
```

Procedure 4.9. Changing a Directory and its Contents Types

The following example demonstrates creating a new directory, and changing the directory's file type along with its contents to a type used by Apache HTTP Server. The configuration in this example is used if you want Apache HTTP Server to use a different document root instead of /var/www/html/:

As the root user, create a new /web/ directory and then 3 empty files (file1, file2, and file3) within this directory. The /web/ directory and files in it are labeled with the default_t type:

```
~]# mkdir /web

~]# touch /web/file{1,2,3}

~]# ls -dZ /web
drwxr-xr-x root root unconfined_u:object_r:default_t:s0 /web

~]# ls -lZ /web
-rw-r--r- root root unconfined_u:object_r:default_t:s0 file1
-rw-r--r- root root unconfined_u:object_r:default_t:s0 file2
-rw-r--r- root root unconfined_u:object_r:default_t:s0 file3
```

2. As root, run the following command to change the type of the /web/ directory and the files in it, to httpd_sys_content_t. The -a option adds a new record, and the -t option defines a type (httpd_sys_content_t). The "/web(/.*)?" regular expression causes semanage to apply changes to /web/, as well as the files in it. Note that running this command does not directly change the type; /web/ and files in it are still labeled with the default_t type:

```
~]# semanage fcontext -a -t httpd_sys_content_t "/web(/.*)?"

~]$ ls -dZ /web
drwxr-xr-x root root unconfined_u:object_r:default_t:s0 /web

-]$ ls -lZ /web
-rw-r--r-- root root unconfined_u:object_r:default_t:s0 file1
-rw-r--r-- root root unconfined_u:object_r:default_t:s0 file2
-rw-r--r-- root root unconfined_u:object_r:default_t:s0 file3
```

The semanage fcontext -a -t httpd_sys_content_t "/web(/.*)?" command adds the following entry to /etc/selinux/targeted/contexts/files/file_contexts.local:

```
/web(/.*)? system_u:object_r:httpd_sys_content_t:s0
```

3. As root, use the restorecon utility to change the type of /web/, as well as all files in it. The -R is for recursive, which means all files and directories under /web/ are labeled with the httpd_sys_content_t type. Since semanage added an entry to file.contexts.local for /web(/.*)?, restorecon changes the types to httpd_sys_content_t:

```
~]# restorecon -R -v /web
restorecon reset /web context unconfined_u:object_r:default_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
restorecon reset /web/file2 context unconfined_u:object_r:default_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
restorecon reset /web/file3 context unconfined_u:object_r:default_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
restorecon reset /web/file1 context unconfined_u:object_r:default_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
```

Note that by default, newly-created files and directories inherit the SELinux type of their parent directories.

Procedure 4.10. Deleting an added Context

The following example demonstrates adding and removing an SELinux context. If the context is part of a regular expression, for example, /web(/.*)?, use quotation marks around the regular expression:

```
~]# semanage fcontext -d "/web(/.*)?"
```

1. To remove the context, as root, run the following command, where *file-name*|*directory-name* is the first part in **file_contexts.local**:

```
~]# semanage fcontext -d file-name|directory-name
```

The following is an example of a context in **file_contexts.local**:

```
/test system_u:object_r:httpd_sys_content_t:s0
```

With the first part being /test. To prevent the /test/ directory from being labeled with the httpd_sys_content_t after running restorecon, or after a file system relabel, run the following command as root to delete the context from file_contexts.local:

```
~]# semanage fcontext -d /test
```

2. As root, use the restorecon utility to restore the default SELinux context.

See the semanage(8) manual page for further information about semanage.



Important

When changing the SELinux context with **semanage fcontext -a**, use the full path to the file or directory to avoid files being mislabeled after a file system relabel, or after the **restorecon** command is run.

4.8. The file_t and default_t Types

When using a file system that supports extended attributes (EA), the file_t type is the default type of a file that has not yet been assigned EA value. This type is only used for this purpose and does not

exist on correctly-labeled file systems, because all files on a system running SELinux should have a proper SELinux context, and the file_t type is never used in file-context configuration⁷.

The default_t type is used on files that do not match any pattern in file-context configuration, so that such files can be distinguished from files that do not have a context on disk, and generally are kept inaccessible to confined domains. For example, if you create a new top-level directory, such as /mydirectory/, this directory may be labeled with the default_t type. If services need access to this directory, you need to update the file-contexts configuration for this location. See Section 4.7.2, "Persistent Changes: semanage fcontext" for details on adding a context to the file-context configuration.

4.9. Mounting File Systems

By default, when a file system that supports extended attributes is mounted, the security context for each file is obtained from the *security.selinux* extended attribute of the file. Files in file systems that do not support extended attributes are assigned a single, default security context from the policy configuration, based on file system type.

Use the **mount -o context** command to override existing extended attributes, or to specify a different, default context for file systems that do not support extended attributes. This is useful if you do not trust a file system to supply the correct attributes, for example, removable media used in multiple systems. The **mount -o context** command can also be used to support labeling for file systems that do not support extended attributes, such as File Allocation Table (FAT) or NFS volumes. The context specified with the **context** option is not written to disk: the original contexts are preserved, and are seen when mounting without **context** if the file system had extended attributes in the first place.

For further information about file system labeling, see James Morris's "Filesystem Labeling in SELinux" article: http://www.linuxjournal.com/article/7426.

4.9.1. Context Mounts

To mount a file system with the specified context, overriding existing contexts if they exist, or to specify a different, default context for a file system that does not support extended attributes, as the root user, use the **mount -o context=SELinux_user:role:type:level** command when mounting the desired file system. Context changes are not written to disk. By default, NFS mounts on the client side are labeled with a default context defined by policy for NFS volumes. In common policies, this default context uses the nfs_t type. Without additional mount options, this may prevent sharing NFS volumes using other services, such as the Apache HTTP Server. The following example mounts an NFS volume so that it can be shared via the Apache HTTP Server:

```
~]# mount server:/export /local/mount/point -o \
context="system_u:object_r:httpd_sys_content_t:s0"
```

Newly-created files and directories on this file system appear to have the SELinux context specified with **-o context**. However, since these changes are not written to disk, the context specified with this option does not persist between mounts. Therefore, this option must be used with the same context specified during every mount to retain the desired context. For information about making context mount persistent, see *Section 4.9.5*, *"Making Context Mounts Persistent"*.

⁷ Files in the /etc/selinux/targeted/contexts/files/ directory define contexts for files and directories. Files in this directory are read by the restorecon and setfiles utilities to restore files and directories to their default contexts.

Type Enforcement is the main permission control used in SELinux targeted policy. For the most part, SELinux users and roles can be ignored, so, when overriding the SELinux context with **-o context**, use the SELinux system_u user and object_r role, and concentrate on the type. If you are not using the MLS policy or multi-category security, use the s0 level.



Note

When a file system is mounted with a **context** option, context changes by users and processes are prohibited. For example, running the **chcon** command on a file system mounted with a **context** option results in a **Operation not supported** error.

4.9.2. Changing the Default Context

As mentioned in Section 4.8, "The file_t and default_t Types", on file systems that support extended attributes, when a file that lacks an SELinux context on disk is accessed, it is treated as if it had a default context as defined by SELinux policy. In common policies, this default context uses the file_t type. If it is desirable to use a different default context, mount the file system with the defcontext option.

The following example mounts a newly-created file system on /dev/sda2 to the newly-created / test/ directory. It assumes that there are no rules in /etc/selinux/targeted/contexts/ files/ that define a context for the /test/ directory:

~]# mount /dev/sda2 /test/ -o defcontext="system_u:object_r:samba_share_t:s0"

In this example:

- the **defcontext** option defines that **system_u:object_r:samba_share_t:s0** is "the default security context for unlabeled files"⁸.
- when mounted, the root directory (/test/) of the file system is treated as if it is labeled with the
 context specified by defcontext (this label is not stored on disk). This affects the labeling for files
 created under /test/: new files inherit the samba_share_t type, and these labels are stored on
 disk.
- files created under /test/ while the file system was mounted with a defcontext option retain their labels.

4.9.3. Mounting an NFS Volume

By default, NFS mounts on the client side are labeled with a default context defined by policy for NFS volumes. In common policies, this default context uses the nfs_t type. Depending on policy configuration, services, such as Apache HTTP Server and MariaDB, may not be able to read files labeled with the nfs_t type. This may prevent file systems labeled with this type from being mounted and then read or exported by other services.

If you would like to mount an NFS volume and read or export that file system with another service, use the **context** option when mounting to override the nfs_t type. Use the following context option to mount NFS volumes so that they can be shared via the Apache HTTP Server:

⁸ Morris, James. "Filesystem Labeling in SELinux". Published 1 October 2004. Accessed 14 October 2008: http://www.linuxjournal.com/article/7426.

```
~]# mount server:/export /local/mount/point -o
context="system_u:object_r:httpd_sys_content_t:s0"
```

Since these changes are not written to disk, the context specified with this option does not persist between mounts. Therefore, this option must be used with the same context specified during every mount to retain the desired context. For information about making context mount persistent, see *Section 4.9.5, "Making Context Mounts Persistent"*.

As an alternative to mounting file systems with **context** options, Booleans can be enabled to allow services access to file systems labeled with the nfs_t type. See *Part II*, "Managing Confined Services" for instructions on configuring Booleans to allow services access to the nfs_t type.

4.9.4. Multiple NFS Mounts

When mounting multiple mounts from the same NFS export, attempting to override the SELinux context of each mount with a different context, results in subsequent mount commands failing. In the following example, the NFS server has a single export, /export/, which has two subdirectories, /web/ and /database/. The following commands attempt two mounts from a single NFS export, and try to override the context for each one:

```
~]# mount server:/export/web /local/web -o context="system_u:object_r:httpd_sys_content_t:s0"
```

```
~]# mount server:/export/database /local/database -o
context="system_u:object_r:mysqld_db_t:s0"
```

The second mount command fails, and the following is logged to /var/log/messages:

```
kernel: SELinux: mount invalid. Same superblock, different security settings for (dev 0:15, type nfs)
```

To mount multiple mounts from a single NFS export, with each mount having a different context, use the **-o nosharecache, context** options. The following example mounts multiple mounts from a single NFS export, with a different context for each mount (allowing a single service access to each one):

```
~]# mount server:/export/web /local/web -o
nosharecache,context="system_u:object_r:httpd_sys_content_t:s0"
```

```
~]# mount server:/export/database /local/database -o \
nosharecache,context="system_u:object_r:mysqld_db_t:s0"
```

In this example, **server:/export/web** is mounted locally to the **/local/web/** directory, with all files being labeled with the httpd_sys_content_t type, allowing Apache HTTP Server access. **server:/export/database** is mounted locally to **/local/database**/, with all files being labeled with the mysqld_db_t type, allowing MariaDB access. These type changes are not written to disk.



Important

The **nosharecache** options allows you to mount the same subdirectory of an export multiple times with different contexts, for example, mounting **/export/web/** multiple times. Do not mount the same subdirectory from an export multiple times with different contexts, as this creates an overlapping mount, where files are accessible under two different contexts.

4.9.5. Making Context Mounts Persistent

To make context mounts persistent across remounting and reboots, add entries for the file systems in the /etc/fstab file or an automounter map, and use the desired context as a mount option. The following example adds an entry to /etc/fstab for an NFS context mount:

server:/export /local/mount/ nfs context="system_u:object_r:httpd_sys_content_t:s0" 0 0

4.10. Maintaining SELinux Labels

These sections describe what happens to SELinux contexts when copying, moving, and archiving files and directories. Also, it explains how to preserve contexts when copying and archiving.

4.10.1. Copying Files and Directories

When a file or directory is copied, a new file or directory is created if it does not exist. That new file or directory's context is based on default-labeling rules, not the original file or directory's context unless options were used to preserve the original context. For example, files created in user home directories are labeled with the user_home_t type:

```
~]$ touch file1
```

```
~]$ ls -Z file1
-rw-rw-r-- user1 group1 unconfined_u:object_r:user_home_t:s0 file1
```

If such a file is copied to another directory, such as /etc/, the new file is created in accordance to default-labeling rules for /etc/. Copying a file without additional options may not preserve the original context:

```
~]$ ls -Z file1
-rw-rw-r-- user1 group1 unconfined_u:object_r:user_home_t:s0 file1
```

```
~]# cp file1 /etc/
```

```
~]$ ls -Z /etc/file1
-rw-r--r-- root root unconfined_u:object_r:etc_t:s0 /etc/file1
```

When **file1** is copied to **/etc/**, if **/etc/file1** does not exist, **/etc/file1** is created as a new file. As shown in the example above, **/etc/file1** is labeled with the etc_t type, in accordance to default-labeling rules.

When a file is copied over an existing file, the existing file's context is preserved, unless the user specified **cp** options to preserve the context of the original file, such as **--preserve=context**. SELinux policy may prevent contexts from being preserved during copies.

Procedure 4.11. Copying Without Preserving SELinux Contexts

This procedure shows that when copying a file with the **cp** command, if no options are given, the type is inherited from the targeted, parent directory.

1. Create a file in a user's home directory. The file is labeled with the user_home_t type:

```
~]$ touch file1

~]$ ls -Z file1
-rw-rw-r-- user1 group1 unconfined_u:object_r:user_home_t:s0 file1
```

2. The /var/www/html/ directory is labeled with the httpd_sys_content_t type, as shown with the following command:

```
~]$ ls -dZ /var/www/html/
drwxr-xr-x root root system_u:object_r:httpd_sys_content_t:s0 /var/www/html/
```

3. When **file1** is copied to **/var/www/html/**, it inherits the httpd_sys_content_t type:

```
~]# cp file1 /var/www/html/

~]$ ls -Z /var/www/html/file1
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 /var/www/html/file1
```

Procedure 4.12. Preserving SELinux Contexts When Copying

This procedure shows how to use the **--preserve=context** option to preserve contexts when copying.

1. Create a file in a user's home directory. The file is labeled with the user_home_t type:

```
~]$ touch file1

~]$ ls -Z file1
-rw-rw-r-- user1 group1 unconfined_u:object_r:user_home_t:s0 file1
```

2. The /var/www/html/ directory is labeled with the httpd_sys_content_t type, as shown with the following command:

```
~]$ ls -dZ /var/www/html/
drwxr-xr-x root root system_u:object_r:httpd_sys_content_t:s0 /var/www/html/
```

3. Using the --preserve=context option preserves SELinux contexts during copy operations. As shown below, the user_home_t type of file1 was preserved when the file was copied to / var/www/html/:

```
~]# cp --preserve=context file1 /var/www/html/
```

```
~]$ ls -Z /var/www/html/file1
-rw-r--r-- root root unconfined_u:object_r:user_home_t:s0 /var/www/html/file1
```

Procedure 4.13. Copying and Changing the Context

This procedure show how to use the **--context** option to change the destination copy's context. The following example is performed in the user's home directory:

1. Create a file in a user's home directory. The file is labeled with the user_home_t type:

```
~]$ touch file1
```

```
~]$ ls -Z file1
-rw-rw-r-- user1 group1 unconfined_u:object_r:user_home_t:s0 file1
```

2. Use the **--context** option to define the SELinux context:

```
~]$ cp --context=system_u:object_r:samba_share_t:s0 file1 file2
```

Without --context, file2 would be labeled with the unconfined_u:object_r:user_home_t context:

```
~]$ ls -Z file1 file2
-rw-rw-r-- user1 group1 unconfined_u:object_r:user_home_t:s0 file1
-rw-rw-r-- user1 group1 system_u:object_r:samba_share_t:s0 file2
```

Procedure 4.14. Copying a File Over an Existing File

This procedure shows that when a file is copied over an existing file, the existing file's context is preserved unless an option is used to preserve contexts.

 As root, create a new file, file1 in the /etc/ directory. As shown below, the file is labeled with the etc_t type:

```
~]# touch /etc/file1
```

```
~]$ ls -Z /etc/file1
-rw-r--r-- root root unconfined_u:object_r:etc_t:s0 /etc/file1
```

2. Create another file, **file2**, in the **/tmp/** directory. As shown below, the file is labeled with the user_tmp_t type:

```
~]$ touch /tmp/file2
```

```
~$ ls -Z /tmp/file2
-rw-r--r-- root root unconfined_u:object_r:user_tmp_t:s0 /tmp/file2
```

3. Overwrite file1 with file2:

```
~]# cp /tmp/file2 /etc/file1
```

4. After copying, the following command shows **file1** labeled with the etc_t type, not the user_tmp_t type from **/tmp/file2** that replaced **/etc/file1**:

```
~]$ ls -Z /etc/file1
-rw-r--r-- root root unconfined_u:object_r:etc_t:s0 /etc/file1
```



Important

Copy files and directories, rather than moving them. This helps ensure they are labeled with the correct SELinux contexts. Incorrect SELinux contexts can prevent processes from accessing such files and directories.

4.10.2. Moving Files and Directories

Files and directories keep their current SELinux context when they are moved. In many cases, this is incorrect for the location they are being moved to. The following example demonstrates moving a file from a user's home directory to the /var/www/html/ directory, which is used by the Apache HTTP Server. Since the file is moved, it does not inherit the correct SELinux context:

Procedure 4.15. Moving Files and Directories

 Change into your home directory and create file in it. The file is labeled with the user_home_t type:

```
~]$ touch file1

~]$ ls -Z file1
-rw-rw-r-- user1 group1 unconfined_u:object_r:user_home_t:s0 file1
```

Run the following command to view the SELinux context of the /var/www/html/ directory:

```
~]$ ls -dZ /var/www/html/
drwxr-xr-x root root system_u:object_r:httpd_sys_content_t:s0 /var/www/html/
```

By default, /var/www/html/ is labeled with the httpd_sys_content_t type. Files and directories created under /var/www/html/ inherit this type, and as such, they are labeled with this type.

3. As root, move **file1** to **/var/www/html/**. Since this file is moved, it keeps its current user_home_t type:

```
~]# mv file1 /var/www/html/
```

```
~]# ls -Z /var/www/html/file1
-rw-rw-r-- user1 group1 unconfined_u:object_r:user_home_t:s0 /var/www/html/file1
```

By default, the Apache HTTP Server cannot read files that are labeled with the user_home_t type. If all files comprising a web page are labeled with the user_home_t type, or another type that the Apache HTTP Server cannot read, permission is denied when attempting to access them via web browsers, such as **Mozilla Firefox**.



Important

Moving files and directories with the **mv** command may result in the incorrect SELinux context, preventing processes, such as the Apache HTTP Server and Samba, from accessing such files and directories.

4.10.3. Checking the Default SELinux Context

Use the matchpathcon utility to check if files and directories have the correct SELinux context. This utility queries the system policy and then provides the default security context associated with the file path. The following example demonstrates using matchpathcon to verify that files in /var/www/html/ directory are labeled correctly:

Procedure 4.16. Checking the Default SELinux Conxtext with matchpathcon

 As the root user, create three files (file1, file2, and file3) in the /var/www/html/ directory. These files inherit the httpd_sys_content_t type from /var/www/html/:

```
~]# touch /var/www/html/file{1,2,3}
```

```
~]# ls -Z /var/www/html/
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 file1
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 file2
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 file3
```

2. As root, change the **file1** type to samba_share_t. Note that the Apache HTTP Server cannot read files or directories labeled with the samba_share_t type.

```
~]# chcon -t samba_share_t /var/www/html/file1
```

 The matchpathcon -V option compares the current SELinux context to the correct, default context in SELinux policy. Run the following command to check all files in the /var/www/html/ directory:

⁹ See the matchpathcon(8) manual page for further information about matchpathcon.

```
-]$ matchpathcon -V /var/www/html/*
/var/www/html/file1 has context unconfined_u:object_r:samba_share_t:s0, should be
system_u:object_r:httpd_sys_content_t:s0
/var/www/html/file2 verified.
/var/www/html/file3 verified.
```

The following output from the **matchpathcon** command explains that **file1** is labeled with the samba_share_t type, but should be labeled with the httpd_sys_content_t type:

```
/var/www/html/file1 has context unconfined_u:object_r:samba_share_t:s0, should be system_u:object_r:httpd_sys_content_t:s0
```

To resolve the label problem and allow the Apache HTTP Server access to **file1**, as root, use the restorecon utility:

```
~]# restorecon -v /var/www/html/file1
restorecon reset /var/www/html/file1 context unconfined_u:object_r:samba_share_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
```

4.10.4. Archiving Files with tar

The tar utility does not retain extended attributes by default. Since SELinux contexts are stored in extended attributes, contexts can be lost when archiving files. Use the **tar --selinux** command to create archives that retain contexts and to restore files from the archives. If a tar archive contains files without extended attributes, or if you want the extended attributes to match the system defaults, use the restorecon utility:

```
~]$ tar -xvf archive.tar | restorecon -f -
```

Note that depending on the directory, you may need to be the root user to run the restorecon.

The following example demonstrates creating a tar archive that retains SELinux contexts:

Procedure 4.17. Creating a tar Archive

1. Change to the /var/www/html/ directory and view its SELinux context:

```
-]$ cd /var/www/html/
html]$ ls -dZ /var/www/html/
drwxr-xr-x. root root system_u:object_r:httpd_sys_content_t:s0 .
```

2. As root, create three files (**file1**, **file2**, and **file3**) in /var/www/html/. These files inherit the httpd_sys_content_t type from /var/www/html/:

```
html]# touch file{1,2,3}
```

```
html]$ ls -Z /var/www/html/
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 file1
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 file2
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 file3
```

 As root, run the following command to create a tar archive named test.tar. Use the -selinux to retain the SELinux context:

```
html]# tar --selinux -cf test.tar file{1,2,3}
```

4. As root, create a new directory named /test/, and then allow all users full access to it:

```
~]# mkdir /test
```

```
~]# chmod 777 /test/
```

5. Copy the test.tar file into /test/:

```
~]$ cp /var/www/html/test.tar /test/
```

6. Change into /test/ directory. Once in this directory, run the following command to extract the tar archive. Specify the --selinux option again otherwise the SELinux context will be changed to default_t:

```
-]$ cd /test/
```

```
test]$ tar --selinux -xvf test.tar
```

7. View the SELinux contexts. The httpd_sys_content_t type has been retained, rather than being changed to default_t, which would have happened had the --selinux not been used:

```
test]$ ls -lZ /test/
-rw-r--r-- user1 group1 unconfined_u:object_r:httpd_sys_content_t:s0 file1
-rw-r--r-- user1 group1 unconfined_u:object_r:httpd_sys_content_t:s0 file2
-rw-r--r-- user1 group1 unconfined_u:object_r:httpd_sys_content_t:s0 file3
-rw-r--r-- user1 group1 unconfined_u:object_r:default_t:s0 test.tar
```

8. If the /test/ directory is no longer required, as root, run the following command to remove it, as well as all files in it:

```
~]# rm -ri /test/
```

See the tar(1) manual page for further information about tar, such as the **--xattrs** option that retains all extended attributes.

4.10.5. Archiving Files with star

The star utility does not retain extended attributes by default. Since SELinux contexts are stored in extended attributes, contexts can be lost when archiving files. Use the **star -xattr -H=exustar** command to create archives that retain contexts. The *star* package is not installed by default. To install **star**, run the **dnf install star** command as the root user.

The following example demonstrates creating a star archive that retains SELinux contexts:

Procedure 4.18. Creating a star Archive

As root, create three files (file1, file2, and file3) in the /var/www/html/. These files inherit the httpd_sys_content_t type from /var/www/html/:

```
~]# touch /var/www/html/file{1,2,3}
```

```
~]# ls -Z /var/www/html/
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 file1
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 file2
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 file3
```

2. Change into /var/www/html/ directory. Once in this directory, as root, run the following command to create a star archive named test.star:

```
~]$ cd /var/www/html
```

```
html]# star -xattr -H=exustar -c -f=test.star file{1,2,3} star: 1 blocks + 0 bytes (total of 10240 bytes = 10.00k).
```

3. As root, create a new directory named /test/, and then allow all users full access to it:

```
~]# mkdir /test
```

```
~]# chmod 777 /test/
```

4. Run the following command to copy the **test.star** file into **/test/**:

```
~]$ cp /var/www/html/test.star /test/
```

5. Change into /test/. Once in this directory, run the following command to extract the star archive:

```
~]$ cd /test/
```

```
test]$ star -x -f=test.star
star: 1 blocks + 0 bytes (total of 10240 bytes = 10.00k).
```

6. View the SELinux contexts. The httpd_sys_content_t type has been retained, rather than being changed to default_t, which would have happened had the -xattr -H=exustar option not been used:

```
~]$ ls -lZ /test/
-rw-r--r-- user1 group1 unconfined_u:object_r:httpd_sys_content_t:s0 file1
-rw-r--r-- user1 group1 unconfined_u:object_r:httpd_sys_content_t:s0 file2
-rw-r--r-- user1 group1 unconfined_u:object_r:httpd_sys_content_t:s0 file3
-rw-r--r-- user1 group1 unconfined_u:object_r:default_t:s0 test.star
```

7. If the /test/ directory is no longer required, as root, run the following command to remove it, as well as all files in it:

```
~]# rm -ri /test/
```

8. If star is no longer required, as root, remove the package:

```
~]# dnf remove star
```

See the star(1) manual page for further information about star.

4.11. Information Gathering Tools

The utilities listed below are command-line tools that provide well-formatted information, such as access vector cache statistics or the number of classes, types, or Booleans.

avcstat

This command provides a short output of the access vector cache statistics since boot. You can watch the statistics in real time by specifying a time interval in seconds. This provides updated statistics since the initial output. The statistics file used is /selinux/avc/cache_stats, and you can specify a different cache file with the -f /path/to/file option.

```
~]# avcstat
lookups hits misses allocs reclaims frees
47517410 47504630 12780 12780 12176 12275
```

seinfo

This utility is useful in describing the break-down of a policy, such as the number of classes, types, Booleans, allow rules, and others. seinfo is a command-line utility that uses a **policy.conf** file, a binary policy file, a modular list of policy packages, or a policy list file as input. You must have the setools-console package installed to use the seinfo utility.

The output of seinfo will vary between binary and source files. For example, the policy source file uses the { } brackets to group multiple rule elements onto a single line. A similar effect happens with attributes, where a single attribute expands into one or many types. Because these are expanded and no longer relevant in the binary policy file, they have a return value of zero in the search results. However, the number of rules greatly increases as each formerly one line rule using brackets is now a number of individual lines.

Some items are not present in the binary policy. For example, neverallow rules are only checked during policy compile, not during runtime, and initial Security Identifiers (SIDs) are not part of the binary policy since they are required prior to the policy being loaded by the kernel during boot.

```
~]# seinfo

Statistics for policy file: /etc/selinux/targeted/policy/policy.24
Policy Version & Type: v.24 (binary, mls)

Classes: 77 Permissions: 229
Sensitivities: 1 Categories: 1024
Types: 3001 Attributes: 244
```

```
Users: 9
                      Roles:
                                       13
Booleans:
               158
                      Cond. Expr.:
                                      193
             262796 Neverallow:
44 Dontaudit:
Allow:
                                       0
Auditallow:
                                    156710
             10760 Type_change:
Type_trans:
                                    38
                      Role allow:
Type_member:
             44
                                       20
Role_trans:
Constraints:
                237
                      Range_trans:
                                      2546
                62
                      Validatetrans:
                                       0
               27
Initial SIDs:
                     Fs_use:
                                       22
Genfscon:
               82 Portcon:
                                       373
Netifcon:
                 0
                      Nodecon:
                                        0
Permissives:
                 22
                                        2
                      Polcap:
```

The seinfo utility can also list the number of types with the domain attribute, giving an estimate of the number of different confined processes:

```
~]# seinfo -adomain -x | wc -l
550
```

Not all domain types are confined. To look at the number of unconfined domains, use the unconfined_domain attribute:

```
~]# seinfo -aunconfined_domain_type -x | wc -l
52
```

Permissive domains can be counted with the **--permissive** option:

```
~]# seinfo --permissive -x | wc -l
31
```

Remove the additional | wc -1 command in the above commands to see the full lists.

sesearch

You can use the sesearch utility to search for a particular rule in the policy. It is possible to search either policy source files or the binary file. For example:

```
~]$ sesearch --role_allow -t httpd_sys_content_t /etc/selinux/targeted/policy/policy.24
Found 20 role allow rules:
   allow system_r sysadm_r;
   allow sysadm_r system_r;
  allow sysadm_r staff_r;
   allow sysadm_r user_r;
   allow system_r git_shell_r;
   allow system_r guest_r;
   allow logadm_r system_r;
  allow system_r logadm_r;
  allow system_r nx_server_r;
  allow system_r staff_r;
   allow staff_r logadm_r;
   allow staff_r sysadm_r;
   allow staff_r unconfined_r;
   allow staff_r webadm_r;
   allow unconfined_r system_r;
   allow system_r unconfined_r;
   allow system_r user_r;
   allow webadm_r system_r;
```

```
allow system_r webadm_r;
allow system_r xguest_r;
```

The sesearch utility can provide the number of *allow* rules:

```
~]# sesearch --allow | wc -l
262798
```

And the number of dontaudit rules:

```
~]# sesearch --dontaudit | wc -l
156712
```

4.12. Multi-Level Security (MLS)

The Multi-Level Security technology refers to a security scheme that enforces the Bell-La Padula Mandatory Access Model. Under MLS, users and processes are called *subjects*, and files, devices, and other passive components of the system are called *objects*. Both subjects and objects are labeled with a security level, which entails a subject's clearance or an object's classification. Each security level is composed of a *sensitivity* and a *category*, for example, an internal release schedule is filed under the internal documents category with a confidential sensitivity.

Figure 4.1, "Levels of clearance" shows levels of clearance as originally designed by the US defense community. Relating to our internal schedule example above, only users that have gained the confidential clearance are allowed to view documents in the confidential category. However, users who only have the confidential clearance are not allowed to view documents that require higher levels or clearance; they are allowed read access only to documents with lower levels of clearance, and write access to documents with higher levels of clearance.



Figure 4.1. Levels of clearance

Figure 4.2, "Allowed data flows using MLS" shows all allowed data flows between a subject running under the "Secret" security level and various objects with different security levels. In simple terms, the Bell-LaPadula model enforces two properties: *no read up* and *no write down*.

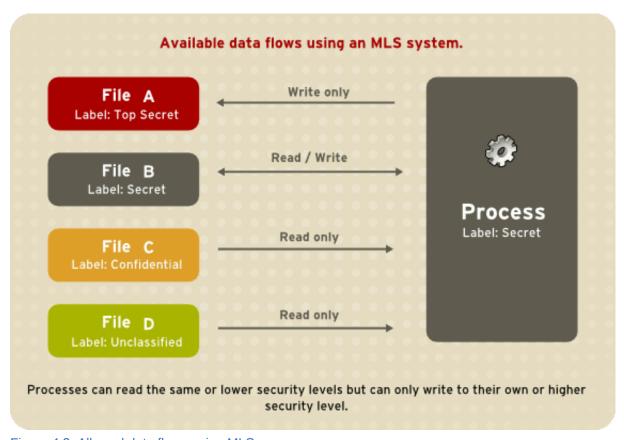


Figure 4.2. Allowed data flows using MLS

4.12.1. MLS and System Privileges

MLS access rules are always combined with conventional access permissions (file permissions). For example, if a user with a security level of "Secret" uses Discretionary Access Control (DAC) to block access to a file by other users, this also blocks access by users with a security level of "Top Secret". It is important to remember that SELinux MLS policy rules are checked *after* DAC rules. A higher security clearance does not automatically give permission to arbitrarily browse a file system.

Users with top-level clearances do not automatically acquire administrative rights on multi-level systems. While they may have access to all information on the computer, this is different from having administrative rights.

4.12.2. Enabling MLS in SELinux



Note

It is not recommended to use the MLS policy on a system that is running the X Window System.

Follow these steps to enable the SELinux MLS policy on your system.

Procedure 4.19. Enabling SELinux MLS Policy

1. Install the selinux-policy-mls package:

```
~]# dnf install selinux-policy-mls
```

2. Before the MLS policy is enabled, each file on the file system must be relabeled with an MLS label. When the file system is relabeled, confined domains may be denied access, which may prevent your system from booting correctly. To prevent this from happening, configure SELINUX=permissive in the /etc/selinux/config file. Also, enable the MLS policy by configuring SELINUXTYPE=mls. Your configuration file should look like this:

```
# This file controls the state of SELinux on the system.
# SELINUX= can take one of these three values:
# enforcing - SELinux security policy is enforced.
# permissive - SELinux prints warnings instead of enforcing.
# disabled - No SELinux policy is loaded.
SELINUX=permissive
# SELINUXTYPE= can take one of these two values:
# targeted - Targeted processes are protected,
# mls - Multi Level Security protection.
SELINUXTYPE=mls
```

3. Make sure SELinux is running in the permissive mode:

```
~]# setenforce 0
```

```
~]$ getenforce
Permissive
```

4. Create the **.autorelabel** file in root's home directory to ensure that files are relabeled upon next reboot:

```
~]# touch /.autorelabel
```

Note that it is necessary to add the ${ extstyle - F}$ option to this file. This can be done by executing the following command:

```
~]# echo "-F" >> /.autorelabel
```

5. Reboot your system. During the next boot, all file systems will be relabeled according to the MLS policy. The label process labels all files with an appropriate SELinux context:

Each * (asterisk) character on the bottom line represents 1000 files that have been labeled. In the above example, eleven * characters represent 11000 files which have been labeled. The time it takes to label all files depends upon the number of files on the system, and the speed of the hard

disk drives. On modern systems, this process can take as little as 10 minutes. Once the labeling process finishes, the system will automatically reboot.

6. In permissive mode, SELinux policy is not enforced, but denials are still logged for actions that would have been denied if running in enforcing mode. Before changing to enforcing mode, as root, run the following command to confirm that SELinux did not deny actions during the last boot. If SELinux did not deny actions during the last boot, this command does not return any output. See *Chapter 10, Troubleshooting* for troubleshooting information if SELinux denied access during boot.

```
~]# grep "SELinux is preventing" /var/log/messages
```

7. If there were no denial messages in the /var/log/messages file, or you have resolved all existing denials, configure SELINUX=enforcing in the /etc/selinux/config file:

```
# This file controls the state of SELinux on the system.
# SELINUX= can take one of these three values:
# enforcing - SELinux security policy is enforced.
# permissive - SELinux prints warnings instead of enforcing.
# disabled - No SELinux policy is loaded.
SELINUX=enforcing
# SELINUXTYPE= can take one of these two values:
# targeted - Targeted processes are protected,
# mls - Multi Level Security protection.
SELINUXTYPE=mls
```

8. Reboot your system and make sure SELinux is running in enforcing mode:

```
~]$ getenforce
Enforcing
```

and the MLS policy is enabled:

```
~]# sestatus |grep mls
Policy from config file: mls
```

4.12.3. Creating a User With a Specific MLS Range

Follow these steps to create a new Linux user with a specific MLS range:

Procedure 4.20. Creating a User With a Specific MLS Range

1. Add a new Linux user using the **useradd** command and map the new Linux user to an existing SELinux user (in this case, user_u):

```
~]# useradd -Z user_u john
```

2. Assign the newly-created Linux user a password:

```
prompt~]# passwd john
```

3. Run the following command as root to view the mapping between SELinux and Linux users. The output should be as follows:

```
~]# semanage login -1
Login Name
                    SELinux User
                                         MLS/MCS Range
                                                              Service
 _default__
                   unconfined_u
                                         s0-s0:c0.c1023
john
                    user_u
                    unconfined_u
                                         s0-s0:c0.c1023
root
system_u
                    system_u
                                         s0-s0:c0.c1023
```

4. Define a specific range for user **john**:

```
~]# semanage login --modify --seuser user_u --range s2:c100 john
```

5. View the mapping between SELinux and Linux users again. Note that the user **john** now has a specific MLS range defined:

```
~]# semanage login -l
Login Name
                     SELinux User
                                          MLS/MCS Range
                                                                Service
 _default__
                     unconfined_u
                                          s0-s0:c0.c1023
john
                     user_u
                                          s2:c100
root
                     unconfined_u
                                          s0-s0:c0.c1023
                                          s0-s0:c0.c1023
system_u
                     system u
```

6. To correct the label on john's home directory if needed, run the following command:

```
~]# chcon -R -l s2:c100 /home/john
```

4.12.4. Setting Up Polyinstantiated Directories

The /tmp/ and /var/tmp/ directories are normally used for temporary storage by all programs, services, and users. Such setup, however, makes these directories vulnerable to race condition attacks, or an information leak based on file names. SELinux offers a solution in the form of polyinstantiated directories. This effectively means that both /tmp/ and /var/tmp/ are instantiated, making them appear private for each user. When instantiation of directories is enabled, each user's /tmp/ and /var/tmp/ directory is automatically mounted under /tmp-inst and /var/tmp/tmp-inst.

Follow these steps to enable polyinstantiation of directories:

Procedure 4.21. Enabling Polyinstantiation Directories

 Uncomment the last three lines in the /etc/security/namespace.conf file to enable instantiation of the /tmp/, /var/tmp/, and users' home directories:

```
~]$ tail -n 3 /etc/security/namespace.conf
/tmp /tmp-inst/ level root,adm
/var/tmp /var/tmp/tmp-inst/ level root,adm
$HOME $HOME/$USER.inst/ level
```

Ensure that in the /etc/pam.d/login file, the pam_namespace. so module is configured for session:

```
~]$ grep namespace /etc/pam.d/login
session required pam_namespace.so
```

3. Reboot your system.

4.13. File Name Transition

The *file name transition* feature allows policy writers to specify the file name when writing policy transition rules. It is possible to write a rule that states: If a process labeled **A_t** creates a specified object class in a directory labeled **B_t** and the specified object class is named **objectname**, it gets the label **C_t**. This mechanism provides more fine-grained control over processes on the system.

Without file name transition, there are three possible ways how to label an object:

- By default, objects inherit labels from parent directories. For example, if the user creates a file in a
 directory labeled etc_t, then the file is labeled also etc_t. However, this method is useless when
 it is desirable to have multiple files within a directory with different labels.
- Policy writers can write a rule in policy that states: If a process with type A_t creates a specified object class in a directory labeled B_t, the object gets the new C_t label. This practice is problematic if a single program creates multiple objects in the same directory where each object requires a separate label. Moreover, these rules provide only partial control, because names of the created objects are not specified.
- Certain applications have SELinux awareness that allow such an application to ask the system what
 the label of a certain path should be. These applications then request the kernel to create the object
 with the required label. Examples of applications with SELinux awareness are the rpm package
 manager, the restorecon utility, or the udev device manager. However, it is not possible to instruct
 every application that creates files or directories with SELinux awareness. It is often necessary to
 relabel objects with the correct label after creating. Otherwise, when a confined domain attempts to
 use the object, AVC messages are returned.

The file name transition feature decreases problems related to incorrect labeling and improves the system to be more secure. Policy writers are able to state properly that a certain application can only create a file with a specified name in a specified directory. The rules take into account the file name, not the file path. This is the basename of the file path. Note that file name transition uses an exact match done by the strcmp() function. Use of regular expressions or wildcard characters is not considered.



Note

File paths can vary in the kernel and file name transition does not use the paths to determine labels. Consequently, this feature only affects initial file creation and does not fix incorrect labels of already created objects.

Example 4.1. Examples of Policy Rules Written with File Name Transition

The example below shows a policy rule with file name transition:

```
filetrans_pattern(unconfined_t, admin_home_t, ssh_home_t, dir, ".ssh")
```

This rule states that if a process with the unconfined_t type creates the \sim /.ssh/ directory in a directory labeled admin_home_t, the \sim /.ssh/ directory gets the label ssh_home_t.

Similar examples of policy rules written with file name transition are presented below:

```
filetrans_pattern(staff_t, user_home_dir_t, httpd_user_content_t, dir, "public_html")
filetrans_pattern(thumb_t, user_home_dir_t, thumb_home_t, file, "missfont.log")
filetrans_pattern(kernel_t, device_t, xserver_misc_device_t, chr_file, "nvidia0")
filetrans_pattern(puppet_t, etc_t, krb5_conf_t, file, "krb5.conf")
```



Note

The file name transition feature affects mainly policy writers, but users can notice that instead of file objects almost always created with the default label of the containing directory, some file objects have a different label as specified in policy.

4.14. Disable ptrace()

The ptrace() system call allows one process to observe and control the execution of another process and change its memory and registers. This call is used primarily by developers during debugging, for example when using the strace utility. When ptrace() is not needed, it can be disabled to improve system security. This can be done by enabling the <code>deny_ptrace</code> Boolean, which denies all processes, even those that are running in <code>unconfined_t</code> domains, from being able to use <code>ptrace()</code> on other processes.

The **deny_ptrace** Boolean is disabled by default. To enable it, run the **setsebool -P deny_ptrace on** command as the root user:

```
~]# setsebool -P deny_ptrace on
```

To verify if this Boolean is enabled, use the following command:

```
~]$ getsebool deny_ptrace deny_ptrace --> on
```

To disable this Boolean, run the **setsebool** -P **deny_ptrace off** command as root:

```
~]# setsebool -P deny_ptrace off
```



Note

The **setsebool** -**P** command makes persistent changes. Do not use the -**P** option if you do not want changes to persist across reboots.

This Boolean influences only packages that are part of Fedora. Consequently, third-party packages could still use the ptrace() system call. To list all domains that are allowed to use ptrace(), run

the following command. Note that the *setools-console* package provides the sesearch utility and that the package is not installed by default.

```
~]# sesearch -A -p ptrace,sys_ptrace -C | grep -v deny_ptrace | cut -d ' ' -f 5
```

4.15. Thumbnail Protection

The thumbnail icons can potentially allow an attacker to break into a locked machine using removable media, such as USB devices or CDs. When the system detects a removable media, the Nautilus file manager executes the thumbnail driver code to display thumbnail icons in an appropriate file browser even if the machine is locked. This behavior is unsafe because if the thumbnail executables were vulnerable, the attacker could use the thumbnail driver code to bypass the lock screen without entering the password.

Therefore, a new SELinux policy is used to prevent such attacks. This policy ensures that all thumbnail drivers are locked when the screen is locked. The thumbnail protection is enabled for both confined users and unconfined users. This policy affects the following applications:

- · /usr/bin/evince-thumbnailer
- /usr/bin/ffmpegthumbnailer
- /usr/bin/gnome-exe-thumbnailer.sh
- · /usr/bin/gnome-nds-thumbnailer
- · /usr/bin/gnome-xcf-thumbnailer
- · /usr/bin/gsf-office-thumbnailer
- · /usr/bin/raw-thumbnailer
- /usr/bin/shotwell-video-thumbnailer
- · /usr/bin/totem-video-thumbnailer
- · /usr/bin/whaaw-thumbnailer
- · /usr/lib/tumbler-1/tumblerd
- /usr/lib64/tumbler-1/tumblerd

The sepolicy Suite

The sepolicy utility provides a suite of features to query the installed SELinux policy. These features are either new or were previously provided by separate utilities, such as sepolgen or setrans. The suite allows you to generate transition reports, man pages, or even new policy modules, thus giving users easier access and better understanding of the SELinux policy.

The *policycoreutils-devel* package provides sepolicy. Run the following command as the root user to install sepolicy:

```
~]# dnf install policycoreutils-devel
```

The sepolicy suite provides the following features that are invoked as command-line parameters:

	Table 5.1.	The se	policy	/ Features
--	------------	--------	--------	------------

Feature	Description
booleans	Query the SELinux Policy to see description of Booleans
communicate	Query the SELinux policy to see if domains can communicate with each other
generate	Generate an SELinux policy module template
gui	Graphical User Interface for SELinux Policy
interface	List SELinux Policy interfaces
manpage	Generate SELinux man pages
network	Query SELinux policy network information
transition	Query SELinux policy and generate a process transition report

5.1. The sepolicy Python Bindings

In previous versions of Fedora, the *setools* package included the sesearch and seinfo utilities. The sesearch utility is used for searching rules in a SELinux policy while the seinfo utility allows you to query various other components in the policy.

In Fedora 24, Python bindings for sesearch and seinfo have been added so that you can use the functionality of these utilities via the sepolicy suite. See the example below:

```
> python
>>> import sepolicy
>>> sepolicy.info(sepolicy.ATTRIBUTE)
Returns a dictionary of all information about SELinux Attributes
>>>sepolicy.search([sepolicy.ALLOW])
Returns a dictionary of all allow rules in the policy.
```

5.2. Generating SELinux Policy Modules: sepolicy generate

In previous versions of Fedora, the sepolgen or selinux-polgengui utilities were used for generating a SELinux policy. These tools have been merged to the sepolicy suite. In Fedora 24, the sepolicy generate command is used to generate an initial SELinux policy module template.

Unlike sepolgen, it is not necessary to run **sepolicy generate** as the root user. This utility also creates an RPM spec file, which can be used to build an RPM package that installs the policy package file (*NAME.pp*) and the interface file (*NAME.if*) to the correct location, provides installation of the SELinux policy into the kernel, and fixes the labeling. The setup script continues to install SELinux policy and sets up the labeling. In addition, a manual page based on the installed policy is generated using the **sepolicy manpage** command. Finally, **sepolicy generate** builds and compiles the SELinux policy and the manual page into an RPM package, ready to be installed on other systems.

When sepolicy generate is executed, the following files are produced:

NAME. te – type enforcing file

This file defines all the types and rules for a particular domain.

NAME. if – interface file

This file defines the default file context for the system. It takes the file types created in the **NAME.te** file and associates file paths to the types. Utilities, such as restorecon and rpm, use these paths to write labels.

NAME_selinux.spec - RPM spec file

This file is an RPM spec file that installs SELinux policy and sets up the labeling. This file also installs the interface file and a man page describing the policy. You can use the **sepolicy manpage -d NAME** command to generate the man page.

NAME . **sh** – helper shell script

This script helps to compile, install, and fix the labeling on the system. It also generates a man page based on the installed policy, compiles, and builds an RPM package suitable to be installed on other systems.

If it is possible to generate an SELinux policy module, sepolicy generate prints out all generated paths from the source domain to the target domain. See the sepolicy-generate(8) manual page for further information about sepolicy generate.

5.3. Prioritizing SELinux Policy Modules

In Fedora 23, the SELinux module storage in **/var/lib/selinux** allows using a priority on SELinux modules. Run the following command as root to show two module directories with a different priority:

```
~]# ls /var/lib/selinux/targeted/active/modules
100 400 disabled
```

The default priority and priority used in the packages of SELinux policy is 100, so you can find most of the SELinux modules organized in the directory /var/lib/selinux/targeted/active/modules/100.

To overwrite an existing module with a modified module, prioritize the new module. The highest priority wins.

Example 5.1. Using SELinux Policy Modules Priority

Prepare a new module with modified file context. Install the module with the **semodule -i** command and set the priority of the module to 400. We use **sandbox.pp** in the following example.

```
~]# semodule -X 400 -i sandbox.pp
```

¹ See Section 5.5, "Generating Manual Pages: **sepolicy manpage**" for more information about sepolicy manpage.

```
~]# semodule --list-modules=full | grep sandbox
400 sandbox pp
100 sandbox pp
```

To return back to the default module, run the **semodule** -r command as root:

```
~]# semodule -X 400 -r sandbox
libsemanage.semanage_direct_remove_key: sandbox module at priority 100 is now active.
```

5.4. Understanding Domain Transitions: sepolicy transition

Previously, the setrans utility was used to examine if transition between two domain or process types is possible and printed out all intermediary types that are used to transition between these domains or processes. In Fedora 24, setrans is provided as part of the sepolicy suite and the sepolicy transition command is now used instead.

The sepolicy transition command queries a SELinux policy and creates a process transition report. The **sepolicy transition** command requires two command-line arguments – a source domain (specified by the **-s** option) and a target domain (specified by the **-t** option). If only the source domain is entered, sepolicy transition lists all possible domains that the source domain can transition to. The following output does not contain all entries. The "@" character means "execute":

```
~]$ sepolicy transition -s httpd_t
httpd_t @ httpd_suexec_exec_t --> httpd_suexec_t
httpd_t @ mailman_cgi_exec_t --> mailman_cgi_t
httpd_t @ abrt_retrace_worker_exec_t --> abrt_retrace_worker_t
httpd_t @ dirsrvadmin_unconfined_script_exec_t --> dirsrvadmin_unconfined_script_t
httpd_t @ httpd_unconfined_script_exec_t --> httpd_unconfined_script_t
```

If the target domain is specified, sepolicy transition examines SELinux policy for all transition paths from the source domain to the target domain and lists these paths. The output below is not complete:

```
~]$ sepolicy transition -s httpd_t -t system_mail_t
httpd_t @ exim_exec_t --> system_mail_t
httpd_t @ courier_exec_t --> system_mail_t
httpd_t @ sendmail_exec_t --> system_mail_t
httpd_t ... httpd_suexec_t @ sendmail_exec_t --> system_mail_t
httpd_t ... httpd_suexec_t @ exim_exec_t --> system_mail_t
httpd_t ... httpd_suexec_t @ courier_exec_t --> system_mail_t
httpd_t ... httpd_suexec_t @ courier_exec_t --> system_mail_t
httpd_t ... httpd_suexec_t ... httpd_mojomojo_script_t @ sendmail_exec_t --> system_mail_t
```

See the sepolicy-transition(8) manual page for further information about sepolicy transition.

5.5. Generating Manual Pages: sepolicy manpage

The sepolicy manpage command generates manual pages based on the SELinux policy that document process domains. As a result, such documentation is always up-to-date. Each name of automatically generated manual pages consists of the process domain name and the _selinux suffix, for example httpd_selinux.

The manual pages include several sections that provide information about various parts of the SELinux policy for confined domains:

Chapter 5. The sepolicy Suite

- The Entrypoints section contains all executable files that need to be executed during a domain transition.
- The **Process Types** section lists all process types that begin with the same prefix as the target domain.
- The **Booleans** section lists Booleans associated with the domain.
- The **Port Types** section contains the port types matching the same prefix as the domain and describes the default port numbers assigned to these port types.
- The **Managed Files** section describes the types that the domain is allowed to write to and the default paths associated with these types.
- The **File Contexts** section contains all file types associated with the domain and describes how to use these file types along with the default path labeling on a system.
- The **Sharing Files** section explains how to use the domain sharing types, such as public_content_t.

See the sepolicy-manpage(8) manual page for further information about sepolicy manpage.

Confining Users

A number of confined SELinux users are available in Fedora. Each Linux user is mapped to an SELinux user using SELinux policy, allowing Linux users to inherit the restrictions placed on SELinux users, for example (depending on the user), not being able to: run the X Window System; use networking; run setuid applications (unless SELinux policy permits it); or run the **su** and **sudo** commands. This helps protect the system from the user. See Section 3.3, "Confined and Unconfined Users" for further information about confined users.

6.1. Linux and SELinux User Mappings

As the root user, run the following command to view the mapping between Linux users and SELinux users:

```
~]# semanage login -l

Login Name SELinux User MLS/MCS Range Service

__default__ unconfined_u s0-s0:c0.c1023 *
root unconfined_u s0-s0:c0.c1023 *
system_u system_u s0-s0:c0.c1023 *
```

In Fedora, Linux users are mapped to the SELinux __default__ login by default (which is in turn mapped to the SELinux unconfined_u user). When a Linux user is created with the useradd command, if no options are specified, they are mapped to the SELinux unconfined_u user. The following defines the default-mapping:

```
__default__ unconfined_u s0-s0:c0.c1023 *
```

6.2. Confining New Linux Users: useradd

Linux users mapped to the SELinux unconfined_u user run in the unconfined_t domain. This is seen by running the id -Z command while logged-in as a Linux user mapped to unconfined_u:

```
~]$ id -Z
unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023
```

When Linux users run in the unconfined_t domain, SELinux policy rules are applied, but policy rules exist that allow Linux users running in the unconfined_t domain almost all access. If unconfined Linux users execute an application that SELinux policy defines can transition from the unconfined_t domain to its own confined domain, unconfined Linux users are still subject to the restrictions of that confined domain. The security benefit of this is that, even though a Linux user is running unconfined, the application remains confined, and therefore, the exploitation of a flaw in the application can be limited by policy.



Note

This does not protect the system from the user. Instead, the user and the system are being protected from possible damage caused by a flaw in the application.

When creating Linux users with the **useradd** command, use the **-Z** option to specify which SELinux user they are mapped to. The following example creates a new Linux user, **useruuser**, and maps that user to the SELinux user_u user. Linux users mapped to the SELinux user_u user run in the user_t domain. In this domain, Linux users are unable to run setuid applications unless SELinux policy permits it (such as passwd), and cannot run the **su** or **sudo** command, preventing them from becoming the root user with these commands.

Procedure 6.1. Confining a New Linux User to user_u SELinux User

1. As root, create a new Linux user (useruuser) that is mapped to the SELinux user_u user.

```
~]# useradd -Z user_u useruuser
```

2. To view the mapping between **useruuser** and user_u, run the following command as root:

```
~]# semanage login -l
Login Name
                     SELinux User
                                          MLS/MCS Range
                                                                Service
                     unconfined_u
                                          s0-s0:c0.c1023
 _default__
                     unconfined u
                                          s0-s0:c0.c1023
root
                                          s0-s0:c0.c1023
system_u
                     system_u
useruuser
                     user_u
```

3. As root, assign a password to the Linux **useruuser** user:

```
~]# passwd useruuser
Changing password for user useruuser.
New password: Enter a password
Retype new password: Enter the same password again
passwd: all authentication tokens updated successfully.
```

4. Log out of your current session, and log in as the Linux **useruuser** user. When you log in, the pam_selinux module maps the Linux user to an SELinux user (in this case, user_u), and sets up the resulting SELinux context. The Linux user's shell is then launched with this context. Run the following command to view the context of a Linux user:

```
~]$ id -Z
user_u:user_r:user_t:s0
```

Log out of the Linux useruuser's session, and log back in with your account. If you do not want
the Linux useruuser user, run the following command as root to remove it, along with its home
directory:

```
~]# userdel -r useruuser
```

6.3. Confining Existing Linux Users: semanage login

If a Linux user is mapped to the SELinux unconfined_u user (the default behavior), and you would like to change which SELinux user they are mapped to, use the **semanage login** command. The following example creates a new Linux user named **newuser**, then maps that Linux user to the SELinux user_u user:

Procedure 6.2. Mapping Linux Users to the SELinux Users

1. As the root user, create a new Linux user (newuser). Since this user uses the default mapping, it does not appear in the semanage login -l output:

```
~]# useradd newuser

~]# semanage login -l

Login Name SELinux User MLS/MCS Range Service

__default__ unconfined_u s0-s0:c0.c1023 *
root unconfined_u s0-s0:c0.c1023 *
system_u system_u s0-s0:c0.c1023 *
```

2. To map the Linux **newuser** user to the SELinux user_u user, run the following command as root:

```
~]# semanage login -a -s user_u newuser
```

The -a option adds a new record, and the -s option specifies the SELinux user to map a Linux user to. The last argument, newuser, is the Linux user you want mapped to the specified SELinux user.

3. To view the mapping between the Linux **newuser** user and user_u, use the semanage utility again:

```
~]# semanage login -l
Login Name
                     SELinux User
                                           MLS/MCS Range
                                                                Service
__default__
                     unconfined_u
                                           s0-s0:c0.c1023
newuser
                     user_u
                                           s0
                     unconfined_u
                                           s0-s0:c0.c1023
root
                                           s0-s0:c0.c1023
system_u
                     system_u
```

4. As root, assign a password to the Linux **newuser** user:

```
~]# passwd newuser
Changing password for user newuser.
New password: Enter a password
Retype new password: Enter the same password again
passwd: all authentication tokens updated successfully.
```

5. Log out of your current session, and log in as the Linux **newuser** user. Run the following command to view the **newuser**'s SELinux context:

```
~]$ id -Z
user_u:user_r:user_t:s0
```

6. Log out of the Linux **newuser**'s session, and log back in with your account. If you do not want the Linux **newuser** user, run the following command as root to remove it, along with its home directory:

```
~]# userdel -r newuser
```

As root, remove the mapping between the Linux **newuser** user and user_u:

```
~]# semanage login -d newuser

-]# semanage login -l

Login Name SELinux User MLS/MCS Range Service

__default__ unconfined_u s0-s0:c0.c1023 *
root unconfined_u s0-s0:c0.c1023 *
system_u system_u s0-s0:c0.c1023 *
```

6.4. Changing the Default Mapping

In Fedora, Linux users are mapped to the SELinux __default__ login by default (which is in turn mapped to the SELinux unconfined_u user). If you would like new Linux users, and Linux users not specifically mapped to an SELinux user to be confined by default, change the default mapping with the semanage login command.

For example, run the following command as root to change the default mapping from unconfined_u to user_u:

```
~]# semanage login -m -S targeted -s "user_u" -r s0 __default__
```

Verify the __default__ login is mapped to user_u:

```
~]# semanage login -l

Login Name SELinux User MLS/MCS Range Service

__default__ user_u s0-s0:c0.c1023 *
root unconfined_u s0-s0:c0.c1023 *
system_u system_u s0-s0:c0.c1023 *
```

If a new Linux user is created and an SELinux user is not specified, or if an existing Linux user logs in and does not match a specific entry from the **semanage login -1** output, they are mapped to user_u, as per the __**default**__ login.

To change back to the default behavior, run the following command as root to map the __default__ login to the SELinux unconfined_u user:

```
~]# semanage login -m -S targeted -s "unconfined_u" -r s0-s0:c0.c1023 __default__
```

6.5. xguest: Kiosk Mode

The *xguest* package provides a kiosk user account. This account is used to secure machines that people walk up to and use, such as those at libraries, banks, airports, information kiosks, and coffee shops. The kiosk user account is very limited: essentially, it only allows users to log in and use **Firefox** to browse Internet websites. Any changes made while logged in with this account, such as creating files or changing settings, are lost when you log out.

To set up the kiosk account:

1. As the root user, install the *xguest* package. Install dependencies as required:

```
~]# dnf install xguest
```

2. In order to allow the kiosk account to be used by a variety of people, the account is not password-protected, and as such, the account can only be protected if SELinux is running in enforcing mode. Before logging in with this account, use the getenforce utility to confirm that SELinux is running in enforcing mode:

```
~]$ getenforce
Enforcing
```

If this is not the case, see Section 4.4, "Permanent Changes in SELinux States and Modes" for information about changing to enforcing mode. It is not possible to log in with this account if SELinux is in permissive mode or disabled.

3. You can only log in to this account via the GNOME Display Manager (GDM). Once the *xguest* package is installed, a **Guest** account is added to the GDM login screen.

6.6. Booleans for Users Executing Applications

Not allowing Linux users to execute applications (which inherit users' permissions) in their home directories and the /tmp/ directory, which they have write access to, helps prevent flawed or malicious applications from modifying files that users own. In Fedora, by default, Linux users in the guest_t and xguest_t domains cannot execute applications in their home directories or /tmp/; however, by default, Linux users in the user_t and staff_t domains can.

Booleans are available to change this behavior, and are configured with the setsebool utility, which must be run as the root user. The **setsebool** -**P** command makes persistent changes. Do not use the -**P** option if you do not want changes to persist across reboots:

quest t

To *allow* Linux users in the guest_t domain to execute applications in their home directories and / tmp/:

```
~]# setsebool -P guest_exec_content on
```

xguest t

To *allow* Linux users in the xguest_t domain to execute applications in their home directories and / tmp/:

```
\sim]# setsebool -P xguest_exec_content on
```

user_t

To *prevent* Linux users in the user_t domain from executing applications in their home directories and /tmp/:

```
~]# setsebool -P user_exec_content off
```

staff t

To *prevent* Linux users in the staff_t domain from executing applications in their home directories and /tmp/:

```
~]# setsebool -P staff_exec_content off
```

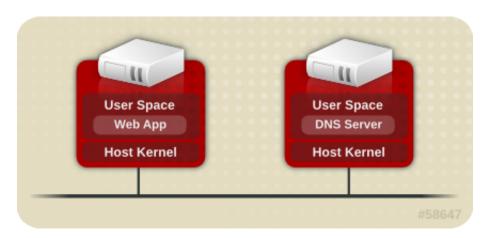
sVirt

sVirt is a technology included in Fedora that integrates SELinux and virtualization. sVirt applies Mandatory Access Control (MAC) to improve security when using virtual machines. The main reasons for integrating these technologies are to improve security and harden the system against bugs in the hypervisor that might be used as an attack vector aimed toward the host or to another virtual machine.

This chapter describes how sVirt integrates with virtualization technologies in Fedora.

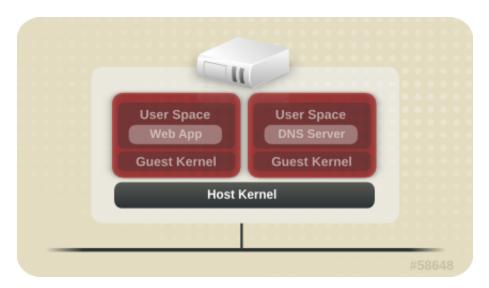
Non-Virtualized Environment

In a non-virtualized environment, hosts are separated from each other physically and each host has a self-contained environment, consisting of services such as a Web server, or a DNS server. These services communicate directly to their own user space, host kernel and physical host, offering their services directly to the network. The following image represents a non-virtualized environment:



Virtualized Environment

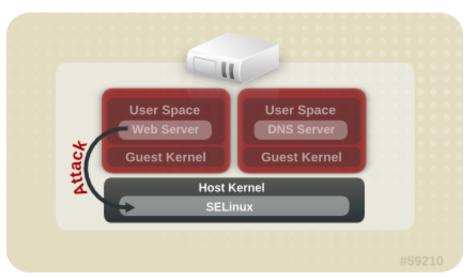
In a virtualized environment, several operating systems can be housed (as "guests") within a single host kernel and physical host. The following image represents a virtualized environment:



7.1. Security and Virtualization

When services are not virtualized, machines are physically separated. Any exploit is usually contained to the affected machine, with the obvious exception of network attacks. When services are grouped together in a virtualized environment, extra vulnerabilities emerge in the system. If there is a security flaw in the hypervisor that can be exploited by a guest instance, this guest may be able to not only attack the host, but also other guests running on that host. This is not theoretical; attacks already exist on hypervisors. These attacks can extend beyond the guest instance and could expose other guests to attack.

sVirt is an effort to isolate guests and limit their ability to launch further attacks if exploited. This is demonstrated in the following image, where an attack cannot break out of the virtual machine and extend to another host instance:



SELinux introduces a pluggable security framework for virtualized instances in its implementation of Mandatory Access Control (MAC). The sVirt framework allows guests and their resources to be uniquely labeled. Once labeled, rules can be applied which can reject access between different quests.

7.2. sVirt Labeling

Like other services under the protection of SELinux, sVirt uses process-based mechanisms and restrictions to provide an extra layer of security over guest instances. Under typical use, you should not even notice that sVirt is working in the background. This section describes the labeling features of sVirt.

As shown in the following output, when using sVirt, each Virtual Machine (VM) process is labeled and runs with a dynamically generated level. Each process is isolated from other VMs with different levels:

```
~]# ps -eZ | grep qemu

system_u:system_r:svirt_t:s0:c87,c520 27950 ? 00:00:17 qemu-kvm

system_u:system_r:svirt_t:s0:c639,c757 27989 ? 00:00:06 qemu-system-x86
```

The actual disk images are automatically labeled to match the processes, as shown in the following output:

```
~]# ls -lZ /var/lib/libvirt/images/*
```

```
system_u:object_r:svirt_image_t:s0:c87,c520 image1
```

The following table outlines the different labels that can be assigned when using sVirt:

Table 7.1. sVirt Labels

Туре	SELinux Context	Description
Virtual Machine Processes	system_u:system_r:svirt_t:MCS1	MCS1 is a randomly selected MCS field. Currently approximately 500,000 labels are supported.
Virtual Machine Image	system_u:object_r:svirt_image_t:	MDS/Aprocesses labeled svirt_t with the same MCS fields are able to read/write these image files and devices.
Virtual Machine Shared Read/ Write Content	system_u:object_r:svirt_image_t:	sall processes labeled svirt_t are allowed to write to the svirt_image_t:s0 files and devices.
Virtual Machine Image	system_u:object_r:virt_content_t:	ssystem default label used when an image exits. No svirt_t virtual processes are allowed to read files/devices with this label.

It is also possible to perform static labeling when using sVirt. Static labels allow the administrator to select a specific label, including the MCS/MLS field, for a virtual machine. Administrators who run statically-labeled virtual machines are responsible for setting the correct label on the image files. The virtual machine will always be started with that label, and the sVirt system will never modify the label of a statically-labeled virtual machine's content. This allows the sVirt component to run in an MLS environment. You can also run multiple virtual machines with different sensitivity levels on a system, depending on your requirements.

Secure Linux Containers

Linux Containers (LXC) is a low-level virtualization feature that allows you to run multiple copies of the same service at the same time on a system. Compared to full virtualization, containers do not require an entire new system to boot, can use less memory, and can use the base operating system in a read-only manner. For example, LXC allow you to run multiple web servers simultaneously, each with their own data while sharing the system data, and even running as the root user. However, running a privileged process within a container could affect other processes running outside of the container or processes running in other containers. Secure Linux containers use the SELinux context, therefore preventing the processes running within them from interacting with each other or with the host.

The **Docker** application is the main utility for managing Linux Containers in Fedora. As an alternative, you can also use the virsh command-line utility provided by the *libvirt* package.

For further details about Linux Containers see the Resource Management and Linux Containers Guide¹.

¹ http://access.redhat.com/site/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/Resource_Management_and_Linux_Containers_Guide/index.html

SELinux systemd Access Control

In Fedora 24, system services are controlled by the systemd daemon. In previous releases of Fedora, daemons could be started in two ways:

 At boot time, the System V init daemon launched an init.rc script and then this script launched the desired daemon. For example, the Apache server, which was started at boot, got the following SELinux label:

```
system_u:system_r:httpd_t:s0
```

 An administrator launched the init.rc script manually, causing the daemon to run. For example, when the service httpd restart command was invoked on the Apache server, the resulting SELinux label looked as follows:

```
unconfined_u:system_r:httpd_t:s0
```

When launched manually, the process adopted the user portion of the SELinux label that started it, making the labeling in the two scenarios above inconsistent. With the systemd daemon, the transitions are very different. As systemd handles all the calls to start and stop daemons on the system, using the init_t type, it can override the user part of the label when a daemon is restarted manually. As a result, the labels in both scenarios above are <code>system_u:system_r:httpd_t:s0</code> as expected and the SELinux policy could be improved to govern which domains are able to control which units.

9.1. SELinux Access Permissions for Services

In previous versions of Fedora, an administrator was able to control, which users or applications were able to start or stop services based on the label of the System V Init script. Now, systemd starts and stops all services, and users and processes communicate with systemd using the systemctl utility. The systemd daemon has the ability to consult the SELinux policy and check the label of the calling process and the label of the unit file that the caller tries to manage, and then ask SELinux whether or not the caller is allowed the access. This approach strengthens access control to critical system capabilities, which include starting and stopping system services.

For example, previously, administrators had to allow NetworkManager to execute systemctl to send a D-Bus message to systemd, which would in turn start or stop whatever service NetworkManager requested. In fact, NetworkManager was allowed to do everything systemctl could do. It was also impossible to setup confined administrators so that they could start or stop just particular services.

To fix these issues, systemd also works as an SELinux Access Manager. It can retrieve the label of the process running systemctl or the process that sent a D-Bus message to systemd. The daemon then looks up the label of the unit file that the process wanted to configure. Finally, systemd can retrieve information from the kernel if the SELinux policy allows the specific access between the process label and the unit file label. This means a compromised application that needs to interact with systemd for a specific service can now be confined via SELinux. Policy writers can also use these fine-grained controls to confine administrators. Policy changes involve a new class called **service**, with the following permissions:

```
class service
{
    start
```

```
stop
status
reload
kill
load
enable
disable
}
```

For example, a policy writer can now allow a domain to get the status of a service or start and stop a service, but not enable or disable a service. Access control operations in SELinux and systemd do not match in all cases. A mapping was defined to line up systemd method calls with SELinux access checks. Table 9.1, "Mapping of systemd unit file method calls on SELinux access checks" maps access checks on unit files while Table 9.2, "Mapping of systemd general system calls on SELinux access checks" covers access checks for the system in general. If no match is found in either table, then the undefined system check is called.

Table 9.1. Mapping of systemd unit file method calls on SELinux access checks

systemd unit file method	SELinux access check
DisableUnitFiles	disable
EnableUnitFiles	enable
GetUnit	status
GetUnitByPID	status
GetUnitFileState	status
Kill	stop
KillUnit	stop
LinkUnitFiles	enable
ListUnits	status
LoadUnit	status
MaskUnitFiles	disable
PresetUnitFiles	enable
ReenableUnitFiles	enable
Reexecute	start
Reload	reload
ReloadOrRestart	start
ReloadOrRestartUnit	start
ReloadOrTryRestart	start
ReloadOrTryRestartUnit	start
ReloadUnit	reload
ResetFailed	stop
ResetFailedUnit	stop
Restart	start
RestartUnit	start
Start	start
StartUnit	start
StartUnitReplace	start
	

systemd unit file method	SELinux access check
Stop	stop
StopUnit	stop
TryRestart	start
TryRestartUnit	start
UnmaskUnitFiles	enable

Table 9.2. Mapping of systemd general system calls on SELinux access checks

systemd general system call	SELinux access check
ClearJobs	reboot
FlushDevices	halt
Get	status
GetAll	status
GetJob	status
GetSeat	status
GetSession	status
GetSessionByPID	status
GetUser	status
Halt	halt
Introspect	status
KExec	reboot
KillSession	halt
KillUser	halt
ListJobs	status
ListSeats	status
ListSessions	status
ListUsers	status
LockSession	halt
PowerOff	halt
Reboot	reboot
SetUserLinger	halt
TerminateSeat	halt
TerminateSession	halt
TerminateUser	halt

Example 9.1. SELinux Policy for a System Service

By using the sesearch utility, you can list policy rules for a system service. For example, calling the sesearch -A -s NetworkManager_t -c service command returns:

```
allow NetworkManager_t dnsmasq_unit_file_t : service { start stop status reload kill
load };
allow NetworkManager_t nscd_unit_file_t : service { start stop status reload kill
load };
```

```
allow NetworkManager_t ntpd_unit_file_t : service { start stop status reload kill
load };
allow NetworkManager_t pppd_unit_file_t : service { start stop status reload kill
load };
allow NetworkManager_t polipo_unit_file_t : service { start stop status reload kill
load };
```

9.2. SELinux and journald

In systemd, the journald daemon (also known as systemd-journal) is the alternative for the syslog utility, which is a system service that collects and stores logging data. It creates and maintains structured and indexed journals based on logging information that is received from the kernel, from user processes using the libc syslog() function, from standard and error output of system services, or using its native API. It implicitly collects numerous metadata fields for each log message in a secure way.

The systemd-journal service can be used with SELinux to increase security. SELinux controls processes by only allowing them to do what they were designed to do; sometimes even less, depending on the security goals of the policy writer. For example, SELinux prevents a compromised ntpd process from doing anything other than handle Network Time. However, the ntpd process sends syslog messages, so that SELinux would allow the compromised process to continue to send those messages. The compromised ntpd could format syslog messages to match other daemons and potentially mislead an administrator, or even worse, a utility that reads the syslog file into compromising the whole system.

The systemd-journal daemon verifies all log messages and, among other things, adds SELinux labels to them. It is then easy to detect inconsistencies in log messages and prevent an attack of this type before it occurs. You can use the journalctl utility to query logs of systemd journals. If no command-line arguments are specified, running this utility lists the full content of the journal, starting from the oldest entries. To see all logs generated on the system, including logs for system components, execute journalctl as root. If you execute it as a non-root user, the output will be limited only to logs related to the currently logged-in user.

Example 9.2. Listing Logs with journalctl

It is possible to use journalctl for listing all logs related to a particular SELinux label. For example, the following command lists all logs logged under the system_u:system_r:policykit_t:s0 label:

```
~]# journalctl _SELINUX_CONTEXT=system_u:system_r:policykit_t:s0
Oct 21 10:22:42 localhost.localdomain polkitd[647]: Started polkitd version 0.112
Oct 21 10:22:44 localhost.localdomain polkitd[647]: Loading rules from directory /etc/
polkit-1/rules.d
Oct 21 10:22:44 localhost.localdomain polkitd[647]: Loading rules from directory /usr/
share/polkit-1/rules.d
Oct 21 10:22:44 localhost.localdomain polkitd[647]: Finished loading, compiling and
executing 5 rules
Oct 21 10:22:44 localhost.localdomain polkitd[647]: Acquired the name
org.freedesktop.PolicyKit1 on the system bus Oct 21 10:23:10 localhost polkitd[647]:
Registered Authentication Agent for unix-session:c1 (system bus name :1.49, object path /
org/freedesktop/PolicyKit1/AuthenticationAgent, locale en_US.UTF-8) (disconnected from
bus)
Oct 21 10:23:35 localhost polkitd[647]: Unregistered Authentication Agent for unix-
session:c1 (system bus name :1.80 [/usr/bin/gnome-shell --mode=classic], object path /org/
freedesktop/PolicyKit1/AuthenticationAgent, locale en_US.utf8)
```

For more information about journalctl, see the journalctl(1) manual page.

Troubleshooting

The following chapter describes what happens when SELinux denies access; the top three causes of problems; where to find information about correct labeling; analyzing SELinux denials; and creating custom policy modules with audit2allow.

10.1. What Happens when Access is Denied

SELinux decisions, such as allowing or disallowing access, are cached. This cache is known as the Access Vector Cache (AVC). Denial messages are logged when SELinux denies access. These denials are also known as "AVC denials", and are logged to a different location, depending on which daemons are running:

Daemon Log Location

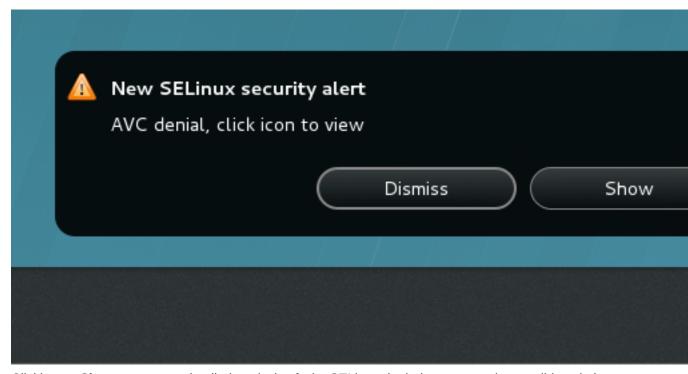
auditd on /var/log/audit/audit.log

auditd off; rsyslogd on /var/log/messages

setroubleshootd, rsyslogd, and /var/log/audit/audit.log. Easier-to-read denial messages

auditd on also sent to /var/log/messages

If you are running the X Window System, have the *setroubleshoot* and *setroubleshoot-server* packages installed, and the <code>setroubleshootd</code> and auditd daemons are running, a warning is displayed when access is denied by SELinux:



Clicking on **Show** presents a detailed analysis of why SELinux denied access, and a possible solution for allowing access. If you are not running the X Window System, it is less obvious when access is denied by SELinux. For example, users browsing your website may receive an error similar to the following:

Forbidden

You don't have permission to access file name on this server

For these situations, if DAC rules (standard Linux permissions) allow access, check /var/log/messages and /var/log/audit/audit.log for "SELinux is preventing" and "denied" errors respectively. This can be done by running the following commands as the root user:

```
~]# grep "SELinux is preventing" /var/log/messages
```

```
~]# grep "denied" /var/log/audit/audit.log
```

10.2. Top Three Causes of Problems

The following sections describe the top three causes of problems: labeling problems, configuring Booleans and ports for services, and evolving SELinux rules.

10.2.1. Labeling Problems

On systems running SELinux, all processes and files are labeled with a label that contains security-relevant information. This information is called the SELinux context. If these labels are wrong, access may be denied. If an application is labeled incorrectly, the process it transitions to may not have the correct label, possibly causing SELinux to deny access, and the process being able to create mislabeled files.

A common cause of labeling problems is when a non-standard directory is used for a service. For example, instead of using <code>/var/www/html/</code> for a website, an administrator wants to use <code>/srv/myweb/</code>. On Fedora, the <code>/srv/</code> directory is labeled with the <code>var_t</code> type. Files and directories created and <code>/srv/</code> inherit this type. Also, newly-created top-level directories (such as <code>/myserver/</code>) may be labeled with the <code>default_t</code> type. SELinux prevents the Apache HTTP Server (<code>httpd</code>) from accessing both of these types. To allow access, SELinux must know that the files in <code>/srv/myweb/</code> are to be accessible to <code>httpd</code>:

```
~]# semanage fcontext -a -t httpd_sys_content_t "/srv/myweb(/.*)?"
```

This **semanage** command adds the context for the /**srv/myweb**/ directory (and all files and directories under it) to the SELinux file-context configuration¹. The semanage utility does not change the context. As root, run the restorecon utility to apply the changes:

```
~]# restorecon -R -v /srv/myweb
```

See Section 4.7.2, "Persistent Changes: semanage fcontext" for further information about adding contexts to the file-context configuration.

10.2.1.1. What is the Correct Context?

The matchpathcon utility checks the context of a file path and compares it to the default label for that path. The following example demonstrates using matchpathcon on a directory that contains incorrectly labeled files:

¹ Files in /etc/selinux/targeted/contexts/files/ define contexts for files and directories. Files in this directory are read by the restorecon and setfiles utilities to restore files and directories to their default contexts.

```
~]$ matchpathcon -V /var/www/html/*
/var/www/html/index.html has context unconfined_u:object_r:user_home_t:s0, should be
system_u:object_r:httpd_sys_content_t:s0
/var/www/html/page1.html has context unconfined_u:object_r:user_home_t:s0, should be
system_u:object_r:httpd_sys_content_t:s0
```

In this example, the **index.html** and **page1.html** files are labeled with the user_home_t type. This type is used for files in user home directories. Using the **mv** command to move files from your home directory may result in files being labeled with the user_home_t type. This type should not exist outside of home directories. Use the restorecon utility to restore such files to their correct type:

```
~]# restorecon -v /var/www/html/index.html
restorecon reset /var/www/html/index.html context unconfined_u:object_r:user_home_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
```

To restore the context for all files under a directory, use the **-R** option:

```
~]# restorecon -R -v /var/www/html/
restorecon reset /var/www/html/page1.html context unconfined_u:object_r:samba_share_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
restorecon reset /var/www/html/index.html context unconfined_u:object_r:samba_share_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
```

See Section 4.10.3, "Checking the Default SELinux Context" for a more detailed example of matchpathcon.

10.2.2. How are Confined Services Running?

Services can be run in a variety of ways. To cater for this, you need to specify how you run your services. This can be achieved via Booleans that allow parts of SELinux policy to be changed at runtime, without any knowledge of SELinux policy writing. This allows changes, such as allowing services access to NFS volumes, without reloading or recompiling SELinux policy. Also, running services on non-default port numbers requires policy configuration to be updated via the **semanage** command.

For example, to allow the Apache HTTP Server to communicate with MariaDB, enable the httpd_can_network_connect_db Boolean:

```
~]# setsebool -P httpd_can_network_connect_db on
```

If access is denied for a particular service, use the getsebool and grep utilities to see if any Booleans are available to allow access. For example, use the **getsebool -a | grep ftp** command to search for FTP related Booleans:

```
~]$ getsebool -a | grep ftp

ftpd_anon_write --> off

ftpd_full_access --> off

ftpd_use_cifs --> off

ftpd_use_nfs --> off

ftp_home_dir --> off

ftpd_connect_db --> off

httpd_enable_ftp_server --> off

tftp_anon_write --> off
```

For a list of Booleans and whether they are on or off, run the **getsebool** -a command. For a list of Booleans, an explanation of what each one is, and whether they are on or off, run the **semanage boolean** -1 command as root. See *Section 4.6, "Booleans"* for information about listing and configuring Booleans.

Port Numbers

Depending on policy configuration, services may only be allowed to run on certain port numbers. Attempting to change the port a service runs on without changing policy may result in the service failing to start. For example, run the **semanage port -1 | grep http** command as root to list http related ports:

```
~]# semanage port -1 | grep http
http_cache_port_t
                                        3128, 8080, 8118
                               tcp
http_cache_port_t
                               udp
                                        3130
http_port_t
                               tcp
                                        80, 443, 488, 8008, 8009, 8443
pegasus_http_port_t
                               tcp
                                        5988
pegasus_https_port_t
                                        5989
                               tcp
```

The http_port_t port type defines the ports Apache HTTP Server can listen on, which in this case, are TCP ports 80, 443, 488, 8008, 8009, and 8443. If an administrator configures **httpd.conf** so that httpd listens on port 9876 (**Listen 9876**), but policy is not updated to reflect this, the following command fails:

```
~]# systemctl start httpd.service
Job for httpd.service failed. See 'systemctl status httpd.service' and 'journalctl -xn' for details.
```

```
~]# systemctl status httpd.service
httpd.service - The Apache HTTP Server
Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled)
Active: failed (Result: exit-code) since Thu 2013-08-15 09:57:05 CEST; 59s ago
Process: 16874 ExecStop=/usr/sbin/httpd $OPTIONS -k graceful-stop (code=exited, status=0/SUCCESS)
Process: 16870 ExecStart=/usr/sbin/httpd $OPTIONS -DFOREGROUND (code=exited, status=1/FAILURE)
```

An SELinux denial message similar to the following is logged to /var/log/audit/audit.log:

```
type=AVC msg=audit(1225948455.061:294): avc: denied { name_bind } for
pid=4997 comm="httpd" src=9876 scontext=unconfined_u:system_r:httpd_t:s0
tcontext=system_u:object_r:port_t:s0 tclass=tcp_socket
```

To allow httpd to listen on a port that is not listed for the http_port_t port type, run the **semanage port** command to add a port to policy configuration²:

```
~]# semanage port -a -t http_port_t -p tcp 9876
```

² The **semanage port -a** command adds an entry to the **/etc/selinux/targeted/modules/active/ports.local** file. Note that by default, this file can only be viewed by root.

The **-a** option adds a new record; the **-t** option defines a type; and the **-p** option defines a protocol. The last argument is the port number to add.

10.2.3. Evolving Rules and Broken Applications

Applications may be broken, causing SELinux to deny access. Also, SELinux rules are evolving – SELinux may not have seen an application running in a certain way, possibly causing it to deny access, even though the application is working as expected. For example, if a new version of PostgreSQL is released, it may perform actions the current policy has not seen before, causing access to be denied, even though access should be allowed.

For these situations, after access is denied, use the audit2allow utility to create a custom policy module to allow access. See *Section 10.3.8*, "Allowing Access: audit2allow" for information about using audit2allow.

10.3. Fixing Problems

The following sections help troubleshoot issues. They go over: checking Linux permissions, which are checked before SELinux rules; possible causes of SELinux denying access, but no denials being logged; manual pages for services, which contain information about labeling and Booleans; permissive domains, for allowing one process to run permissive, rather than the whole system; how to search for and view denial messages; analyzing denials; and creating custom policy modules with audit2allow.

10.3.1. Linux Permissions

When access is denied, check standard Linux permissions. As mentioned in *Chapter 1, Introduction*, most operating systems use a Discretionary Access Control (DAC) system to control access, allowing users to control the permissions of files that they own. SELinux policy rules are checked after DAC rules. SELinux policy rules are not used if DAC rules deny access first.

If access is denied and no SELinux denials are logged, use the following command to view the standard Linux permissions:

```
~]$ ls -l /var/www/html/index.html
-rw-r---- 1 root root 0 2009-05-07 11:06 index.html
```

In this example, index.htm1 is owned by the root user and group. The root user has read and write permissions (-rw), and members of the root group have read permissions (-r-). Everyone else has no access (---). By default, such permissions do not allow httpd to read this file. To resolve this issue, use the **chown** command to change the owner and group. This command must be run as root:

```
~]# chown apache:apache /var/www/html/index.html
```

This assumes the default configuration, in which httpd runs as the Linux Apache user. If you run httpd with a different user, replace **apache**: apache with that user.

See the *Fedora Documentation Project "Permissions"* draft for information about managing Linux permissions.

³ http://fedoraproject.org/wiki/Docs/Drafts/AdministrationGuide/Permissions

10.3.2. Possible Causes of Silent Denials

In certain situations, AVC denial messages may not be logged when SELinux denies access. Applications and system library functions often probe for more access than required to perform their tasks. To maintain least privilege without filling audit logs with AVC denials for harmless application probing, the policy can silence AVC denials without allowing a permission by using **dontaudit** rules. These rules are common in standard policy. The downside of **dontaudit** is that, although SELinux denies access, denial messages are not logged, making troubleshooting more difficult.

To temporarily disable **dontaudit** rules, allowing all denials to be logged, run the following command as root:

```
~]# semodule -DB
```

The **-D** option disables **dontaudit** rules; the **-B** option rebuilds policy. After running **semodule -DB**, try exercising the application that was encountering permission problems, and see if SELinux denials — relevant to the application — are now being logged. Take care in deciding which denials should be allowed, as some should be ignored and handled via **dontaudit** rules. If in doubt, or in search of guidance, contact other SELinux users and developers on an SELinux list, such as *fedora-selinux-list*⁴.

To rebuild policy and enable **dontaudit** rules, run the following command as root:

```
~]# semodule -B
```

This restores the policy to its original state. For a full list of **dontaudit** rules, run the **sesearch -- dontaudit** command. Narrow down searches using the **-s** *domain* option and the **grep** command.
For example:

```
~]$ sesearch --dontaudit -s smbd_t | grep squid
dontaudit smbd_t squid_port_t : tcp_socket name_bind ;
dontaudit smbd_t squid_port_t : udp_socket name_bind ;
```

See Section 10.3.6, "Raw Audit Messages" and Section 10.3.7, "sealert Messages" for information about analyzing denials.

10.3.3. Manual Pages for Services

Manual pages for services contain valuable information, such as what file type to use for a given situation, and Booleans to change the access a service has (such as httpd accessing NFS volumes). This information may be in the standard manual page or in the manual page that can be automatically generated from the SELinux policy for every service domain using the sepolicy manpage utility. Such manual pages are named in the <code>service-name_selinux</code> format. Such manual pages are also shipped with the <code>selinux-policy-devel</code> package. The <code>selinux-policy-doc</code> package includes the HTML version of the <code>service-name_selinux</code> manual pages.

For example, the httpd_selinux(8) manual page has information about what file type to use for a given situation, as well as Booleans to allow scripts, sharing files, accessing directories inside user home directories, and so on. Other manual pages with SELinux information for services include:

• Samba: the samba_selinux(8) manual page for example describes that enabling the samba_enable_home_dirs Boolean allows Samba to share users home directories.

⁴ http://www.redhat.com/mailman/listinfo/fedora-selinux-list

• NFS: the nfsd_selinux(8) manual page describes SELinux nfsd policy that allows users to setup their nfsd processes in as secure a method as possible.

The information in manual pages helps you configure the correct file types and Booleans, helping to prevent SELinux from denying access.

See Section 5.5, "Generating Manual Pages: **sepolicy manpage**" for further information about sepolicy manpage.

10.3.4. Permissive Domains

When SELinux is running in permissive mode, SELinux does not deny access, but denials are logged for actions that would have been denied if running in enforcing mode. Previously, it was not possible to make a single domain permissive (remember: processes run in domains). In certain situations, this led to making the whole system permissive to troubleshoot issues.

Permissive domains allow an administrator to configure a single process (domain) to run permissive, rather than making the whole system permissive. SELinux checks are still performed for permissive domains; however, the kernel allows access and reports an AVC denial for situations where SELinux would have denied access.

Permissive domains have the following uses:

- They can be used for making a single process (domain) run permissive to troubleshoot an issue without putting the entire system at risk by making it permissive.
- They allow an administrator to create policies for new applications. Previously, it was recommended that a minimal policy be created, and then the entire machine put into permissive mode, so that the application could run, but SELinux denials still logged. The audit2allow could then be used to help write the policy. This put the whole system at risk. With permissive domains, only the domain in the new policy can be marked permissive, without putting the whole system at risk.

10.3.4.1. Making a Domain Permissive

To make a domain permissive, run the **semanage permissive** -a *domain* command, where *domain* is the domain you want to make permissive. For example, run the following command as root to make the httpd_t domain (the domain the Apache HTTP Server runs in) permissive:

```
~]# semanage permissive -a httpd_t
```

To view a list of domains you have made permissive, run the **semodule -1 | grep permissive** command as root. For example:

```
~]# semodule -1 | grep permissive
permissive_httpd_t 1.0
permissivedomains 1.0.0
```

If you no longer want a domain to be permissive, run the **semanage permissive** -d **domain** command as root. For example:

```
~]# semanage permissive -d httpd_t
```

10.3.4.2. Disabling Permissive Domains

The permissivedomains.pp module contains all of the permissive domain declarations that are presented on the system. To disable all permissive domains, run the following command as root:

```
~]# semodule -d permissivedomains
```

10.3.4.3. Denials for Permissive Domains

The **SYSCALL** message is different for permissive domains. The following is an example AVC denial (and the associated system call) from the Apache HTTP Server:

```
type=AVC msg=audit(1226882736.442:86): avc: denied { getattr } for pid=2427 comm="httpd" path="/var/www/html/file1" dev=dm-0 ino=284133 scontext=unconfined_u:system_r:httpd_t:s0 tcontext=unconfined_u:object_r:samba_share_t:s0 tclass=file
```

type=SYSCALL msg=audit(1226882736.442:86): arch=40000003 syscall=196 success=no exit=-13 a0=b9a1e198 a1=bfc2921c a2=54dff4 a3=2008171 items=0 ppid=2425 pid=2427 auid=502 uid=48 gid=48 euid=48 suid=48 fsuid=48 egid=48 sgid=48 fsgid=48 tty=(none) ses=4 comm="httpd" exe="/usr/sbin/httpd" subj=unconfined_u:system_r:httpd_t:s0 key=(null)

By default, the httpd_t domain is not permissive, and as such, the action is denied, and the SYSCALL message contains success=no. The following is an example AVC denial for the same situation, except the semanage permissive -a httpd_t command has been run to make the httpd_t domain permissive:

```
type=AVC msg=audit(1226882925.714:136): avc: denied { read } for pid=2512
comm="httpd" name="file1" dev=dm-0 ino=284133 scontext=unconfined_u:system_r:httpd_t:s0
tcontext=unconfined_u:object_r:samba_share_t:s0 tclass=file

type=SYSCALL msg=audit(1226882925.714:136): arch=40000003 syscall=5 success=yes exit=11
a0=b962a1e8 a1=8000 a2=0 a3=8000 items=0 ppid=2511 pid=2512 auid=502 uid=48 gid=48 euid=48
suid=48 fsuid=48 egid=48 sgid=48 fsgid=48 tty=(none) ses=4 comm="httpd" exe="/usr/sbin/httpd" subj=unconfined_u:system_r:httpd_t:s0 key=(null)
```

In this case, although an AVC denial was logged, access was not denied, as shown by **success=yes** in the **SYSCALL** message.

See Dan Walsh's "Permissive Domains" blog entry for further information about permissive domains.

10.3.5. Searching For and Viewing Denials

This section assumes the *setroubleshoot*, *setroubleshoot-server*, *dbus* and *audit* packages are installed, and that the auditd, rsyslogd, and setroubleshootd daemons are running. See *Section 4.2*, "Which Log File is Used" for information about starting these daemons. A number of utilites are available for searching for and viewing SELinux AVC messages, such as ausearch, aureport, and sealert.

⁵ http://danwalsh.livejournal.com/24537.html

ausearch

The *audit* package provides the **ausearch** utility that can query the audit daemon logs based for events based on different search criteria. The ausearch utility accesses **/var/log/audit/audit.log**, and as such, must be run as the root user:

Searching For Command

all denials ausearch -m avc

denials for that today ausearch -m avc -ts today denials from the last 10 minutes ausearch -m avc -ts recent

To search for SELinux AVC messages for a particular service, use the **-c** *comm-name* option, where *comm-name* is the executable's name, for example, httpd for the Apache HTTP Server, and smbd for Samba:

```
~]# ausearch -m avc -c httpd
```

```
~]# ausearch -m avc -c smbd
```

With each **ausearch** command, it is advised to use either the **--interpret** (**-i**) option for easier readability, or the **--raw** (**-r**) option for script processing. See the ausearch(8) manual page for further **ausearch** options.

aureport

The *audit* package provides the aureport utility, which produces summary reports of the audit system logs. ⁷ The aureport utility accesses /var/log/audit/audit.log, and as such, must be run as the root user. To view a list of SELinux denial messages and how often each one occurred, run the aureport -a command. The following is example output that includes two denials:

sealert

The setroubleshoot-server package provides the sealert utility, which reads denial messages translated by setroubleshoot-server. Denials are assigned IDs, as seen in /var/log/messages. The following is an example denial from messages:

⁶ See the ausearch(8) manual page for further information about ausearch.

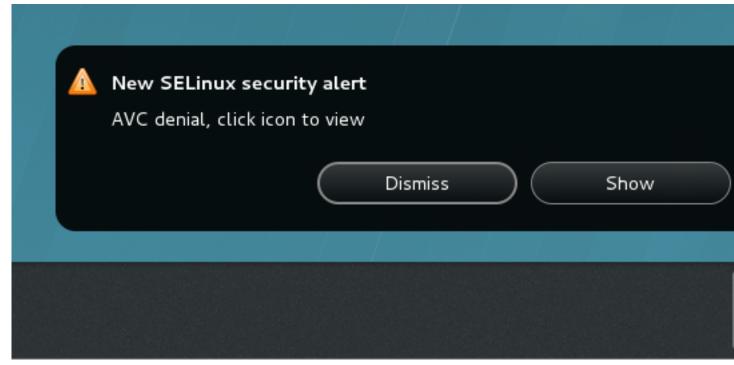
 $^{^{7}}$ See the aureport(8) manual page for further information about aureport.

⁸ See the sealert(8) manual page for further information about **sealert**.

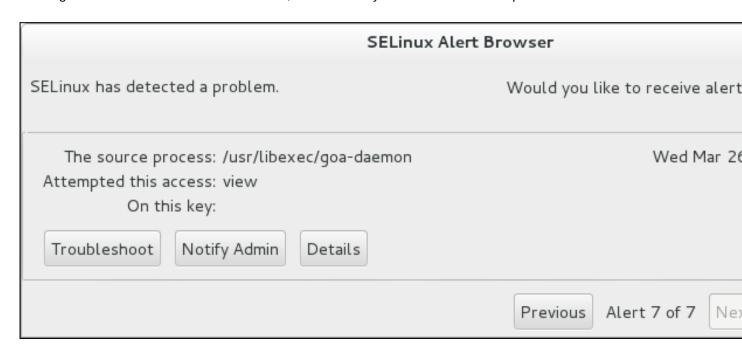
```
setroubleshoot: SELinux is preventing /usr/sbin/httpd from name_bind access on the tcp_socket. For complete SELinux messages. run sealert -l 8c123656-5dda-4e5d-8791-9e3bd03786b7
```

In this example, the denial ID is **8c123656-5dda-4e5d-8791-9e3bd03786b7**. The **-1** option takes an ID as an argument. Running the **sealert -1 8c123656-5dda-4e5d-8791-9e3bd03786b7** command presents a detailed analysis of why SELinux denied access, and a possible solution for allowing access.

If you are running the X Window System, have the *setroubleshoot* and *setroubleshoot-server* packages installed, and the <code>setroubleshootd</code>, dbus and auditd daemons are running, a warning is displayed when access is denied by SELinux:



Clicking on **Show** launches the sealert GUI, which allows you to troubleshoot the problem:



Alternatively, run the **sealert** -**b** command to launch the **sealert** GUI. To view a detailed analysis of all denial messages, run the **sealert** -**1** * command.

10.3.6. Raw Audit Messages

Raw audit messages are logged to <code>/var/log/audit/audit.log</code>. The following is an example AVC denial message (and the associated system call) that occurred when the Apache HTTP Server (running in the <code>httpd_t</code> domain) attempted to access the <code>/var/www/html/file1</code> file (labeled with the <code>samba_share_t</code> type):

```
type=AVC msg=audit(1226874073.147:96): avc: denied { getattr } for pid=2465 comm="httpd"
path="/var/www/html/file1" dev=dm-0 ino=284133 scontext=unconfined_u:system_r:httpd_t:s0
tcontext=unconfined_u:object_r:samba_share_t:s0 tclass=file
```

type=SYSCALL msg=audit(1226874073.147:96): arch=40000003 syscall=196 success=no exit=-13
a0=b98df198 a1=bfec85dc a2=54dff4 a3=2008171 items=0 ppid=2463 pid=2465 auid=502 uid=48
gid=48 euid=48 suid=48 fsuid=48 egid=48 sgid=48 fsgid=48 tty=(none) ses=6 comm="httpd"
exe="/usr/sbin/httpd" subj=unconfined_u:system_r:httpd_t:s0 key=(null)

{ getattr }

The item in the curly brackets indicates the permission that was denied. The **getattr** entry indicates the source process was trying to read the target file's status information. This occurs before reading files. This action is denied due to the file being accessed having a wrong label. Commonly seen permissions include **getattr**, **read**, and **write**.

comm="httpd"

The executable that launched the process. The full path of the executable is found in the **exe**= section of the system call (SYSCALL) message, which in this case, is **exe**="/usr/sbin/httpd".

path="/var/www/html/file1"

The path to the object (target) the process attempted to access.

```
scontext="unconfined_u:system_r:httpd_t:s0"
```

The SELinux context of the process that attempted the denied action. In this case, it is the SELinux context of the Apache HTTP Server, which is running in the httpd_t domain.

```
tcontext="unconfined_u:object_r:samba_share_t:s0"
```

The SELinux context of the object (target) the process attempted to access. In this case, it is the SELinux context of **file1**. Note that the samba_share_t type is not accessible to processes running in the httpd_t domain.

In certain situations, the **tcontext** may match the **scontext**, for example, when a process attempts to execute a system service that will change characteristics of that running process, such as the user ID. Also, the **tcontext** may match the **scontext** when a process tries to use more resources (such as memory) than normal limits allow, resulting in a security check to see if that process is allowed to break those limits.

From the system call (SYSCALL) message, two items are of interest:

- **success=no**: indicates whether the denial (AVC) was enforced or not. **success=no** indicates the system call was not successful (SELinux denied access). **success=yes** indicates the system call was successful. This can be seen for permissive domains or unconfined domains, such as unconfined_service_t and kernel_t.
- exe="/usr/sbin/httpd": the full path to the executable that launched the process, which in this case, is exe="/usr/sbin/httpd".

An incorrect file type is a common cause for SELinux denying access. To start troubleshooting, compare the source context (**scontext**) with the target context (**tcontext**). Should the process (**scontext**) be accessing such an object (**tcontext**)? For example, the Apache HTTP Server (httpd_t) should only be accessing types specified in the httpd_selinux(8) manual page, such as httpd_sys_content_t, public_content_t, and so on, unless configured otherwise.

10.3.7. sealert Messages

Denials are assigned IDs, as seen in /var/log/messages. The following is an example AVC denial (logged to messages) that occurred when the Apache HTTP Server (running in the httpd_t domain) attempted to access the /var/www/html/file1 file (labeled with the samba_share_t type):

```
hostname setroubleshoot: SELinux is preventing httpd (httpd_t) "getattr" to /var/www/html/file1 (samba_share_t). For complete SELinux messages. run sealert -l 84e0b04d-d0ad-4347-8317-22e74f6cd020
```

As suggested, run the **sealert -1 84e0b04d-d0ad-4347-8317-22e74f6cd020** command to view the complete message. This command only works on the local machine, and presents the same information as the **sealert** GUI:

```
~]$ sealert -1 84e0b04d-d0ad-4347-8317-22e74f6cd020
Summary:
SELinux is preventing httpd (httpd_t) "getattr" to /var/www/html/file1
(samba_share_t).
Detailed Description:
SELinux denied access to /var/www/html/file1 requested by httpd.
/var/www/html/file1 has a context used for sharing by different program. If you
would like to share /var/www/html/file1 from httpd also, you need to change its
file context to public_content_t. If you did not intend to this access, this
could signal a intrusion attempt.
Allowing Access:
You can alter the file context by executing chcon -t public_content_t
'/var/www/html/file1'
Fix Command:
chcon -t public_content_t '/var/www/html/file1'
Additional Information:
                              unconfined_u:system_r:httpd_t:s0
Source Context
Target Context
                              unconfined_u:object_r:samba_share_t:s0
Target Objects
                              /var/www/html/file1 [ file ]
Source
                              httpd
Source Path
                              /usr/sbin/httpd
Port
                              <Unknown>
Host
                              hostname
                              httpd-2.2.10-2
Source RPM Packages
Target RPM Packages
Policy RPM
                              selinux-policy-3.5.13-11.fc12
Selinux Enabled
                              True
Policy Type
                              targeted
MLS Enabled
                              True
Enforcing Mode
                              Enforcing
```

```
Plugin Name
                                                                                                   public_content
Host Name
                                                                                                   hostname
Platform
                                                                                                   Linux hostname 2.6.27.4-68.fc12.i686 #1 SMP Thu Oct
30 00:49:42 EDT 2008 1686 1686
Alert Count
First Seen
                                                                                                   Wed Nov
                                                                                                                               5 18:53:05 2008
Last Seen
                                                                                                   Wed Nov 5 01:22:58 2008
Local ID
                                                                                                   84e0b04d-d0ad-4347-8317-22e74f6cd020
Line Numbers
Raw Audit Messages
node=hostname type=AVC msg=audit(1225812178.788:101): avc: denied { getattr }
   for pid=2441 comm="httpd" path="/var/www/html/file1" dev=dm-0 ino=284916
   scontext=unconfined\_u: system\_r: httpd\_t: s0 tcontext=unconfined\_u: object\_r: samba\_share\_t: s0 tcontext=unconfined\_u: object\_r: s0 tcontext=unconfined\_u: s0 tcontex
   tclass=file
node=hostname type=SYSCALL msg=audit(1225812178.788:101): arch=40000003 syscall=196
   success=no exit=-13 a0=b8e97188 a1=bf87aaac a2=54dff4 a3=2008171 items=0 ppid=2439 pid=2441
   auid=502 uid=48 gid=48 euid=48 suid=48 fsuid=48 egid=48 sgid=48 fsgid=48 tty=(none) ses=3
   comm="httpd" exe="/usr/sbin/httpd" subj=unconfined_u:system_r:httpd_t:s0 key=(null)
```

Summary

A brief summary of the denied action. This is the same as the denial in /var/log/messages. In this example, the httpd process was denied access to a file (file1), which is labeled with the samba_share_t type.

Detailed Description

A more verbose description. In this example, **file1** is labeled with the samba_share_t type. This type is used for files and directories that you want to export via Samba. The description suggests changing the type to a type that can be accessed by the Apache HTTP Server and Samba, if such access is desired.

Allowing Access

A suggestion for how to allow access. This may be relabeling files, enabling a Boolean, or making a local policy module. In this case, the suggestion is to label the file with a type accessible to both the Apache HTTP Server and Samba.

Fix Command

A suggested command to allow access and resolve the denial. In this example, it gives the command to change the **file1** type to public_content_t, which is accessible to the Apache HTTP Server and Samba.

Additional Information

Information that is useful in bug reports, such as the policy package name and version (**selinux-policy-3.5.13-11.fc12**), but may not help towards solving why the denial occurred.

Raw Audit Messages

The raw audit messages from /var/log/audit/audit.log that are associated with the denial. See Section 10.3.6, "Raw Audit Messages" for information about each item in the AVC denial.

10.3.8. Allowing Access: audit2allow



Warning

Do not use the example in this section in production. It is used only to demonstrate the use of the audit2allow utility.

The audit2allow utility gathers information from logs of denied operations and then generates SELinux policy allow rules. After analyzing denial messages as per Section 10.3.7, "sealert Messages", and if no label changes or Booleans allowed access, use audit2allow to create a local policy module. When access is denied by SELinux, running audit2allow generates Type Enforcement rules that allow the previously denied access.

The following example demonstrates using audit2allow to create a policy module:

 A denial message and the associated system call are logged to the /var/log/audit/ audit.log file:

```
type=AVC msg=audit(1226270358.848:238): avc: denied { write }
for pid=13349 comm="certwatch" name="cache" dev=dm-0 ino=218171
scontext=system_u:system_r:certwatch_t:s0 tcontext=system_u:object_r:var_t:s0
tclass=dir

type=SYSCALL msg=audit(1226270358.848:238): arch=40000003 syscall=39 success=no exit=-13
a0=39a2bf a1=3ff a2=3a0354 a3=94703c8 items=0 ppid=13344 pid=13349 auid=4294967295
uid=0 gid=0 euid=0 suid=0 fsuid=0 egid=0 sgid=0 fsgid=0 tty=(none) ses=4294967295
comm="certwatch" exe="/usr/bin/certwatch" subj=system_u:system_r:certwatch_t:s0
key=(null)
```

In this example, **certwatch** was denied the write access to a directory labeled with the var_t type. Analyze the denial message as per *Section 10.3.7*, "sealert Messages". If no label changes or Booleans allowed access, use audit2allow to create a local policy module.

2. Run the following command to produce a human-readable description of why the access was denied. The audit2allow utility reads /var/log/audit/audit.log, and as such, must be run as the root user:

```
~]# audit2allow -w -a

type=AVC msg=audit(1226270358.848:238): avc: denied { write }

for pid=13349 comm="certwatch" name="cache" dev=dm-0 ino=218171

scontext=system_u:system_r:certwatch_t:s0 tcontext=system_u:object_r:var_t:s0

tclass=dir

Was caused by:

Missing type enforcement (TE) allow rule.

You can use audit2allow to generate a loadable module to allow this access.
```

The -a command-line option causes all audit logs to be read. The -w option produces the human-readable description. As shown, access was denied due to a missing Type Enforcement rule.

 $^{^{9}}$ See the audit2allow(1) manual page for more information about audit2allow.

3. Run the following command to view the Type Enforcement rule that allows the denied access:



Important

Missing Type Enforcement rules are usually caused by bugs in the SELinux policy, and should be reported in *Red Hat Bugzilla*¹⁰. For Fedora, create bugs against the **Fedora** product, and select the **selinux-policy** component. Include the output of the **audit2allow -w -a** and **audit2allow -a** commands in such bug reports.

4. To use the rule displayed by **audit2allow -a**, run the following command as root to create a custom module. The **-M** option creates a Type Enforcement file (**.te**) with the name specified with **-M**, in your current working directory:

5. Also, audit2allow compiles the Type Enforcement rule into a policy package (.pp):

```
~]# 1s
mycertwatch.pp mycertwatch.te
```

To install the module, run the following command as the root:

```
~]# semodule -i mycertwatch.pp
```

¹⁰ https://bugzilla.redhat.com/



Important

Modules created with audit2allow may allow more access than required. It is recommended that policy created with audit2allow be posted to the upstream SELinux list for review. If you believe there is a bug in the policy, please create a bug in Red Hat Bugzilla¹¹.

If you have multiple denial messages from multiple processes, but only want to create a custom policy for a single process, use the grep utility to narrow down the input for audit2allow. The following example demonstrates using grep to only send denial messages related to certwatch through audit2allow:

~]# grep certwatch /var/log/audit/audit.log | audit2allow -R -M mycertwatch2 To make this policy package active, execute: semodule -i mycertwatch2.pp

¹¹ https://bugzilla.redhat.com/

Further Information

11.1. Contributors

- Dominick Grift¹ Technical Editor
- Murray McAllister² Product Security
- James Morris³ Technical Editor
- Eric Paris⁴ Technical Editor
- Scott Radvan⁵ Red Hat Customer Content Services
- Daniel Walsh⁶ Red Hat Security Engineering

11.2. Other Resources

The National Security Agency (NSA)

NSA was the original developer of SELinux. Researchers in NSA's National Information Assurance Research Laboratory (NIARL) designed and implemented flexible mandatory access controls in the major subsystems of the Linux kernel and implemented the new operating system components provided by the Flask architecture, namely the security server and the access vector cache.⁷

- Main SELinux website: http://www.nsa.gov/research/selinux/index.shtml.
- SELinux documentation: http://www.nsa.gov/research/selinux/docs.shtml.
- SELinux background: http://www.nsa.gov/research/selinux/background.shtml.

Tresys Technology

*Tresys Technology*⁸ are the upstream for:

- SELinux userland libraries and tools⁹.
- SELinux Reference Policy¹⁰.

SELinux News

- News: http://selinuxnews.org/wp/.
- Planet SELinux (blogs): http://selinuxnews.org/planet/.

¹ mailto:domg472@gmail.com

² mailto:mmcallis@redhat.com

³ mailto:jmorris@redhat.com

⁴ mailto:eparis@parisplace.org

mailto:sradvan@redhat.com

⁶ mailto:dwalsh@redhat.com

⁷ See the NSA *Contributors to SELinux* [http://www.nsa.gov/research/selinux/contrib.shtml] page for more information.

⁸ http://www.tresys.com/

http://userspace.selinuxproject.org/trac/

http://oss.tresys.com/projects/refpolicy

SELinux Project Wiki

- Main page: http://selinuxproject.org/page/Main_Page.
- User resources, including links to documentation, mailing lists, websites, and tools: http://selinuxproject.org/page/User_Resources.

Fedora

- Main page: http://fedoraproject.org/wiki/SELinux.
- Troubleshooting: http://fedoraproject.org/wiki/SELinux/Troubleshooting.
- Fedora SELinux FAQ: http://fedoraproject.org/wiki/SELinux_FAQ.

The UnOfficial SELinux FAQ

http://www.crypt.gen.nz/selinux/faq.html

The SELinux Notebook - The Foundations - 3rd Edition

http://www.freetechbooks.com/the-selinux-notebook-the-foundations-t785.html

IRC

On Freenode¹¹:

- #selinux
- · #fedora-selinux
- #security

96

¹¹ http://freenode.net/





Introduction

This part of the book focuses more on practical tasks and provides information how to set up and configure various services. For each service, there are listed the most common types and Booleans with the specifications. Also included are real-world examples of configuring those services and demonstrations of how SELinux complements their operation.

When SELinux is in enforcing mode, the default policy used in Fedora, is the targeted policy. Processes that are targeted run in a confined domain, and processes that are not targeted run in an unconfined domain. See *Chapter 3, Targeted Policy* for more information about targeted policy and confined and unconfined processes.

The Apache HTTP Server

The Apache HTTP Server provides an open-source HTTP server with the current HTTP standards.¹

In Fedora, the *httpd* package provides the Apache HTTP Server. Run the following command to see if the *httpd* package is installed:

```
~]$ rpm -q httpd
package httpd is not installed
```

If it is not installed and you want to use the Apache HTTP Server, use the DNF utility as the root user to install it:

~]# dnf install httpd

13.1. The Apache HTTP Server and SELinux

When SELinux is enabled, the Apache HTTP Server (httpd) runs confined by default. Confined processes run in their own domains, and are separated from other confined processes. If a confined process is compromised by an attacker, depending on SELinux policy configuration, an attacker's access to resources and the possible damage they can do is limited. The following example demonstrates the httpd processes running in their own domain. This example assumes the httpd, setroubleshoot, setroubleshoot-server and policycoreutils-python packages are installed:

Run the getenforce command to confirm SELinux is running in enforcing mode:

```
~]$ getenforce
Enforcing
```

The command returns **Enforcing** when SELinux is running in enforcing mode.

2. Run the following command as root to start httpd:

```
~]# systemctl start httpd.service
```

Confirm that the service is running. The output should include the information below (only the time stamp will differ):

```
~]# systemctl status httpd.service
httpd.service - The Apache HTTP Server
Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled)
Active: active (running) since Mon 2013-08-05 14:00:55 CEST; 8s ago
```

3. To view the httpd processes, execute the following command:

¹ Refer to the *Apache HTTP Server Project* [http://httpd.apache.org/] page for more information.

```
~]$ ps -eZ | grep httpd
system_u:system_r:httpd_t:s0
                              19780 ?
                                             00:00:00 httpd
system_u:system_r:httpd_t:s0 19781 ?
                                             00:00:00 httpd
system_u:system_r:httpd_t:s0 19782 ?
                                             00:00:00 httpd
system_u:system_r:httpd_t:s0 19783 ?
                                             00:00:00 httpd
system_u:system_r:httpd_t:s0
                              19784 ?
                                             00:00:00 httpd
system_u:system_r:httpd_t:s0
                              19785 ?
                                             00:00:00 httpd
```

The SELinux context associated with the httpd processes is

system_u:system_r:httpd_t:s0. The second last part of the context, httpd_t, is the type. A type defines a domain for processes and a type for files. In this case, the httpd processes are running in the httpd_t domain.

SELinux policy defines how processes running in confined domains (such as httpd_t) interact with files, other processes, and the system in general. Files must be labeled correctly to allow httpd access to them. For example, httpd can read files labeled with the httpd_sys_content_t type, but cannot write to them, even if Linux (DAC) permissions allow write access. Booleans must be enabled to allow certain behavior, such as allowing scripts network access, allowing httpd access to NFS and CIFS volumes, and httpd being allowed to execute Common Gateway Interface (CGI) scripts.

When the /etc/httpd/conf/httpd.conf file is configured so httpd listens on a port other than TCP ports 80, 443, 488, 8008, 8009, or 8443, the semanage port command must be used to add the new port number to SELinux policy configuration. The following example demonstrates configuring httpd to listen on a port that is not already defined in SELinux policy configuration for httpd, and, as a consequence, httpd failing to start. This example also demonstrates how to then configure the SELinux system to allow httpd to successfully listen on a non-standard port that is not already defined in the policy. This example assumes the httpd package is installed. Run each command in the example as the root user:

1. Run the following command to confirm httpd is not running:

```
~]# systemctl status httpd.service
httpd.service - The Apache HTTP Server
Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled)
Active: inactive (dead)
```

If the output differs, stop the process:

```
~]# systemctl stop httpd.service
```

2. Use the semanage utility to view the ports SELinux allows httpd to listen on:

```
~]# semanage port -1 | grep -w http_port_t
http_port_t tcp 80, 443, 488, 8008, 8009, 8443
```

3. Edit the /etc/httpd/conf/httpd.conf file as root. Configure the Listen option so it lists a port that is not configured in SELinux policy configuration for httpd. In this example, httpd is configured to listen on port 12345:

```
# Change this to Listen on specific IP addresses as shown below to
# prevent Apache from glomming onto all bound IP addresses (0.0.0.0)
#
```

```
#Listen 12.34.56.78:80
Listen 127.0.0.1:12345
```

4. Run the following command to start httpd:

```
~]# systemctl start httpd.service
Job for httpd.service failed. See 'systemctl status httpd.service' and 'journalctl -xn' for details.
```

An SELinux denial message similar to the following is logged:

```
setroubleshoot: SELinux is preventing the httpd (httpd_t) from binding to port 12345. For complete SELinux messages. run sealert -l f18bca99-db64-4c16-9719-1db89f0d8c77
```

5. For SELinux to allow httpd to listen on port 12345, as used in this example, the following command is required:

```
~]# semanage port -a -t http_port_t -p tcp 12345
```

6. Start httpd again and have it listen on the new port:

```
~]# systemctl start httpd.service
```

- 7. Now that SELinux has been configured to allow httpd to listen on a non-standard port (TCP 12345 in this example), httpd starts successfully on this port.
- 8. To prove that httpd is listening and communicating on TCP port 12345, open a telnet connection to the specified port and issue a HTTP GET command, as follows:

```
~]# telnet localhost 12345
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
GET / HTTP/1.0

HTTP/1.1 200 OK
Date: Wed, 02 Dec 2009 14:36:34 GMT
Server: Apache/2.2.13 (Red Hat)
Accept-Ranges: bytes
Content-Length: 3985
Content-Type: text/html; charset=UTF-8
[...continues...]
```

13.2. Types

Type Enforcement is the main permission control used in SELinux targeted policy. All files and processes are labeled with a type: types define a domain for processes and a type for files. SELinux policy rules define how types access each other, whether it be a domain accessing a type, or a domain accessing another domain. Access is only allowed if a specific SELinux policy rule exists that allows it.

The following example creates a new file in the /var/www/html/ directory, and shows the file inheriting the httpd_sys_content_t type from its parent directory (/var/www/html/):

1. Run the following command to view the SELinux context of /var/www/html/:

```
~]$ ls -dZ /var/www/html drwxr-xr-x root root system_u:object_r:httpd_sys_content_t:s0 /var/www/html
```

This shows /var/www/html/ is labeled with the httpd_sys_content_t type.

2. Create a new file by using the touch utility as root:

```
~]# touch /var/www/html/file1
```

3. Run the following command to view the SELinux context:

```
~]$ ls -Z /var/www/html/file1
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 /var/www/html/file1
```

The **1s** -**Z** command shows **file1** labeled with the httpd_sys_content_t type. SELinux allows httpd to read files labeled with this type, but not write to them, even if Linux permissions allow write access. SELinux policy defines what types a process running in the httpd_t domain (where httpd runs) can read and write to. This helps prevent processes from accessing files intended for use by another process.

For example, httpd can access files labeled with the httpd_sys_content_t type (intended for the Apache HTTP Server), but by default, cannot access files labeled with the samba_share_t type (intended for Samba). Also, files in user home directories are labeled with the user_home_t type: by default, this prevents httpd from reading or writing to files in user home directories.

The following lists some of the types used with httpd. Different types allow you to configure flexible access:

```
httpd_sys_content_t
```

Use this type for static web content, such as .html files used by a static website. Files labeled with this type are accessible (read only) to httpd and scripts executed by httpd. By default, files and directories labeled with this type cannot be written to or modified by httpd or other processes. Note that by default, files created in or copied into the /var/www/html/ directory are labeled with the httpd_sys_content_t type.

```
httpd_sys_script_exec_t
```

Use this type for scripts you want httpd to execute. This type is commonly used for Common Gateway Interface (CGI) scripts in the /var/www/cgi-bin/ directory. By default, SELinux policy prevents httpd from executing CGI scripts. To allow this, label the scripts with the httpd_sys_script_exec_t type and enable the httpd_enable_cgi Boolean. Scripts labeled with httpd_sys_script_exec_t run in the httpd_sys_script_t domain when executed by httpd. The httpd_sys_script_t domain has access to other system domains, such as postgresql_t and mysqld_t.

```
httpd_sys_rw_content_t
```

Files labeled with this type can be written to by scripts labeled with the httpd_sys_script_exec_t type, but cannot be modified by scripts labeled with any other type. You must use the httpd_sys_rw_content_t type to label files that will be read from and written to by scripts labeled with the httpd_sys_script_exec_t type.

httpd_sys_ra_content_t

Files labeled with this type can be appended to by scripts labeled with the httpd_sys_script_exec_t type, but cannot be modified by scripts labeled with any other type. You must use the httpd_sys_ra_content_t type to label files that will be read from and appended to by scripts labeled with the httpd_sys_script_exec_t type.

httpd_unconfined_script_exec_t

Scripts labeled with this type run without SELinux protection. Only use this type for complex scripts, after exhausting all other options. It is better to use this type instead of disabling SELinux protection for httpd, or for the entire system.



Note

To see more of the available types for httpd, run the following command:

~]\$ grep httpd /etc/selinux/targeted/contexts/files/file_contexts

Procedure 13.1. Changing the SELinux Context

The type for files and directories can be changed with the **chcon** command. Changes made with **chcon** do not survive a file system relabel or the **restorecon** command. SELinux policy controls whether users are able to modify the SELinux context for any given file. The following example demonstrates creating a new directory and an **index.html** file for use by httpd, and labeling that file and directory to allow httpd access to them:

 Use the mkdir utility as root to create a top-level directory structure to store files to be used by httpd:

```
~]# mkdir -p /my/website
```

2. Files and directories that do not match a pattern in file-context configuration may be labeled with the default_t type. This type is inaccessible to confined services:

```
~]$ ls -dZ /my
drwxr-xr-x root root unconfined_u:object_r:default_t:s0 /my
```

3. Run the following command as root to change the type of the /my/ directory and subdirectories, to a type accessible to httpd. Now, files created under /my/website/ inherit the httpd_sys_content_t type, rather than the default_t type, and are therefore accessible to httpd:

```
~]# chcon -R -t httpd_sys_content_t /my/
~]# touch /my/website/index.html
~]# ls -Z /my/website/index.html
-rw-r--r-- root root unconfined_u:object_r:httpd_sys_content_t:s0 /my/website/
index.html
```

See Section 4.7.1, "Temporary Changes: chcon" for further information about chcon.

Use the **semanage fcontext** command (**semanage** is provided by the *policycoreutils-python* package) to make label changes that survive a relabel and the **restorecon** command. This command adds changes to file-context configuration. Then, run **restorecon**, which reads file-context configuration, to apply the label change. The following example demonstrates creating a new directory and an **index.html** file for use by httpd, and persistently changing the label of that directory and file to allow httpd access to them:

 Use the mkdir utility as root to create a top-level directory structure to store files to be used by httpd:

```
~]# mkdir -p /my/website
```

2. Run the following command as root to add the label change to file-context configuration:

```
~]# semanage fcontext -a -t httpd_sys_content_t "/my(/.*)?"
```

The "/my(/.*)?" expression means the label change applies to the /my/ directory and all files and directories under it.

3. Use the touch utility as root to create a new file:

```
~]# touch /my/website/index.html
```

4. Run the following command as root to apply the label changes (**restorecon** reads file-context configuration, which was modified by the **semanage** command in step 2):

```
-]# restorecon -R -v /my/
restorecon reset /my context unconfined_u:object_r:default_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
restorecon reset /my/website context unconfined_u:object_r:default_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
restorecon reset /my/website/index.html context unconfined_u:object_r:default_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
```

See Section 4.7.2, "Persistent Changes: semanage fcontext" for further information on semanage.

13.3. Booleans

SELinux is based on the least level of access required for a service to run. Services can be run in a variety of ways; therefore, you need to specify how you run your services. This can be achieved using Booleans that allow parts of SELinux policy to be changed at runtime, without any knowledge of SELinux policy writing. This allows changes, such as allowing services access to NFS volumes, without reloading or recompiling SELinux policy.

To modify the state of a Boolean, use the **setsebool** command. For example, to enable the httpd_anon_write Boolean, run the following command as the root user:

```
~]# setsebool -P httpd_anon_write on
```

To disable a Boolean, using the same example, simply change **on** to **off** in the command, as shown below:

~]# setsebool -P httpd_anon_write off



Note

Do not use the -P option if you do not want **setsebool** changes to persist across reboots.

Below is a description of common Booleans available that cater for the way httpd is running:

httpd_anon_write

When disabled, this Boolean allows httpd to only have read access to files labeled with the public_content_rw_t type. Enabling this Boolean allows httpd to write to files labeled with the public_content_rw_t type, such as a public directory containing files for a public file transfer service.

httpd_mod_auth_ntlm_winbind

Enabling this Boolean allows access to NTLM and Winbind authentication mechanisms via the mod_auth_ntlm_winbind module in httpd.

httpd_mod_auth_pam

Enabling this Boolean allows access to PAM authentication mechanisms via the mod_auth_pam module in httpd.

httpd_sys_script_anon_write

This Boolean defines whether or not HTTP scripts are allowed write access to files labeled with the public_content_rw_t type, as used in a public file transfer service.

httpd_builtin_scripting

This Boolean defines access to httpd scripting. Having this Boolean enabled is often required for PHP content.

httpd_can_network_connect

When disabled, this Boolean prevents HTTP scripts and modules from initiating a connection to a network or remote port. Enable this Boolean to allow this access.

httpd_can_network_connect_db

When disabled, this Boolean prevents HTTP scripts and modules from initiating a connection to database servers. Enable this Boolean to allow this access.

httpd_can_network_relay

Enable this Boolean when httpd is being used as a forward or reverse proxy.

httpd_can_sendmail

When disabled, this Boolean prevents HTTP modules from sending mail. This can prevent spam attacks should a vulnerability be found in httpd. Enable this Boolean to allow HTTP modules to send mail.

httpd_dbus_avahi

When disabled, this Boolean denies httpd access to the avahi service via D-Bus. Enable this Boolean to allow this access.

httpd_enable_cgi

When disabled, this Boolean prevents httpd from executing CGI scripts. Enable this Boolean to allow httpd to execute CGI scripts (CGI scripts must be labeled with the httpd_sys_script_exec_t type).

httpd_enable_ftp_server

Enabling this Boolean allows httpd to listen on the FTP port and act as an FTP server.

httpd_enable_homedirs

When disabled, this Boolean prevents httpd from accessing user home directories. Enable this Boolean to allow httpd access to user home directories; for example, content in /home/*/.

httpd_execmem

When enabled, this Boolean allows httpd to execute programs that require memory addresses that are both executable and writeable. Enabling this Boolean is not recommended from a security standpoint as it reduces protection against buffer overflows, however certain modules and applications (such as Java and Mono applications) require this privilege.

httpd_ssi_exec

This Boolean defines whether or not server side include (SSI) elements in a web page can be executed.

httpd_tty_comm

This Boolean defines whether or not httpd is allowed access to the controlling terminal. Usually this access is not required, however in cases such as configuring an SSL certificate file, terminal access is required to display and process a password prompt.

httpd_unified

When enabled, this Boolean allows httpd_t complete access to all of the httpd types (that is to execute, read, or write sys_content_t). When disabled, there is separation in place between web content that is read-only, writeable, or executable. Disabling this Boolean ensures an extra level of security but adds the administrative overhead of having to individually label scripts and other web content based on the file access that each should have.

httpd_use_cifs

Enable this Boolean to allow httpd access to files on CIFS volumes that are labeled with the cifs_t type, such as file systems mounted via Samba.

httpd_use_nfs

Enable this Boolean to allow httpd access to files on NFS volumes that are labeled with the nfs_t type, such as file systems mounted using NFS.



Note

Due to the continuous development of the SELinux policy, the list above might not contain all Booleans related to the service at all times. To list them, run the following command:

```
~]$ getsebool -a | grep service_name
```

Run the following command to view description of a particular Boolean:

```
~]$ sepolicy booleans -b boolean_name
```

Note that the additional *policycoreutils-devel* package providing the sepolicy utility is required.

13.4. Configuration examples

The following examples provide real-world demonstrations of how SELinux complements the Apache HTTP Server and how full function of the Apache HTTP Server can be maintained.

13.4.1. Running a static site

To create a static website, label the .html files for that website with the httpd_sys_content_t type. By default, the Apache HTTP Server cannot write to files that are labeled with the httpd_sys_content_t type. The following example creates a new directory to store files for a read-only website:

Use the mkdir utility as root to create a top-level directory:

```
~]# mkdir /mywebsite
```

 As root, create a /mywebsite/index.html file. Copy and paste the following content into / mywebsite/index.html:

```
<html>
<h2>index.html from /mywebsite/</h2>
</html>
```

3. To allow the Apache HTTP Server read only access to /mywebsite/, as well as files and subdirectories under it, label the directory with the httpd_sys_content_t type. Run the following command as root to add the label change to file-context configuration:

```
~]# semanage fcontext -a -t httpd_sys_content_t "/mywebsite(/.*)?"
```

4. Use the restorecon utility as root to make the label changes:

```
~]# restorecon -R -v /mywebsite
restorecon reset /mywebsite context unconfined_u:object_r:default_t:s0-
>system_u:object_r:httpd_sys_content_t:s0
```

```
restorecon reset /mywebsite/index.html context unconfined_u:object_r:default_t:s0->system_u:object_r:httpd_sys_content_t:s0
```

5. For this example, edit the /etc/httpd/conf/httpd.conf file as root. Comment out the existing DocumentRoot option. Add a DocumentRoot "/mywebsite" option. After editing, these options should look as follows:

```
#DocumentRoot "/var/www/html"
DocumentRoot "/mywebsite"
```

6. Run the following command as root to see the status of the Apache HTTP Server. If the server is stopped, start it:

```
~]# systemctl status httpd.service
httpd.service - The Apache HTTP Server
Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled)
Active: inactive (dead)
```

```
~]# systemctl start httpd.service
```

If the server is running, restart the service by executing the following command as root (this also applies any changes made to **httpd.conf**):

```
~]# systemctl status httpd.service
httpd.service - The Apache HTTP Server
Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled)
Active: active (running) since Wed 2014-02-05 13:16:46 CET; 2s ago
```

```
~]# systemctl restart httpd.service
```

7. Use a web browser to navigate to http://localhost/index.html. The following is displayed:

```
index.html from /mywebsite/
```

13.4.2. Sharing NFS and CIFS volumes

By default, NFS mounts on the client side are labeled with a default context defined by policy for NFS volumes. In common policies, this default context uses the nfs_t type. Also, by default, Samba shares mounted on the client side are labeled with a default context defined by policy. In common policies, this default context uses the cifs_t type.

Depending on policy configuration, services may not be able to read files labeled with the nfs_t or cifs_t types. This may prevent file systems labeled with these types from being mounted and then read or exported by other services. Booleans can be enabled or disabled to control which services are allowed to access the nfs_t and cifs_t types.

Enable the httpd_use_nfs Boolean to allow httpd to access and share NFS volumes (labeled with the nfs_t type):

```
~]# setsebool -P httpd_use_nfs on
```

Enable the httpd_use_cifs Boolean to allow httpd to access and share CIFS volumes (labeled with the cifs_t type):

```
~]# setsebool -P httpd_use_cifs on
```



Note

Do not use the **-P** option if you do not want **setsebool** changes to persist across reboots.

13.4.3. Sharing files between services

Type Enforcement helps prevent processes from accessing files intended for use by another process. For example, by default, Samba cannot read files labeled with the httpd_sys_content_t type, which are intended for use by the Apache HTTP Server. Files can be shared between the Apache HTTP Server, FTP, rsync, and Samba, if the desired files are labeled with the public_content_t or public_content_rw_t type.

The following example creates a directory and files, and allows that directory and files to be shared (read only) through the Apache HTTP Server, FTP, rsync, and Samba:

1. Use the mkdir utility as root to create a new top-level directory to share files between multiple services:

```
~]# mkdir /shares
```

2. Files and directories that do not match a pattern in file-context configuration may be labeled with the default_t type. This type is inaccessible to confined services:

```
~]$ ls -dZ /shares
drwxr-xr-x root root unconfined_u:object_r:default_t:s0 /shares
```

 As root, create a /shares/index.html file. Copy and paste the following content into / shares/index.html:

```
<html>
<body>
Hello
</body>
</html>
```

4. Labeling /shares/ with the public_content_t type allows read-only access by the Apache HTTP Server, FTP, rsync, and Samba. Run the following command as root to add the label change to file-context configuration:

```
~]# semanage fcontext -a -t public_content_t "/shares(/.*)?"
```

5. Use the restorecon utility as root to apply the label changes:

```
~]# restorecon -R -v /shares/
restorecon reset /shares context unconfined_u:object_r:default_t:s0-
>system_u:object_r:public_content_t:s0
restorecon reset /shares/index.html context unconfined_u:object_r:default_t:s0-
>system_u:object_r:public_content_t:s0
```

To share /shares/ through Samba:

1. Confirm the *samba*, *samba-common*, and *samba-client* packages are installed (version numbers may differ):

```
~]$ rpm -q samba samba-common samba-client
samba-3.4.0-0.41.el6.3.i686
samba-common-3.4.0-0.41.el6.3.i686
samba-client-3.4.0-0.41.el6.3.i686
```

If any of these packages are not installed, install them by running the following command as root:

```
~]# dnf install package-name
```

2. Edit the /etc/samba/smb.conf file as root. Add the following entry to the bottom of this file to share the /shares/ directory through Samba:

```
[shares]
comment = Documents for Apache HTTP Server, FTP, rsync, and Samba
path = /shares
public = yes
writeable = no
```

A Samba account is required to mount a Samba file system. Run the following command as root to create a Samba account, where *username* is an existing Linux user. For example, smbpasswd -a testuser creates a Samba account for the Linux testuser user:

```
~]# smbpasswd -a testuser

New SMB password: Enter a password

Retype new SMB password: Enter the same password again

Added user testuser.
```

If you run the above command, specifying a user name of an account that does not exist on the system, it causes a **Cannot locate Unix account for 'username'!** error.

4. Start the Samba service:

```
~]# systemctl start smb.service
```

5. Run the following command to list the available shares, where *username* is the Samba account added in step 3. When prompted for a password, enter the password assigned to the Samba account in step 3 (version numbers may differ):

```
~]$ smbclient -U username -L localhost
Enter username's password:
Domain=[HOSTNAME] OS=[Unix] Server=[Samba 3.4.0-0.41.el6]
Sharename
               Type
                         Comment
                ----
               Disk
                         Documents for Apache HTTP Server, FTP, rsync, and Samba
shares
IPC$
               IPC
                         IPC Service (Samba Server Version 3.4.0-0.41.el6)
username
               Disk
                         Home Directories
Domain=[HOSTNAME] OS=[Unix] Server=[Samba 3.4.0-0.41.el6]
                     Comment
Server
Workgroup
                     Master
```

6. User the mkdir utility to create a new directory. This directory will be used to mount the **shares** Samba share:

```
~]# mkdir /test/
```

7. Run the following command as root to mount the **shares** Samba share to **/test/**, replacing *username* with the user name from step 3:

```
~]# mount //localhost/shares /test/ -o user=username
```

Enter the password for *username*, which was configured in step 3.

8. View the content of the file, which is being shared through Samba:

```
~]$ cat /test/index.html
<html>
<body>
Hello
</body>
</html>
```

To share /shares/ through the Apache HTTP Server:

1. Confirm the *httpd* package is installed (version number may differ):

```
~]$ rpm -q httpd
httpd-2.2.11-6.i386
```

If this package is not installed, use the DNF utility as root to install it:

```
~]# dnf install httpd
```

2. Change into the /var/www/html/ directory. Run the following command as root to create a link (named shares) to the /shares/ directory:

```
html]# ln -s /shares/ shares
```

3. Start the Apache HTTP Server:

```
~]# systemctl start httpd.service
```

4. Use a web browser to navigate to http://localhost/shares. The /shares/index.html file is displayed.

By default, the Apache HTTP Server reads an **index.html** file if it exists. If **/shares/** did not have **index.html**, and instead had **file1**, **file2**, and **file3**, a directory listing would occur when accessing **http://localhost/shares**:

1. Remove the index.html file:

```
~]# rm -i /shares/index.html
```

2. Use the touch utility as root to create three files in /shares/:

```
~]# touch /shares/file{1,2,3}

~]# ls -Z /shares/
-rw-r--r-- root root system_u:object_r:public_content_t:s0 file1
-rw-r--r-- root root unconfined_u:object_r:public_content_t:s0 file2
-rw-r--r-- root root unconfined_u:object_r:public_content_t:s0 file3
```

3. Run the following command as root to see the status of the Apache HTTP Server:

```
~]# systemctl status httpd.service
httpd.service - The Apache HTTP Server
Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled)
Active: inactive (dead)
```

If the server is stopped, start it:

```
~]# systemctl start httpd.service
```

4. Use a web browser to navigate to http://localhost/shares. A directory listing is displayed:

Index of /shares

	<u>Name</u>	Last modified	Size Description
Pare	nt Directory	Į.	-
file1		25-Feb-2009 10:11	0
file2		25-Feb-2009 10:11	0
file3		25-Feb-2009 10:11	0

13.4.4. Changing port numbers

Depending on policy configuration, services may only be allowed to run on certain port numbers. Attempting to change the port a service runs on without changing policy may result in the service failing to start. Use the semanage utility as the root user to list the ports SELinux allows httpd to listen on:

```
~]# semanage port -1 | grep -w http_port_t
http_port_t tcp 80, 443, 488, 8008, 8009, 8443
```

By default, SELinux allows http to listen on TCP ports 80, 443, 488, 8008, 8009, or 8443. If / etc/httpd/conf/httpd.conf is configured so that httpd listens on any port not listed for http_port_t, httpd fails to start.

To configure httpd to run on a port other than TCP ports 80, 443, 488, 8008, 8009, or 8443:

 Edit the /etc/httpd/conf/httpd.conf file as root so the Listen option lists a port that is not configured in SELinux policy for httpd. The following example configures httpd to listen on the 10.0.0.1 IP address, and on TCP port 12345:

```
# Change this to Listen on specific IP addresses as shown below to
# prevent Apache from glomming onto all bound IP addresses (0.0.0.0)
#
#Listen 12.34.56.78:80
Listen 10.0.0.1:12345
```

2. Run the following command as the root user to add the port to SELinux policy configuration:

```
~]# semanage port -a -t http_port_t -p tcp 12345
```

3. Confirm that the port is added:

```
~]# semanage port -1 | grep -w http_port_t
http_port_t tcp 12345, 80, 443, 488, 8008, 8009, 8443
```

If you no longer run httpd on port 12345, use the semanage utility as root to remove the port from policy configuration:

~]# semanage port -d -t http_port_t -p tcp 12345

Samba

Samba is an open-source implementation of the Server Message Block (SMB) and Common Internet File System (CIFS) protocols that provides file and print services between clients across various operating systems.¹

In Fedora, the *samba* package provides the Samba server. Run the following command to see if the *samba* package is installed:

```
~]$ rpm -q samba
package samba is not installed
```

If it is not installed and you want to use Samba, use the DNF utility as the root user to install it:

```
~]# dnf install samba
```

14.1. Samba and SELinux

When SELinux is enabled, the Samba server (smbd) runs confined by default. Confined services run in their own domains, and are separated from other confined services. The following example demonstrates the smbd process running in its own domain. This example assumes the *samba* package is installed:

1. Run the **getenforce** command to confirm SELinux is running in enforcing mode:

```
~]$ getenforce
Enforcing
```

The command returns **Enforcing** when SELinux is running in enforcing mode.

2. Run the following command as root to start smbd:

```
~]# systemctl start smb.service
```

Confirm that the service is running. The output should include the information below (only the time stamp will differ):

```
~]# systemctl status smb.service
smb.service - Samba SMB Daemon
Loaded: loaded (/usr/lib/systemd/system/smb.service; disabled)
Active: active (running) since Mon 2013-08-05 12:17:26 CEST; 2h 22min ago
```

3. To view the smbd processes, execute the following command:

```
~]$ ps -eZ | grep smb
```

¹ Refer to the official *Samba* [http://samba.org] website for more information.

```
      system_u:system_r:smbd_t:s0
      9653 ?
      00:00:00 smbd

      system_u:system_r:smbd_t:s0
      9654?
      00:00:00 smbd
```

The SELinux context associated with the smbd processes is **system_u:system_r:smbd_t:s0**. The second last part of the context, smbd_t, is the type. A type defines a domain for processes and a type for files. In this case, the smbd processes are running in the smbd_t domain.

Files must be labeled correctly to allow smbd to access and share them. For example, smbd can read and write to files labeled with the samba_share_t type, but by default, cannot access files labeled with the httpd_sys_content_t type, which is intended for use by the Apache HTTP Server. Booleans must be enabled to allow certain behavior, such as allowing home directories and NFS volumes to be exported through Samba, as well as to allow Samba to act as a domain controller.

14.2. Types

Label files with the samba_share_t type to allow Samba to share them. Only label files you have created, and do not relabel system files with the samba_share_t type: Booleans can be enabled to share such files and directories. SELinux allows Samba to write to files labeled with the samba_share_t type, as long as the /etc/samba/smb.conf file and Linux permissions are set accordingly.

The samba_etc_t type is used on certain files in the /etc/samba/ directory, such as smb.conf. Do not manually label files with the samba_etc_t type. If files in this directory are not labeled correctly, run the restorecon -R -v /etc/samba command as the root user to restore such files to their default contexts. If /etc/samba/smb.conf is not labeled with the samba_etc_t type, starting the Samba service may fail and an SELinux denial message may be logged. The following is an example denial message when /etc/samba/smb.conf was labeled with the httpd_sys_content_t type:

```
setroubleshoot: SELinux is preventing smbd (smbd_t) "read" to ./smb.conf (httpd_sys_content_t). For complete SELinux messages. run sealert -l deb33473-1069-482b-bb50-e4cd05ab18af
```

14.3. Booleans

SELinux is based on the least level of access required for a service to run. Services can be run in a variety of ways; therefore, you need to specify how you run your services. Use the following Booleans to set up SELinux:

smbd_anon_write

Having this Boolean enabled allows smbd to write to a public directory, such as an area reserved for common files that otherwise has no special access restrictions.

samba_create_home_dirs

Having this Boolean enabled allows Samba to create new home directories independently. This is often done by mechanisms such as PAM.

samba_domain_controller

When enabled, this Boolean allows Samba to act as a domain controller, as well as giving it permission to execute related commands such as **useradd**, **groupadd**, and **passwd**.

samba_enable_home_dirs

Enabling this Boolean allows Samba to share users' home directories.

samba_export_all_ro

Export any file or directory, allowing read-only permissions. This allows files and directories that are not labeled with the samba_share_t type to be shared through Samba. When the samba_export_all_ro Boolean is enabled, but the samba_export_all_rw Boolean is disabled, write access to Samba shares is denied, even if write access is configured in /etc/samba/smb.conf, as well as Linux permissions allowing write access.

samba_export_all_rw

Export any file or directory, allowing read and write permissions. This allows files and directories that are not labeled with the samba_share_t type to be exported through Samba. Permissions in /etc/samba/smb.conf and Linux permissions must be configured to allow write access.

samba_run_unconfined

Having this Boolean enabled allows Samba to run unconfined scripts in the **/var/lib/samba/scripts/** directory.

samba_share_fusefs

This Boolean must be enabled for Samba to share fusefs file systems.

samba_share_nfs

Disabling this Boolean prevents smbd from having full access to NFS shares through Samba. Enabling this Boolean will allow Samba to share NFS volumes.

use_samba_home_dirs

Enable this Boolean to use a remote server for Samba home directories.

virt_use_samba

Allow virtual machine access to CIFS files.



Note

Due to the continuous development of the SELinux policy, the list above might not contain all Booleans related to the service at all times. To list them, run the following command:

```
~]$ getsebool -a | grep service_name
```

Run the following command to view description of a particular Boolean:

```
~]$ sepolicy booleans -b boolean_name
```

Note that the additional policycoreutils-devel package providing the sepolicy utility is required.

14.4. Configuration examples

The following examples provide real-world demonstrations of how SELinux complements the Samba server and how full function of the Samba server can be maintained.

14.4.1. Sharing directories you create

The following example creates a new directory, and shares that directory through Samba:

1. Confirm that the samba, samba-common, and samba-client packages are installed:

```
~]$ rpm -q samba samba-common samba-client
package samba is not installed
package samba-common is not installed
package samba-client is not installed
```

If any of these packages are not installed, install them by using the DNF utility as root:

```
~]# dnf install package-name
```

2. Use the mkdir utility as root to create a new top-level directory to share files through Samba:

```
~]# mkdir /myshare
```

3. Use the touch utility root to create an empty file. This file is used later to verify the Samba share mounted correctly:

```
~]# touch /myshare/file1
```

4. SELinux allows Samba to read and write to files labeled with the samba_share_t type, as long as the /etc/samba/smb.conf file and Linux permissions are set accordingly. Run the following command as root to add the label change to file-context configuration:

```
~]# semanage fcontext -a -t samba_share_t "/myshare(/.*)?"
```

5. Use the restorecon utility as root to apply the label changes:

```
~]# restorecon -R -v /myshare
restorecon reset /myshare context unconfined_u:object_r:default_t:s0-
>system_u:object_r:samba_share_t:s0
restorecon reset /myshare/file1 context unconfined_u:object_r:default_t:s0-
>system_u:object_r:samba_share_t:s0
```

6. Edit /etc/samba/smb.conf as root. Add the following to the bottom of this file to share the / myshare/ directory through Samba:

```
[myshare]
comment = My share
path = /myshare
public = yes
writeable = no
```

7. A Samba account is required to mount a Samba file system. Run the following command as root to create a Samba account, where *username* is an existing Linux user. For example, **smbpasswd** -a testuser creates a Samba account for the Linux testuser user:

```
~]# smbpasswd -a testuser
New SMB password: Enter a password
```

```
Retype new SMB password: Enter the same password again
Added user testuser.
```

If you run the above command, specifying a user name of an account that does not exist on the system, it causes a **Cannot locate Unix account for 'username'!** error.

8. Start the Samba service:

```
~]# systemctl start smb.service
```

9. Run the following command to list the available shares, where *username* is the Samba account added in step 7. When prompted for a password, enter the password assigned to the Samba account in step 7 (version numbers may differ):

```
~]$ smbclient -U username -L localhost
Enter username's password:
Domain=[HOSTNAME] OS=[Unix] Server=[Samba 3.4.0-0.41.el6]
Sharename
                Туре
                          Comment
               Disk
IPC
Disk
myshare
IPC$
username
                          My share
                          IPC Service (Samba Server Version 3.4.0-0.41.el6)
                          Home Directories
Domain=[HOSTNAME] OS=[Unix] Server=[Samba 3.4.0-0.41.el6]
Server
                     Comment
-----
Workgroup
                     Master
```

10. Use the mkdir utility as root to create a new directory. This directory will be used to mount the **myshare** Samba share:

```
~]# mkdir /test/
```

11. Run the following command as root to mount the **myshare** Samba share to **/test/**, replacing *username* with the user name from step 7:

```
~]# mount //localhost/myshare /test/ -o user=username
```

Enter the password for *username*, which was configured in step 7.

12. Run the following command to view the **file1** file created in step 3:

```
~]$ ls /test/
file1
```

14.4.2. Sharing a website

It may not be possible to label files with the samba_share_t type, for example, when wanting to share a website in the /var/www/html/ directory. For these cases, use the

samba_export_all_ro Boolean to share any file or directory (regardless of the current label), allowing read only permissions, or the samba_export_all_rw Boolean to share any file or directory (regardless of the current label), allowing read and write permissions.

The following example creates a file for a website in /var/www/html/, and then shares that file through Samba, allowing read and write permissions. This example assumes the httpd, samba, samba-common, samba-client, and wget packages are installed:

 As the root user, create a /var/www/html/file1.html file. Copy and paste the following content into this file:

```
<html>
<h2>File being shared through the Apache HTTP Server and Samba.</h2>
</html>
```

2. Run the following command to view the SELinux context of **file1.html**:

```
~]$ ls -Z /var/www/html/file1.html
-rw-r--r--. root root unconfined_u:object_r:httpd_sys_content_t:s0 /var/www/html/
file1.html
```

The file is labeled with the httpd_sys_content_t. By default, the Apache HTTP Server can access this type, but Samba cannot.

3. Start the Apache HTTP Server:

```
~]# systemctl start httpd.service
```

4. Change into a directory your user has write access to, and run the following command. Unless there are changes to the default configuration, this command succeeds:

```
~]$ wget http://localhost/file1.html
Resolving localhost... 127.0.0.1
Connecting to localhost|127.0.0.1|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 84 [text/html]
Saving to: `file1.html.1'

100%[=============] 84 ----K/s in 0s

`file1.html.1' saved [84/84]
```

5. Edit /etc/samba/smb.conf as root. Add the following to the bottom of this file to share the / var/www/html/ directory through Samba:

```
[website]
comment = Sharing a website
path = /var/www/html/
public = no
writeable = no
```

6. The /var/www/html/ directory is labeled with the httpd_sys_content_t type. By default, Samba cannot access files and directories labeled with the this type, even if Linux permissions allow it. To allow Samba access, enable the samba_export_all_ro Boolean:

```
~]# setsebool -P samba_export_all_ro on
```

Do not use the **-P** option if you do not want the change to persist across reboots. Note that enabling the samba_export_all_ro Boolean allows Samba to access any type.

7. Start the Samba service:

~]# systemctl start smb.service

File Transfer Protocol

File Transfer Protocol (FTP) is one of the oldest and most commonly used protocols found on the Internet today. Its purpose is to reliably transfer files between computer hosts on a network without requiring the user to log directly into the remote host or have knowledge of how to use the remote system. It allows users to access files on remote systems using a standard set of simple commands.

The Very Secure FTP Daemon (vsftpd) is designed from the ground up to be fast, stable, and, most importantly, secure. Its ability to handle large numbers of connections efficiently and securely is why vsftpd is the only stand-alone FTP distributed with Fedora.

In Fedora, the *vsftpd* package provides the Very Secure FTP daemon. Run the following command to see if *vsftpd* is installed:

```
~]$ rpm -q vsftpd
package vsftpd is not installed
```

If you want an FTP server and the *vsftpd* package is not installed, use the DNF utility as the root user to install it:

```
~]# dnf install vsftpd
```

15.1. FTP and SELinux

The vsftpd FTP daemon runs confined by default. SELinux policy defines how vsftpd interacts with files, processes, and with the system in general. For example, when an authenticated user logs in via FTP, they cannot read from or write to files in their home directories: SELinux prevents vsftpd from accessing user home directories by default. Also, by default, vsftpd does not have access to NFS or CIFS volumes, and anonymous users do not have write access, even if such write access is configured in the /etc/vsftpd/vsftpd.conf file. Booleans can be enabled to allow the previously mentioned access.

The following example demonstrates an authenticated user logging in, and an SELinux denial when trying to view files in their home directory. This example assumes that the *vsftpd* package is installed, that the SELinux targeted policy is used, and that SELinux is running in enforcing mode.

In Fedora, vsftpd only allows anonymous users to log in by default. To allow authenticated users
to log in, edit /etc/vsftpd/vsftpd.conf as root. Make sure the local_enable=YES option
is uncommented:

```
# Uncomment this to allow local users to log in.
local_enable=YES
```

2. Start the vsftpd service:

```
~]# systemctl start vsftpd.service
```

Confirm that the service is running. The output should include the information below (only the time stamp will differ):

```
~]# systemctl status vsftpd.service
```

```
vsftpd.service - Vsftpd ftp daemon
Loaded: loaded (/usr/lib/systemd/system/vsftpd.service; disabled)
Active: active (running) since Tue 2013-08-06 14:42:07 CEST; 6s ago
```

If the service was running before editing **vsftpd.conf**, restart the service to apply the configuration changes:

```
~]# systemctl restart vsftpd.service
```

3. Run the following command as the user you are currently logged in with. When prompted for your name, make sure your user name is displayed. If the correct user name is displayed, press **Enter**, otherwise, enter the correct user name:

```
~]$ ftp localhost
Connected to localhost (127.0.0.1).
220 (vsFTPd 2.1.0)
Name (localhost:username):
331 Please specify the password.
Password: Enter your password
500 OOPS: cannot change directory:/home/username
Login failed.
ftp>
```

An SELinux denial message similar to the following is logged:

```
setroubleshoot: SELinux is preventing the ftp daemon from reading users home directories (username). For complete SELinux messages. run sealert -l c366e889-2553-4c16-b73f-92f36a1730ce
```

5. Access to home directories has been denied by SELinux. This can be fixed by activating the ftp_home_dir Boolean. Enable this Boolean by running the following command as root:

```
~]# setsebool -P ftp_home_dir=1
```



Note

Do not use the -P option if you do not want changes to persist across reboots.

Try to log in again. Now that SELinux is allowing access to home directories using the ftp_home_dir Boolean, logging in will succeed.

15.2. Types

By default, anonymous users have read access to files in the /var/ftp/ directory when they log in via FTP. This directory is labeled with the public_content_t type, allowing only read access, even if write access is configured in /etc/vsftpd/vsftpd.conf. The public_content_t type is accessible to other services, such as Apache HTTP Server, Samba, and NFS.

Use one of the following types to share files through FTP:

public_content_t

Label files and directories you have created with the public_content_t type to share them read-only through vsftpd. Other services, such as Apache HTTP Server, Samba, and NFS, also have access to files labeled with this type. Files labeled with the public_content_t type cannot be written to, even if Linux permissions allow write access. If you require write access, use the public_content_rw_t type.

public_content_rw_t

Label files and directories you have created with the public_content_rw_t type to share them with read and write permissions through vsftpd. Other services, such as Apache HTTP Server, Samba, and NFS, also have access to files labeled with this type. Remember that Booleans for each service must be enabled before they can write to files labeled with this type.

15.3. Booleans

SELinux is based on the least level of access required for a service to run. Services can be run in a variety of ways; therefore, you need to specify how you run your services. Use the following Booleans to set up SELinux:

ftpd_anon_write

When disabled, this Boolean prevents vsftpd from writing to files and directories labeled with the public_content_rw_t type. Enable this Boolean to allow users to upload files using FTP. The directory where files are uploaded to must be labeled with the public_content_rw_t type and Linux permissions must be set accordingly.

ftpd_full_access

When this Boolean is enabled, only Linux (DAC) permissions are used to control access, and authenticated users can read and write to files that are not labeled with the public_content_t or public_content_rw_t types.

ftpd_use_cifs

Having this Boolean enabled allows vsftpd to access files and directories labeled with the cifs_t type; therefore, having this Boolean enabled allows you to share file systems mounted via Samba through vsftpd.

ftpd_use_nfs

Having this Boolean enabled allows vsftpd to access files and directories labeled with the nfs_t type; therefore, this Boolean allows you to share file systems mounted using NFS through vsftpd.

ftp_home_dir

Having this Boolean enabled allows authenticated users to read and write to files in their home directories. When this Boolean is disabled, attempting to download a file from a home directory results in an error such as **550 Failed to open file**. An SELinux denial message is logged.

ftpd_connect_db

Allow FTP daemons to initiate a connection to a database.

httpd_enable_ftp_server

Allow the httpd daemon to listen on the FTP port and act as a FTP server.

tftp_anon_write

Having this Boolean enabled allows TFTP access to a public directory, such as an area reserved for common files that otherwise has no special access restrictions.



Note

Due to the continuous development of the SELinux policy, the list above might not contain all Booleans related to the service at all times. To list them, run the following command:

```
~]$ getsebool -a | grep service_name
```

Run the following command to view description of a particular Boolean:

```
~]$ sepolicy booleans -b boolean_name
```

Note that the additional *policycoreutils-devel* package providing the sepolicy utility is required.

15.4. Configuration Examples

15.4.1. Uploading to an FTP site

The following example creates an FTP site that allows a dedicated user to upload files. It creates the directory structure and the required SELinux configuration changes:

1. Run the following command as the root user to enable access to FTP home directories:

```
~]# setsebool ftp_home_dir=1
```

2. Use the mkdir utility as root to create a new top-level directory:

```
~]# mkdir -p /myftp/pub
```

3. Set Linux permissions on the /myftp/pub/ directory to allow a Linux user write access. This example changes the owner and group from root to owner user1 and group root. Replace user1 with the user you want to give write access to:

```
~]# chown user1:root /myftp/pub
~]# chmod 775 /myftp/pub
```

The **chown** command changes the owner and group permissions. It changes the mode, allowing the **user1** user read, write, and execute permissions, and members of the root group read, write, and execute permissions. Everyone else has read and execute permissions; this is required to allow the Apache HTTP Server to read files from this directory.

4. When running SELinux, files and directories must be labeled correctly to allow access. Setting Linux permissions is not enough. Files labeled with the public_content_t type allow them to be read by FTP, Apache HTTP Server, Samba, and rsync. Files labeled with the public_content_rw_t type can be written to by FTP. Other services, such as Samba, require Booleans to be set before they can write to files labeled with the public_content_rw_t type. Label the top-level directory (/myftp/) with the public_content_t type, to prevent copied or newly-created files under this directory from being written to or modified by services. Run the following command as root to add the label change to file-context configuration:

```
~]# semanage fcontext -a -t public_content_t /myftp
```

5. Use the restorecon utility to apply the label change:

```
~]# restorecon -R -v /myftp/
restorecon reset /myftp context unconfined_u:object_r:default_t:s0-
>system_u:object_r:public_content_t:s0
```

6. Confirm /myftp/ is labeled with the public_content_t type, and /myftp/pub/ is labeled with the default_t type:

```
~]$ ls -dZ /myftp/
drwxr-xr-x. root root system_u:object_r:public_content_t:s0 /myftp/
~]$ ls -dZ /myftp/pub/
drwxrwxr-x. user1 root unconfined_u:object_r:default_t:s0 /myftp/pub/
```

7. FTP must be allowed to write to a directory before users can upload files through FTP. SELinux allows FTP to write to directories labeled with the public_content_rw_t type. This example uses /myftp/pub/ as the directory FTP can write to. Run the following command as root to add the label change to file-context configuration:

```
~]# semanage fcontext -a -t public_content_rw_t "/myftp/pub(/.*)?"
```

8. Use restorecon as root to apply the label change:

```
~]# restorecon -R -v /myftp/pub
restorecon reset /myftp/pub context system_u:object_r:default_t:s0-
>system_u:object_r:public_content_rw_t:s0
```

9. The ftpd_anon_write Boolean must be enabled to allow vsftpd to write to files that are labeled with the public_content_rw_t type. Run the following command as root to enable this Boolean:

```
~]# setsebool -P ftpd_anon_write on
```



Note

Do not use the **-P** option if you do not want changes to persist across reboots.

The following example demonstrates logging in via FTP and uploading a file. This example uses the **user1** user from the previous example, where **user1** is the dedicated owner of the /myftp/pub/directory:

1. Change into your home directory. Then, create a directory to store files to upload via FTP:

```
~]$ cd ~/
```

```
~]$ mkdir myftp
```

2. Change into the ~/myftp/ directory:

```
~]$ cd ~/myftp
```

In this directory, create an **ftpupload** file. Copy the following contents into this file:

```
File upload via FTP from a home directory.
```

3. Confirm that the ftpd_anon_write Boolean is enabled:

```
~]$ getsebool ftpd_anon_write
ftpd_anon_write --> on
```

If this Boolean is disabled, run the following command as root to enable it:

```
~]# setsebool -P ftpd_anon_write on
```



Note

Do not use the **-P** option if you do not want the change to persist across reboots.

4. Start the vsftpd service:

```
~]# systemctl start vsftpd.service
```

5. Run the following command. When prompted for a user name, enter the user name of the user who has write access, then, enter the correct password for that user:

```
~]$ ftp localhost
Connected to localhost (127.0.0.1).
220 (vsFTPd 2.1.0)
Name (localhost:username):
331 Please specify the password.
Password: Enter the correct password
230 Login successful.
Remote system type is UNIX.
Using binary mode to transfer files.
ftp> cd myftp
250 Directory successfully changed.
ftp> put ftpupload
```

```
local: ftpupload remote: ftpupload
227 Entering Passive Mode (127,0,0,1,241,41).
150 Ok to send data.
226 File receive OK.
ftp> 221 Goodbye.
```

The upload succeeds as the $ftpd_anon_write$ Boolean is enabled.

Network File System

A Network File System (NFS) allows remote hosts to mount file systems over a network and interact with those file systems as though they are mounted locally. This enables system administrators to consolidate resources onto centralized servers on the network.¹

In Fedora, the *nfs-utils* package is required for full NFS support. Run the following command to see if the *nfs-utils* is installed:

```
~]$ rpm -q nfs-utils
package nfs-utils is not installed
```

If it is not installed and you want to use NFS, use the DNF utility as root to install it:

```
~]# dnf install nfs-utils
```

16.1. NFS and SELinux

When running SELinux, the NFS daemons are confined by default except the nfsd process, which is labeled with the unconfined kernel_t domain type. The SELinux policy allows NFS to share files by default. Also, passing SELinux labels between a client and the server is supported, which provides better security control of confined domains accessing NFS volumes. For example, when a home directory is set up on an NFS volume, it is possible to specify confined domains that are able to access only the home directory and not other directories on the volume. Similarly, applications, such as Secure Virtualization, can set the label of an image file on an NFS volume, thus increasing the level of separation of virtual machines.

The support for labeled NFS is disabled by default. To enable it, see Section 16.4.1, "Enabling SELinux Labeled NFS Support".

16.2. Types

By default, mounted NFS volumes on the client side are labeled with a default context defined by policy for NFS. In common policies, this default context uses the nfs_t type. The root user is able to override the default type using the **mount -context** option. The following types are used with NFS. Different types allow you to configure flexible access:

```
var_lib_nfs_t
```

This type is used for existing and new files copied to or created in the /var/lib/nfs/ directory. This type should not need to be changed in normal operation. To restore changes to the default settings, run the restorecon -R -v /var/lib/nfs command as the root user.

```
nfsd_exec_t
```

The /usr/sbin/rpc.nfsd file is labeled with the nfsd_exec_t, as are other system executables and libraries related to NFS. Users should not label any files with this type. nfsd_exec_t will transition to nfsd_t.

¹ See the Network File System (NFS) chapter in the Storage Administration Guide [https://access.redhat.com/site/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/Storage_Administration_Guide/ch-nfs.html] for more information.

16.3. Booleans

SELinux is based on the least level of access required for a service to run. Services can be run in a variety of ways; therefore, you need to specify how you run your services. Use the following Booleans to set up SELinux:

ftpd_use_nfs

When enabled, this Boolean allows the ftpd daemon to access NFS volumes.

cobbler_use_nfs

When enabled, this Boolean allows the cobblerd daemon to access NFS volumes.

git_system_use_nfs

When enabled, this Boolean allows the Git system daemon to read system shared repositories on NFS volumes.

httpd_use_nfs

When enabled, this Boolean allows the httpd daemon to access files stored on NFS volumes.

samba_share_nfs

When enabled, this Boolean allows the smbd daemon to share NFS volumes. When disabled, this Boolean prevents smbd from having full access to NFS shares via Samba.

sanlock_use_nfs

When enabled, this Boolean allows the sanlock daemon to manage NFS volumes.

sge_use_nfs

When enabled, this Boolean allows the sge scheduler to access NFS volumes.

use_nfs_home_dirs

When enabled, this Boolean adds support for NFS home directories.

virt_use_nfs

When enabled, this Boolean allows confident virtual guests to manage files on NFS volumes.

xen use nfs

When enabled, this Boolean allows Xen to manage files on NFS volumes.

git_cgi_use_nfs

When enabled, this Boolean allows the Git Common Gateway Interface (CGI) to access NFS volumes.



Note

Due to the continuous development of the SELinux policy, the list above might not contain all Booleans related to the service at all times. To list them, run the following command:

```
~]$ getsebool -a | grep service_name
```

Run the following command to view description of a particular Boolean:

```
~]$ sepolicy booleans -b boolean_name
```

Note that the additional policycoreutils-devel package providing the sepolicy utility is required.

16.4. Configuration Examples

16.4.1. Enabling SELinux Labeled NFS Support

The following example demonstrates how to enable SELinux labeled NFS support. This example assumes that the *nfs-utils* package is installed, that the SELinux targeted policy is used, and that SELinux is running in enforcing mode.



Note

Steps 1-3 are supposed to be performed on the NFS server, nfs-srv.

1. If the NFS server is running, stop it:

```
[nfs-srv]# systemctl stop nfs
```

Confirm that the server is stopped:

```
[nfs-srv]# systemctl status nfs
nfs-server.service - NFS Server
Loaded: loaded (/usr/lib/systemd/system/nfs-server.service; disabled)
Active: inactive (dead)
```

Edit the /etc/sysconfig/nfs file to set the RPCNFSDARGS flag to "-V 4.2":

```
# Optional arguments passed to rpc.nfsd. See rpc.nfsd(8)
RPCNFSDARGS="-V 4.2"
```

3. Start the server again and confirm that it is running. The output will contain information below, only the time stamp will differ:

```
[nfs-srv]# systemctl start nfs
```

```
[nfs-srv]# systemctl status nfs
nfs-server.service - NFS Server
Loaded: loaded (/usr/lib/systemd/system/nfs-server.service; disabled)
Active: active (exited) since Wed 2013-08-28 14:07:11 CEST; 4s ago
```

4. On the client side, mount the NFS server:

```
[nfs-client]# mount -o v4.2 server:mntpoint localmountpoint
```

5. All SELinux labels are now successfully passed from the server to the client:

```
[nfs-srv]$ ls -Z file
-rw-rw-r--. user user unconfined_u:object_r:svirt_image_t:s0 file
[nfs-client]$ ls -Z file
-rw-rw-r--. user user unconfined_u:object_r:svirt_image_t:s0 file
```



Note

If you enable labeled NFS support for home directories or other content, the content will be labeled the same as it was on an EXT file system. Also note that mounting systems with different versions of NFS or an attempt to mount a server that does not support labeled NFS could cause errors to be returned.

Berkeley Internet Name Domain

BIND performs name resolution services via the named daemon. BIND lets users locate computer resources and services by name instead of numerical addresses.

In Fedora, the *bind* package provides a DNS server. Run the following command to see if the *bind* package is installed:

```
~]$ rpm -q bind
package bind is not installed
```

If it is not installed, use the DNF utility as the root user to install it:

```
~]# dnf install bind
```

17.1. BIND and SELinux

The default permissions on the /var/named/slaves/, /var/named/dynamic/ and /var/named/data/ directories allow zone files to be updated via zone transfers and dynamic DNS updates. Files in /var/named/ are labeled with the named_zone_t type, which is used for master zone files.

For a slave server, configure the /etc/named.conf file to place slave zones in /var/named/slaves/. The following is an example of a domain entry in /etc/named.conf for a slave DNS server that stores the zone file for testdomain.com in /var/named/slaves/:

```
zone "testdomain.com" {
  type slave;
  masters { IP-address; };
  file "/var/named/slaves/db.testdomain.com";
  };
```

If a zone file is labeled named_zone_t, the named_write_master_zones Boolean must be enabled to allow zone transfers and dynamic DNS to update the zone file. Also, the mode of the parent directory has to be changed to allow the named user or group read, write and execute access.

If zone files in /var/named/ are labeled with the named_cache_t type, a file system relabel or running restorecon -R /var/ will change their type to named_zone_t.

17.2. Types

The following types are used with BIND. Different types allow you to configure flexible access:

```
named_zone_t
```

Used for master zone files. Other services cannot modify files of this type. The named daemon can only modify files of this type if the named_write_master_zones Boolean is enabled.

```
named cache t
```

By default, named can write to files labeled with this type, without additional Booleans being set. Files copied or created in the /var/named/slaves/,/var/named/dynamic/ and /var/named/data/ directories are automatically labeled with the named_cache_t type.

named_var_run_t

Files copied or created in the /var/run/bind/, /var/run/named/, and /var/run/unbound/ directories are automatically labeled with the named_var_run_t type.

named_conf_t

BIND-related configuration files, usually stored in the **/etc/** directory, are automatically labeled with the named_conf_t type.

named_exec_t

BIND-related executable files, usually stored in the **/usr/sbin/** directory, are automatically labeled with the named_exec_t type.

named_log_t

BIND-related log files, usually stored in the **/var/log/** directory, are automatically labeled with the named_log_t type.

named_unit_file_t

Executable BIND-related files in the /usr/lib/systemd/system/ directory are automatically labeled with the named_unit_file_t type.

17.3. Booleans

SELinux is based on the least level of access required for a service to run. Services can be run in a variety of ways; therefore, you need to specify how you run your services. Use the following Booleans to set up SELinux:

named_write_master_zones

When disabled, this Boolean prevents named from writing to zone files or directories labeled with the named_zone_t type. The daemon does not usually need to write to zone files; but in the case that it needs to, or if a secondary server needs to write to zone files, enable this Boolean to allow this action.

named_tcp_bind_http_port

When enabled, this Boolean allows BIND to bind an Apache port.



Note

Due to the continuous development of the SELinux policy, the list above might not contain all Booleans related to the service at all times. To list them, run the following command:

```
~]$ getsebool -a | grep service_name
```

Run the following command to view description of a particular Boolean:

```
~]$ sepolicy booleans -b boolean_name
```

Note that the additional *policycoreutils-devel* package providing the sepolicy utility is required.

17.4. Configuration Examples

17.4.1. Dynamic DNS

BIND allows hosts to update their records in DNS and zone files dynamically. This is used when a host computer's IP address changes frequently and the DNS record requires real-time modification.

Use the /var/named/dynamic/ directory for zone files you want updated via dynamic DNS. Files created in or copied into this directory inherit Linux permissions that allow named to write to them. As such files are labeled with the named_cache_t type, SELinux allows named to write to them.

If a zone file in /var/named/dynamic/ is labeled with the named_zone_t type, dynamic DNS updates may not be successful for a certain period of time as the update needs to be written to a journal first before being merged. If the zone file is labeled with the named_zone_t type when the journal attempts to be merged, an error such as the following is logged:

```
named[PID]: dumping master file: rename: /var/named/dynamic/zone-name: permission denied
```

Also, the following SELinux denial message is logged:

```
setroubleshoot: SELinux is preventing named (named_t) "unlink" to zone-name (named_zone_t)
```

To resolve this labeling issue, use the restorecon utility as root:

```
~]# restorecon -R -v /var/named/dynamic
```

Concurrent Versioning System

The Concurrent Versioning System (CVS) is a free revision-control system. It is used to monitor and keep track of modifications to a central set of files which are usually accessed by several different users. It is commonly used by programmers to manage a source code repository and is widely used by open source developers.

In Fedora, the cvs package provides CVS. Run the following command to see if the cvs package is installed:

```
~]$ rpm -q cvs
package cvs is not installed
```

If it is not installed and you want to use CVS, use the DNF utility as root to install it:

~]# dnf install cvs

18.1. CVS and SELinux

The cvs daemon runs labeled with the cvs_t type. By default in Fedora, CVS is only allowed to read and write certain directories. The label cvs_data_t defines which areas cvs has read and write access to. When using CVS with SELinux, assigning the correct label is essential for clients to have full access to the area reserved for CVS data.

18.2. Types

The following types are used with CVS. Different types allow you to configure flexible access:

cvs_data_t

This type is used for data in a CVS repository. CVS can only gain full access to data if it has this type.

cvs_exec_t

This type is used for the /usr/bin/cvs binary.

18.3. Booleans

SELinux is based on the least level of access required for a service to run. Services can be run in a variety of ways; therefore, you need to specify how you run your services. Use the following Booleans to set up SELinux:

cvs_read_shadow

This Boolean allows the cvs daemon to access the /etc/shadow file for user authentication.



Note

Due to the continuous development of the SELinux policy, the list above might not contain all Booleans related to the service at all times. To list them, run the following command:

```
~]$ getsebool -a | grep service_name
```

Run the following command to view description of a particular Boolean:

```
~]$ sepolicy booleans -b boolean_name
```

Note that the additional *policycoreutils-devel* package providing the sepolicy utility is required.

18.4. Configuration Examples

18.4.1. Setting up CVS

This example describes a simple CVS setup and an SELinux configuration which allows remote access. Two hosts are used in this example; a CVS server with a host name of cvs-srv with an IP address of 192.168.1.1 and a client with a host name of cvs-client and an IP address of 192.168.1.100. Both hosts are on the same subnet (192.168.1.0/24). This is an example only and assumes that the cvs and xinetd packages are installed, that the SELinux targeted policy is used, and that SELinux is running in enforced mode.

This example will show that even with full DAC permissions, SELinux can still enforce policy rules based on file labels and only allow access to certain areas that have been specifically labeled for access by CVS.



Note

Steps 1-9 are supposed be performed on the CVS server, cvs-srv.

1. This example requires the *cvs* and *xinetd* packages. Confirm that the packages are installed:

```
[cvs-srv]$ rpm -q cvs xinetd
package cvs is not installed
package xinetd is not installed
```

If they are not installed, use the DNF utility as root to install it:

```
[cvs-srv]# dnf install cvs xinetd
```

2. Run the following command as root to create a group named **CVS**:

```
[cvs-srv]# groupadd CVS
```

This can by also done by using the system-config-users utility.

- 3. Create a user with a user name of **cvsuser** and make this user a member of the CVS group. This can be done using system-config-users.
- 4. Edit the /etc/services file and make sure that the CVS server has uncommented entries looking similar to the following:

```
cvspserver 2401/tcp  # CVS client/server operations
cvspserver 2401/udp  # CVS client/server operations
```

5. Create the CVS repository in the root area of the file system. When using SELinux, it is best to have the repository in the root file system so that recursive labels can be given to it without affecting any other subdirectories. For example, as root, create a /cvs/ directory to house the repository:

```
[root@cvs-srv]# mkdir /cvs
```

6. Give full permissions to the /cvs/ directory to all users:

```
[root@cvs-srv]# chmod -R 777 /cvs
```



Warning

This is an example only and these permissions should not be used in a production system.

7. Edit the /etc/xinetd.d/cvs file and make sure that the CVS section is uncommented and configured to use the /cvs/ directory. The file should look similar to:

```
service cvspserver
{
  disable = no
  port = 2401
  socket_type = stream
  protocol = tcp
  wait = no
  user = root
  passenv = PATH
  server = /usr/bin/cvs
  env = HOME=/cvs
  server_args = -f --allow-root=/cvs pserver
# bind = 127.0.0.1
```

8. Start the xinetd daemon:

```
[cvs-srv]# systemctl start xinetd.service
```

- 9. Add a rule which allows inbound connections through TCP on port 2401 by using the system-config-firewall utility.
- 10. On the client side, run the following command as the **cvsuser** user:

```
[cvsuser@cvs-client]$ cvs -d /cvs init
```

11. At this point, CVS has been configured but SELinux will still deny logins and file access. To demonstrate this, set the \$CVSR00T variable on cvs-client and try to log in remotely. The following step is supposed to be performed on cvs-client:

```
[cvsuser@cvs-client]$ export CVSROOT=:pserver:cvsuser@192.168.1.1:/cvs
[cvsuser@cvs-client]$
[cvsuser@cvs-client]$ cvs login
Logging in to :pserver:cvsuser@192.168.1.1:2401/cvs
CVS password: ********
cvs [login aborted]: unrecognized auth response from 192.168.100.1: cvs pserver: cannot open /cvs/CVSROOT/config: Permission denied
```

SELinux has blocked access. In order to get SELinux to allow this access, the following step is supposed to be performed on cvs-srv:

12. Change the context of the /cvs/ directory as root in order to recursively label any existing and new data in the /cvs/ directory, giving it the cvs_data_t type:

```
[root@cvs-srv]# semanage fcontext -a -t cvs_data_t '/cvs(/.*)?'
[root@cvs-srv]# restorecon -R -v /cvs
```

13. The client, cvs-client should now be able to log in and access all CVS resources in this repository:

```
[cvsuser@cvs-client]$ export CVSROOT=:pserver:cvsuser@192.168.1.1:/cvs
[cvsuser@cvs-client]$
[cvsuser@cvs-client]$ cvs login
Logging in to :pserver:cvsuser@192.168.1.1:2401/cvs
CVS password: ********
[cvsuser@cvs-client]$
```

Squid Caching Proxy

Squid is a high-performance proxy caching server for web clients, supporting FTP, Gopher, and HTTP data objects. It reduces bandwidth and improves response times by caching and reusing frequently-requested web pages.¹

In Fedora, the *squid* package provides the Squid Caching Proxy. Run the following command to see if the *squid* package is installed:

```
~]$ rpm -q squid
package squid is not installed
```

If it is not installed and you want to use squid, use the DNF utility as root to install it:

```
~]# dnf install squid
```

19.1. Squid Caching Proxy and SELinux

When SELinux is enabled, Squid runs confined by default. Confined processes run in their own domains, and are separated from other confined processes. If a confined process is compromised by an attacker, depending on SELinux policy configuration, an attacker's access to resources and the possible damage they can do is limited. The following example demonstrates the Squid processes running in their own domain. This example assumes the *squid* package is installed:

1. Run the **getenforce** command to confirm SELinux is running in enforcing mode:

```
~]$ getenforce
Enforcing
```

The command returns **Enforcing** when SELinux is running in enforcing mode.

2. Run the following command as the root user to start the squid daemon:

```
~]# systemctl start squid.service
```

Confirm that the service is running. The output should include the information below (only the time stamp will differ):

```
~]# systemctl status squid.service
squid.service - Squid caching proxy
Loaded: loaded (/usr/lib/systemd/system/squid.service; disabled)
Active: active (running) since Mon 2013-08-05 14:45:53 CEST; 2s ago
```

Run the following command to view the squid processes:

```
~]$ ps -eZ | grep squid
```

¹ See the *Squid Caching Proxy* [http://www.squid-cache.org/] project page for more information.

```
system_u:system_r:squid_t:s0 27018 ? 00:00:00 squid system_u:system_r:squid_t:s0 27020 ? 00:00:00 log_file_daemon
```

The SELinux context associated with the squid processes is

system_u:system_r:squid_t:s0. The second last part of the context, squid_t, is the type. A type defines a domain for processes and a type for files. In this case, the Squid processes are running in the squid_t domain.

SELinux policy defines how processes running in confined domains, such as squid_t, interact with files, other processes, and the system in general. Files must be labeled correctly to allow squid access to them.

When the /etc/squid/squid.conf file is configured so squid listens on a port other than the default TCP ports 3128, 3401 or 4827, the semanage port command must be used to add the required port number to the SELinux policy configuration. The following example demonstrates configuring squid to listen on a port that is not initially defined in SELinux policy configuration for it, and, as a consequence, the server failing to start. This example also demonstrates how to then configure the SELinux system to allow the daemon to successfully listen on a non-standard port that is not already defined in the policy. This example assumes the squid package is installed. Run each command in the example as the root user:

Connfirm the squid daemon is not running:

```
~]# systemctl status squid.service
squid.service - Squid caching proxy
Loaded: loaded (/usr/lib/systemd/system/squid.service; disabled)
Active: inactive (dead)
```

If the output differs, run stop the process:

```
~]# systemctl stop squid.service
```

2. Run the following command to view the ports SELinux allows squid to listen on:

```
~]# semanage port -l | grep -w -i squid_port_t
squid_port_t tcp 3401, 4827
squid_port_t udp 3401, 4827
```

3. Edit /etc/squid/squid.conf as root. Configure the http_port option so it lists a port that is not configured in SELinux policy configuration for squid. In this example, the daemon is configured to listen on port 10000:

```
# Squid normally listens to port 3128
http_port 10000
```

4. Run the **setsebool** command to make sure the squid_connect_any Boolean is set to off. This ensures squid is only permitted to operate on specific ports:

```
~]# setsebool -P squid_connect_any 0
```

5. Start the squid daemon:

```
~]# systemctl start squid.service
Job for squid.service failed. See 'systemctl status squid.service' and 'journalctl -xn' for details.
```

An SELinux denial message similar to the following is logged:

```
localhost setroubleshoot: SELinux is preventing the squid (squid_t) from binding to port 10000. For complete SELinux messages. run sealert -1 97136444-4497-4fff-a7a7-c4d8442db982
```

6. For SELinux to allow squid to listen on port 10000, as used in this example, the following command is required:

```
~]# semanage port -a -t squid_port_t -p tcp 10000
```

7. Start squid again and have it listen on the new port:

```
~]# systemctl start squid.service
```

8. Now that SELinux has been configured to allow Squid to listen on a non-standard port (TCP 10000 in this example), it starts successfully on this port.

19.2. Types

Type Enforcement is the main permission control used in SELinux targeted policy. All files and processes are labeled with a type: types define a domain for processes and a type for files. SELinux policy rules define how types access each other, whether it be a domain accessing a type, or a domain accessing another domain. Access is only allowed if a specific SELinux policy rule exists that allows it.

The following types are used with Squid. Different types allow you to configure flexible access:

```
httpd_squid_script_exec_t
```

This type is used for utilities such as **cachemgr.cgi**, which provides a variety of statistics about Squid and its configuration.

```
squid_cache_t
```

Use this type for data that is cached by Squid, as defined by the cache_dir directive in /etc/squid/squid.conf. By default, files created in or copied into the /var/cache/squid/ and /var/spool/squid/ directories are labeled with the squid_cache_t type. Files for the squidGuard² URL redirector plug-in for squid created in or copied to the /var/squidGuard/ directory are also labeled with the squid_cache_t type. Squid is only able to use files and directories that are labeled with this type for its cached data.

² http://www.squidguard.org/

squid_conf_t

This type is used for the directories and files that Squid uses for its configuration. Existing files, or those created in or copied to the /etc/squid/ and /usr/share/squid/ directories are labeled with this type, including error messages and icons.

squid_exec_t

This type is used for the squid binary, /usr/sbin/squid.

squid_log_t

This type is used for logs. Existing files, or those created in or copied to /var/log/squid/ or /var/log/squidGuard/ must be labeled with this type.

```
squid_initrc_exec_t
```

This type is used for the initialization file required to start squid which is located at /etc/rc.d/init.d/squid.

```
squid_var_run_t
```

This type is used by files in the /var/run/ directory, especially the process id (PID) named / var/run/squid.pid which is created by Squid when it runs.

19.3. Booleans

SELinux is based on the least level of access required for a service to run. Services can be run in a variety of ways; therefore, you need to specify how you run your services. Use the following Booleans to set up SELinux:

squid_connect_any

When enabled, this Boolean permits Squid to initiate a connection to a remote host on any port.

squid_use_tproxy

When enabled, this Boolean allows Squid to run as a transparent proxy.



Note

Due to the continuous development of the SELinux policy, the list above might not contain all Booleans related to the service at all times. To list them, run the following command:

```
~]$ getsebool -a | grep service_name
```

Run the following command to view description of a particular Boolean:

```
~]$ sepolicy booleans -b boolean_name
```

Note that the additional *policycoreutils-devel* package providing the sepolicy utility is required.

19.4. Configuration Examples

19.4.1. Squid Connecting to Non-Standard Ports

The following example provides a real-world demonstration of how SELinux complements Squid by enforcing the above Boolean and by default only allowing access to certain ports. This example will then demonstrate how to change the Boolean and show that access is then allowed.

Note that this is an example only and demonstrates how SELinux can affect a simple configuration of Squid. Comprehensive documentation of Squid is beyond the scope of this document. See the official *Squid documentation*³ for further details. This example assumes that the Squid host has two network interfaces, Internet access, and that any firewall has been configured to allow access on the internal interface using the default TCP port on which Squid listens (TCP 3128).

1. Confirm that the *squid* is installed:

```
~]$ rpm -q squid
package squid is not installed
```

If the package is not installed, use the DNF utility as root to install it:

```
~]# dnf install squid
```

Edit the main configuration file, /etc/squid/squid.conf, and confirm that the cache_dir directive is uncommented and looks similar to the following:

```
cache_dir ufs /var/spool/squid 100 16 256
```

This line specifies the default settings for the cache_dir directive to be used in this example; it consists of the Squid storage format (ufs), the directory on the system where the cache resides (/var/spool/squid), the amount of disk space in megabytes to be used for the cache (100), and finally the number of first-level and second-level cache directories to be created (16 and 256 respectively).

- 3. In the same configuration file, make sure the http_access allow localnet directive is uncommented. This allows traffic from the **localnet** ACL which is automatically configured in a default installation of Squid on Fedora. It will allow client machines on any existing RFC1918 network to have access through the proxy, which is sufficient for this simple example.
- 4. In the same configuration file, make sure the visible_hostname directive is uncommented and is configured to the host name of the machine. The value should be the fully qualified domain name (FQDN) of the host:

visible_hostname	squid.example.com

³ http://www.squid-cache.org/Doc/

5. As root, run the following command to start the squid daemon. As this is the first time squid has started, this command will initialise the cache directories as specified above in the cache_dir directive and will then start the daemon:

```
~]# systemctl start squid.service
```

Ensure that squid starts successfully. The output will include the information below, only the time stamp will differ:

```
~]# systemctl status squid.service
squid.service - Squid caching proxy
Loaded: loaded (/usr/lib/systemd/system/squid.service; disabled)
Active: active (running) since Thu 2014-02-06 15:00:24 CET; 6s ago
```

6. Confirm that the squid process ID (PID) has started as a confined service, as seen here by the squid_var_run_t value:

```
~]# ls -lz /var/run/squid.pid
-rw-r--r--. root squid unconfined_u:object_r:squid_var_run_t:s0 /var/run/squid.pid
```

- 7. At this point, a client machine connected to the **localnet** ACL configured earlier is successfully able to use the internal interface of this host as its proxy. This can be configured in the settings for all common web browsers, or system-wide. Squid is now listening on the default port of the target machine (TCP 3128), but the target machine will only allow outgoing connections to other services on the Internet via common ports. This is a policy defined by SELinux itself. SELinux will deny access to non-standard ports, as shown in the next step:
- 8. When a client makes a request using a non-standard port through the Squid proxy such as a website listening on TCP port 10000, a denial similar to the following is logged:

```
SELinux is preventing the squid daemon from connecting to network port 10000
```

To allow this access, the squid_connect_any Boolean must be modified, as it is disabled by default:

~]# setsebool -P squid_connect_any on



Note

Do not use the **-P** option if you do not want **setsebool** changes to persist across reboots.

10. The client will now be able to access non-standard ports on the Internet as Squid is now permitted to initiate connections to any port, on behalf of its clients.

MariaDB (a replacement for MySQL)

The MariaDB database is a multi-user, multi-threaded SQL database server that consists of the MariaDB server daemon (mysqld) and many client programs and libraries.¹

In Fedora, the *mariadb-server* package provides MariaDB. Run the following command to see if the *mariadb-server* package is installed:

```
~]$ rpm -q mariadb-server package mariadb-server is not installed
```

If it is not installed, use the DNF utility as root to install it:

```
~]# dnf install mariadb-server
```

20.1. MariaDB and SELinux

When MariaDB is enabled, it runs confined by default. Confined processes run in their own domains, and are separated from other confined processes. If a confined process is compromised by an attacker, depending on SELinux policy configuration, an attacker's access to resources and the possible damage they can do is limited. The following example demonstrates the MariaDB processes running in their own domain. This example assumes the *mariadb-server* package is installed:

Run the getenforce command to confirm SELinux is running in enforcing mode:

```
~]$ getenforce
Enforcing
```

The command returns **Enforcing** when SELinux is running in enforcing mode.

2. Run the following command as the root user to start mariadb:

```
~]# systemctl start mariadb.service
```

Confirm that the service is running. The output should include the information below (only the time stamp will differ):

```
~]# systemctl status mariadb.service
mariadb.service - MariaDB database server
Loaded: loaded (/usr/lib/systemd/system/mariadb.service; disabled)
Active: active (running) since Mon 2013-08-05 11:20:11 CEST; 3h 28min ago
```

3. Run the following command to view the mysgld processes:

```
~]$ ps -eZ | grep mysqld
system_u:system_r:mysqld_safe_t:s0 12831 ? 00:00:00 mysqld_safe
system_u:system_r:mysqld_t:s0 13014 ? 00:00:00 mysqld
```

¹ See the *MariaDB* [https://mariadb.org/] project page for more information.

The SELinux context associated with the mysqld processes is

system_u:system_r:mysqld_t:s0. The second last part of the context, mysqld_t, is the type. A type defines a domain for processes and a type for files. In this case, the mysqld processes are running in the mysqld_t domain.

20.2. Types

Type Enforcement is the main permission control used in SELinux targeted policy. All files and processes are labeled with a type: types define a domain for processes and a type for files. SELinux policy rules define how types access each other, whether it be a domain accessing a type, or a domain accessing another domain. Access is only allowed if a specific SELinux policy rule exists that allows it.

The following types are used with mysqld. Different types allow you to configure flexible access:

mysqld_db_t

This type is used for the location of the MariaDB database. In Fedora, the default location for the database is the <code>/var/lib/mysql/</code> directory, however this can be changed. If the location for the MariaDB database is changed, the new location must be labeled with this type. See the example in <code>Section 20.4.1</code>, "MariaDB Changing Database Location" for instructions on how to change the default database location and how to label the new section appropriately.

mysqld_etc_t

This type is used for the MariaDB main configuration file /etc/my.cnf and any other configuration files in the /etc/mysql/ directory.

mysqld_exec_t

This type is used for the mysqld binary located at /usr/libexec/mysqld, which is the default location for the MariaDB binary on Fedora. Other systems may locate this binary at /usr/sbin/mysqld which should also be labeled with this type.

mysqld_unit_file_t

This type is used for executable MariaDB-related files located in the /usr/lib/systemd/system/ directory by default in Fedora.

mysqld_log_t

Logs for MariaDB need to be labeled with this type for proper operation. All log files in the /var/log/ directory matching the mysql.* wildcard must be labeled with this type.

```
mysqld_var_run_t
```

This type is used by files in the /var/run/mariadb/ directory, specifically the process id (PID) named /var/run/mariadb/mariadb.pid which is created by the mysqld daemon when it runs. This type is also used for related socket files such as /var/lib/mysql/mysql.sock. Files such as these must be labeled correctly for proper operation as a confined service.

20.3. Booleans

SELinux is based on the least level of access required for a service to run. Services can be run in a variety of ways; therefore, you need to specify how you run your services. Use the following Booleans to set up SELinux:

selinuxuser_mysql_connect_enabled

When enabled, this Boolean allows users to connect to the local MariaDB server.

exim_can_connect_db

When enabled, this Boolean allows the exim mailer to initiate connections to a database server.

ftpd_connect_db

When enabled, this Boolean allows ftp daemons to initiate connections to a database server.

httpd_can_network_connect_db

Enabling this Boolean is required for a web server to communicate with a database server.



Note

Due to the continuous development of the SELinux policy, the list above might not contain all Booleans related to the service at all times. To list them, run the following command:

```
~]$ getsebool -a | grep service_name
```

Run the following command to view description of a particular Boolean:

```
~]$ sepolicy booleans -b boolean_name
```

Note that the additional *policycoreutils-devel* package providing the sepolicy utility is required.

20.4. Configuration Examples

20.4.1. MariaDB Changing Database Location

When using Fedora, the default location for MariaDB to store its database is /var/lib/mysql/. This is where SELinux expects it to be by default, and hence this area is already labeled appropriately for you, using the mysqld_db_t type.

The location where the database is stored can be changed depending on individual environment requirements or preferences, however it is important that SELinux is aware of this new location; that it is labeled accordingly. This example explains how to change the location of a MariaDB database and then how to label the new location so that SELinux can still provide its protection mechanisms to the new area based on its contents.

Note that this is an example only and demonstrates how SELinux can affect MariaDB. Comprehensive documentation of MariaDB is beyond the scope of this document. See the official *MariaDB* documentation² for further details. This example assumes that the *mariadb-server* and *setroubleshoot-server* packages are installed, that the auditd service is running, and that there is a valid database in the default location of /var/lib/mysql/.

1. View the SELinux context of the default database location for mysql:

```
~]# ls -lZ /var/lib/mysql
drwx-----. mysql mysql system_u:object_r:mysqld_db_t:s0 mysql
```

This shows mysqld_db_t which is the default context element for the location of database files. This context will have to be manually applied to the new database location that will be used in this example in order for it to function properly.

² https://mariadb.com/kb/en/mariadb-documentation/

Run the following command and enter the mysqld root password to show the available databases:

```
~]# mysqlshow -u root -p
Enter password: ******

| Databases |
+-----+
| information_schema |
| mysql |
| test |
| wikidb |
+-----+
```

3. Stop the mysqld daemon:

```
~]# systemctl stop mariadb.service
```

4. Create a new directory for the new location of the database(s). In this example, /mysql/ is used:

```
~]# mkdir -p /mysql
```

5. Copy the database files from the old location to the new location:

```
~]# cp -R /var/lib/mysql/* /mysql/
```

6. Change the ownership of this location to allow access by the mysql user and group. This sets the traditional Unix permissions which SELinux will still observe:

```
~]# chown -R mysql:mysql /mysql
```

7. Run the following command to see the initial context of the new directory:

```
~]# ls -lz /mysql
drwxr-xr-x. mysql mysql unconfined_u:object_r:usr_t:s0 mysql
```

The context usr_t of this newly created directory is not currently suitable to SELinux as a location for MariaDB database files. Once the context has been changed, MariaDB will be able to function properly in this area.

8. Open the main MariaDB configuration file /etc/my.cnf with a text editor and modify the datadir option so that it refers to the new location. In this example the value that should be entered is /mysql:

```
[mysqld]
datadir=/mysql
```

Save this file and exit.

 Start mysqld. The service should fail to start, and a denial message will be logged to the /var/ log/messages file:

```
~]# systemctl start mariadb.service
Job for mariadb.service failed. See 'systemctl status postgresql.service' and
'journalctl -xn' for details.
```

However, if the audit daemon is running alongside the setroubleshoot service, the denial will be logged to the /var/log/audit/audit.log file instead:

```
SELinux is preventing /usr/libexec/mysqld "write" access on /mysql. For complete SELinux messages. run sealert -l b3f01aff-7fa6-4ebe-ad46-abaef6f8ad71
```

The reason for this denial is that /mysql/ is not labeled correctly for MariaDB data files. SELinux is stopping MariaDB from having access to the content labeled as usr_t. Perform the following steps to resolve this problem:

10. Run the following command to add a context mapping for /mysq1/. Note that the semanage utility is not installed by default. If it is missing on your system, install the *policycoreutils-python* package.

```
~]# semanage fcontext -a -t mysqld_db_t "/mysql(/.*)?"
```

11. This mapping is written to the /etc/selinux/targeted/contexts/files/file_contexts.local file:

```
~]# grep -i mysql /etc/selinux/targeted/contexts/files/file_contexts.local
/mysql(/.*)? system_u:object_r:mysqld_db_t:s0
```

12. Now use the restorecon utility to apply this context mapping to the running system:

```
~]# restorecon -R -v /mysql
```

13. Now that the /mysql/ location has been labeled with the correct context for MariaDB, mysqld starts:

```
~]# systemctl start mariadb.service
```

14. Confirm the context has changed for /mysql/:

```
~]$ ls -lZ /mysql drwxr-xr-x. mysql mysql system_u:object_r:mysqld_db_t:s0 mysql
```

15. The location has been changed and labeled, and mysqld has started successfully. At this point all running services should be tested to confirm normal operation.

PostgreSQL

PostgreSQL is an Object-Relational database management system (DBMS).¹

In Fedora, the *postgresql-server* package provides PostgreSQL. Run the following command to see if the *postgresql-server* package is installed:

```
~]# rpm -q postgresql-server
```

If it is not installed, use the DNF utility as root to install it:

```
~]# dnf install postgresql-server
```

21.1. PostgreSQL and SELinux

When PostgreSQL is enabled, it runs confined by default. Confined processes run in their own domains, and are separated from other confined processes. If a confined process is compromised by an attacker, depending on SELinux policy configuration, an attacker's access to resources and the possible damage they can do is limited. The following example demonstrates the PostgreSQL processes running in their own domain. This example assumes the *postgresql-server* package is installed:

1. Run the **getenforce** command to confirm SELinux is running in enforcing mode:

```
~]$ getenforce
Enforcing
```

The command returns **Enforcing** when SELinux is running in enforcing mode.

Run the following command as the root user to start postgresql:

```
~]# systemctl start postgresql.service
```

Confirm that the service is running. The output should include the information below (only the time stamp will differ):

```
~]# systemctl start postgresql.service
postgresql.service - PostgreSQL database server
Loaded: loaded (/usr/lib/systemd/system/postgresql.service; disabled)
Active: active (running) since Mon 2013-08-05 14:57:49 CEST; 12s
```

3. Run the following command to view the postgresql processes:

```
~]$ ps -eZ | grep postgres
system_u:system_r:postgresql_t:s0 395 ? 00:00:00 postmaster
system_u:system_r:postgresql_t:s0 397 ? 00:00:00 postmaster
system_u:system_r:postgresql_t:s0 399 ? 00:00:00 postmaster
```

¹ See the *PostgreSQL* [http://www.postgresql.org/about/] project page for more information.

```
system_u:system_r:postgresql_t:s0 400 ? 00:00:00 postmaster system_u:system_r:postgresql_t:s0 401 ? 00:00:00 postmaster system_u:system_r:postgresql_t:s0 402 ? 00:00:00 postmaster
```

The SELinux context associated with the postgresql processes is system_u:system_r:postgresql_t:s0. The second last part of the context, postgresql_t, is the type. A type defines a domain for processes and a type for files. In this case, the postgresql processes are running in the postgresql_t domain.

21.2. Types

Type Enforcement is the main permission control used in SELinux targeted policy. All files and processes are labeled with a type: types define a domain for processes and a type for files. SELinux policy rules define how types access each other, whether it be a domain accessing a type, or a domain accessing another domain. Access is only allowed if a specific SELinux policy rule exists that allows it.

The following types are used with postgresql. Different types allow you to configure flexible access. Note that in the list below are used several regular expression to match the whole possible locations:

```
postgresql_db_t
```

This type is used for several locations. The locations labeled with this type are used for data files for PostgreSQL:

- /usr/lib/pgsql/test/regres
- /usr/share/jonas/pgsql
- /var/lib/pgsql/data
- /var/lib/postgres(ql)?

```
postgresql_etc_t
```

This type is used for configuration files in the /etc/postgresql/ directory.

```
postgresql_exec_t
```

This type is used for several locations. The locations labeled with this type are used for binaries for PostgreSQL:

- /usr/bin/initdb(.sepgsql)?
- /usr/bin/(se)?postgres
- /usr/lib(64)?/postgresql/bin/.*
- /usr/lib(64)?/pgsql/test/regress/pg_regress

```
systemd_unit_file_t
```

This type is used for the executable PostgreSQL-related files located in the /usr/lib/systemd/system/ directory.

```
postgresql_log_t
```

This type is used for several locations. The locations labeled with this type are used for log files:

- /var/lib/pgsql/logfile
- /var/lib/pgsql/pgstartup.log
- /var/lib/sepgsql/pgstartup.log
- /var/log/postgresql

- /var/log/postgres.log.*
- /var/log/rhdb/rhdb
- /var/log/sepostgresql.log.*

postgresql_var_run_t

This type is used for run-time files for PostgreSQL, such as the process id (PID) in the /var/run/postgresql/ directory.

21.3. Booleans

SELinux is based on the least level of access required for a service to run. Services can be run in a variety of ways; therefore, you need to specify how you run your services. Use the following Booleans to set up SELinux:

selinuxuser_postgresql_connect_enabled

Having this Boolean enabled allows any user domain (as defined by PostgreSQL) to make connections to the database server.



Note

Due to the continuous development of the SELinux policy, the list above might not contain all Booleans related to the service at all times. To list them, run the following command:

~]\$ getsebool -a | grep service_name

Run the following command to view description of a particular Boolean:

~]\$ sepolicy booleans -b boolean_name

Note that the additional *policycoreutils-devel* package providing the sepolicy utility is required.

21.4. Configuration Examples

21.4.1. PostgreSQL Changing Database Location

When using Fedora, the default location for PostgreSQL to store its database is /var/lib/pgsql/data/. This is where SELinux expects it to be by default, and hence this area is already labeled appropriately for you, using the postgresql_db_t type.

The area where the database is located can be changed depending on individual environment requirements or preferences, however it is important that SELinux is aware of this new location; that it is labeled accordingly. This example explains how to change the location of a PostgreSQL database and then how to label the new location so that SELinux can still provide its protection mechanisms to the new area based on its contents.

Note that this is an example only and demonstrates how SELinux can affect PostgreSQL. Comprehensive documentation of PostgreSQL is beyond the scope of this document. See the official

² http://www.postgresql.org/docs/

*PostgreSQL documentation*² for further details. This example assumes that the *postgresql-server* package is installed.

1. View the SELinux context of the default database location for postgresql:

```
~]# ls -lZ /var/lib/pgsql
drwx-----. postgres postgres system_u:object_r:postgresql_db_t:s0 data
```

This shows postgresq1_db_t which is the default context element for the location of database files. This context will have to be manually applied to the new database location that will be used in this example in order for it to function properly.

 Create a new directory for the new location of the database(s). In this example, /opt/ postgresql/data/ is used. If you use a different location, replace the text in the following steps with your location:

```
~]# mkdir -p /opt/postgresq1/data
```

Perform a directory listing of the new location. Note that the initial context of the new directory
is usr_t. This context is not sufficient for SELinux to offer its protection mechanisms to
PostgreSQL. Once the context has been changed, it will be able to function properly in the new
area.

```
~]# ls -lz /opt/postgresql/
drwxr-xr-x. root root unconfined_u:object_r:usr_t:s0 data
```

4. Change the ownership of the new location to allow access by the postgres user and group. This sets the traditional Unix permissions which SELinux will still observe.

```
~]# chown -R postgres:postgres /opt/postgresql
```

5. Open the PostgreSQL init file /etc/rc.d/init.d/postgresql with a text editor and modify the PGDATA and PGLOG variables to point to the new location:

```
~]# vi /etc/rc.d/init.d/postgresql
PGDATA=/opt/postgresql/data
PGLOG=/opt/postgresql/data/pgstartup.log
```

Save this file and exit the text editor.

6. Initialize the database in the new location:

```
~]$ su - postgres -c "initdb -D /opt/postgresql/data"
```

7. Having changed the database location, starting the service will fail at this point:

```
~]# systemctl start postgresql.service
```

```
Job for postgresql.service failed. See 'systemctl status postgresql.service' and 'journalctl -xn' for details.
```

SELinux has caused the service to not start. This is because the new location is not properly labeled. The following steps explain how to label the new location (/opt/postgresq1/) and start the postgresql service properly:

8. Use the semanage utility to add a context mapping for **/opt/postgresql/** and any other directories/files within it:

```
~]# semanage fcontext -a -t postgresql_db_t "/opt/postgresql(/.*)?"
```

9. This mapping is written to the /etc/selinux/targeted/contexts/files/ file_contexts.local file:

```
~]# grep -i postgresql /etc/selinux/targeted/contexts/files/file_contexts.local
/opt/postgresql(/.*)? system_u:object_r:postgresql_db_t:s0
```

10. Now use the restorecon utility to apply this context mapping to the running system:

```
~]# restorecon -R -v /opt/postgresql
```

11. Now that the **/opt/postgresql/** location has been labeled with the correct context for PostgreSQL, the postgresql service will start successfully:

```
~]# systemctl start postgresql.service
```

12. Confirm the context is correct for /opt/postgresql/:

```
~]$ ls -lz /opt
drwxr-xr-x. root root system_u:object_r:postgresql_db_t:s0 postgresql
```

13. Check with the **ps** command that the postgresql process displays the new location:

```
~]# ps aux | grep -i postmaster

postgres 21564 0.3 0.3 42308 4032 ? S 10:13 0:00 /usr/bin/postmaster -p 5432 -D /opt/postgresql/data/
```

14. The location has been changed and labeled, and postgresql has started successfully. At this point all running services should be tested to confirm normal operation.

rsync

The rsync utility performs fast file transfer and it is used for synchronizing data between systems. 1

When using Fedora, the *rsync* package provides rsync. Run the following command to see if the *rsync* package is installed:

```
~]$ rpm -q rsync
package rsync is not installed
```

If it is not installed, use the DNF utility as root to install it:

```
~]# dnf install rsync
```

22.1. rsync and SELinux

SELinux requires files to have an extended attribute to define the file type. Policy governs the access daemons have to these files. If you want to share files using the rsync daemon, you must label the files and directories with the public_content_t type. Like most services, correct labeling is required for SELinux to perform its protection mechanisms over rsync.²

22.2. Types

Type Enforcement is the main permission control used in SELinux targeted policy. All files and processes are labeled with a type: types define a domain for processes and a type for files. SELinux policy rules define how types access each other, whether it be a domain accessing a type, or a domain accessing another domain. Access is only allowed if a specific SELinux policy rule exists that allows it.

The following types are used with rsync. Different types all you to configure flexible access:

```
public_content_t
```

This is a generic type used for the location of files (and the actual files) to be shared using rsync. If a special directory is created to house files to be shared with rsync, the directory and its contents need to have this label applied to them.

```
rsync_exec_t
```

This type is used for the /usr/bin/rsync system binary.

```
rsync_log_t
```

This type is used for the rsync log file, located at /var/log/rsync.log by default. To change the location of the file rsync logs to, use the --log-file=FILE option to the rsync command at run-time.

```
rsync_var_run_t
```

This type is used for the rsyncd lock file, located at /var/run/rsyncd.lock. This lock file is used by the rsync server to manage connection limits.

¹ See the *Rsync* [http://www.samba.org/rsync/] project page for more information.

² See the rsync_selinux(8) manual page for more information about rsync and SELinux.

rsync_data_t

This type is used for files and directories which you want to use as rsync domains and isolate them from the access scope of other services. Also, the public_content_t is a general SELinux context type, which can be used when a file or a directory interacts with multiple services (for example, FTP and NFS directory as an rsync domain).

rsync_etc_t

This type is used for rsync-related files in the /etc/ directory.

22.3. Booleans

SELinux is based on the least level of access required for a service to run. Services can be run in a variety of ways; therefore, you need to specify how you run your services. Use the following Booleans to set up SELinux:

rsync_anon_write

Having this Boolean enabled allows rsync in the rsync_t domain to manage files, links and directories that have a type of public_content_rw_t. Often these are public files used for public file transfer services. Files and directories must be labeled this type.

rsync_client

Having this Boolean enabled allows rsync to initiate connections to ports defined as rsync_port_t, as well as allowing the daemon to manage files, links, and directories that have a type of rsync_data_t. Note that rsync must be in the rsync_t domain in order for SELinux to enact its control over it. The configuration example in this chapter demonstrates rsync running in the rsync_t domain.

rsync_export_all_ro

Having this Boolean enabled allows rsync in the rsync_t domain to export NFS and CIFS volumes with read-only access to clients.



Note

Due to the continuous development of the SELinux policy, the list above might not contain all Booleans related to the service at all times. To list them, run the following command:

~]\$ getsebool -a | grep service_name

Run the following command to view description of a particular Boolean:

~]\$ sepolicy booleans -b boolean_name

Note that the additional *policycoreutils-devel* package providing the sepolicy utility is required.

22.4. Configuration Examples

22.4.1. Rsync as a daemon

When using Fedora, rsync can be used as a daemon so that multiple clients can directly communicate with it as a central server, in order to house centralized files and keep them synchronized. The following example will demonstrate running rsync as a daemon over a network socket in the correct

domain, and how SELinux expects this daemon to be running on a pre-defined (in SELinux policy) TCP port. This example will then show how to modify SELinux policy to allow the rsync daemon to run normally on a non-standard port.

This example will be performed on a single system to demonstrate SELinux policy and its control over local daemons and processes. Note that this is an example only and demonstrates how SELinux can affect rsync. Comprehensive documentation of rsync is beyond the scope of this document. See the official *rsync documentation*³ for further details. This example assumes that the *rsync*, *setroubleshootserver* and *audit* packages are installed, that the SELinux targeted policy is used and that SELinux is running in enforcing mode.

Procedure 22.1. Getting rsync to launch as rsync t

1. Run the **getenforce** command to confirm SELinux is running in enforcing mode:

```
~]$ getenforce
Enforcing
```

The command returns **Enforcing** when SELinux is running in enforcing mode.

2. Run the **which** command to confirm that the rsync binary is in the system path:

```
~]$ which rsync
/usr/bin/rsync
```

3. When running rsync as a daemon, a configuration file should be used and saved as /etc/rsyncd.conf. Note that the following configuration file used in this example is very simple and is not indicative of all the possible options that are available, rather it is just enough to demonstrate the rsync daemon:

4. Now that a simple configuration file exists for rsync to operate in daemon mode, you can start it by running the following command:

```
~]# systemctl start rsyncd.service
```

Ensure that rsyncd was successfully started (the output is supposed to look similar to the one below, only the time stamp will differ):

```
~]# systemctl status rsyncd.service
rsyncd.service - fast remote file copy program daemon
```

³ http://www.samba.org/rsync/documentation.html

```
Loaded: loaded (/usr/lib/systemd/system/rsyncd.service; disabled)
Active: active (running) since Thu 2014-02-27 09:46:24 CET; 2s ago
Main PID: 3220 (rsync)
CGroup: /system.slice/rsyncd.service

—3220 /usr/bin/rsync --daemon --no-detach
```

SELinux can now enforce its protection mechanisms over the rsync daemon as it is now running in the rsync_t domain:

This example demonstrated how to get rsyncd running in the rsync_t domain. Rsync can also be run as a socket-activated service. In that case, the rsyncd is not executed until a client tries to connect to the service. To enable rsyncd to run as a socket-activated service, follow the steps above. To start rsyncd as a socket-activated service, run the following command as root:

```
~]# systemctl start rsyncd.socket
```

The next example shows how to get this daemon successfully running on a non-default port. TCP port 10000 is used in the next example.

Procedure 22.2. Running the rsync daemon on a non-default port

 Modify the /etc/rsyncd.conf file and add the port = 10000 line at the top of the file in the global configuration area (that is, before any file areas are defined). The new configuration file will look like:

2. After launching the rsync daemon with this new setting, a denial message similar to the following is logged by SELinux:

```
Jul 22 10:46:59 localhost setroubleshoot: SELinux is preventing the rsync (rsync_t) from binding to port 10000. For complete SELinux messages, run sealert -l c371ab34-639e-45ae-9e42-18855b5c2de8
```

Use the semanage utility to add TCP port 10000 to the SELinux policy in rsync_port_t:

```
~]# semanage port -a -t rsync_port_t -p tcp 10000
```

4. Now that TCP port 10000 has been added to the SELinux policy for rsync_port_t, rsyncd will start and operate normally on this port:

```
~]# systemctl start rsyncd.service

-]# netstat -lnp | grep 10000
tcp 0 00.0.0.0:10000 0.0.0.0:* LISTEN 9910/rsync
```

SELinux has had its policy modified and is now permitting rsyncd to operate on TCP port 10000.

Postfix

Postfix is an open-source Mail Transport Agent (MTA), which supports protocols like LDAP, SMTP AUTH (SASL), and TLS.¹

In Fedora, the *postfix* package provides Postfix. Run the following command to see if the *postfix* package is installed:

```
~]$ rpm -q postfix
package postfix is not installed
```

If it is not installed, use the DNF utility root to install it:

```
~]# dnf install postfix
```

23.1. Postfix and SELinux

When Postfix is enabled, it runs confined by default. Confined processes run in their own domains, and are separated from other confined processes. If a confined process is compromised by an attacker, depending on SELinux policy configuration, an attacker's access to resources and the possible damage they can do is limited. The following example demonstrates the Postfix and related processes running in their own domain. This example assumes the *postfix* package is installed and that the Postfix service has been started:

1. Run the **getenforce** command to confirm SELinux is running in enforcing mode:

```
~]$ getenforce
Enforcing
```

The command returns **Enforcing** when SELinux is running in enforcing mode.

2. Run the following command as the root user to start postfix:

```
~]# systemctl start postfix.service
```

Confirm that the service is running. The output should include the information below (only the time stamp will differ):

```
~]# systemctl status postfix.service
postfix.service - Postfix Mail Transport Agent
  Loaded: loaded (/usr/lib/systemd/system/postfix.service; disabled)
  Active: active (running) since Mon 2013-08-05 11:38:48 CEST; 3h 25min ago
```

Run following command to view the postfix processes:

```
~]$ ps -eZ | grep postfix
```

¹ Refer to the *Postfix* [http://www.postfix.org/] project page for more information.

```
system_u:system_r:postfix_master_t:s0 1651 ? 00:00:00 master
system_u:system_r:postfix_pickup_t:s0 1662 ? 00:00:00 pickup
system_u:system_r:postfix_qmgr_t:s0 1663 ? 00:00:00 qmgr
```

In the output above, the SELinux context associated with the Postfix master process is system_u:system_r:postfix_master_t:s0. The second last part of the context, postfix_master_t, is the type for this process. A type defines a domain for processes and a type for files. In this case, the master process is running in the postfix_master_t domain.

23.2. Types

Type Enforcement is the main permission control used in SELinux targeted policy. All files and processes are labeled with a type: types define a domain for processes and a type for files. SELinux policy rules define how types access each other, whether it be a domain accessing a type, or a domain accessing another domain. Access is only allowed if a specific SELinux policy rule exists that allows it.

The following types are used with Postfix. Different types all you to configure flexible access:

```
postfix_etc_t
```

This type is used for configuration files for Postfix in the /etc/postfix/ directory.

```
postfix_data_t
```

This type is used for Postfix data files in the /var/lib/postfix/ directory.

```
postfix_var_run_t
```

This type is used for Postfix files stored in the /run/ directory.

```
postfix_initrc_exec_t
```

The Postfix executable files are labeled with the postfix_initrc_exec_t type. When executed, they transition to the postfix_initrc_t domain.

```
postfix_spool_t
```

This type is used for Postfix files stored in the /var/spool/ directory.



Note

To see the full list of files and their types for Postfix, run the following command:

~]\$ grep postfix /etc/selinux/targeted/contexts/files/file_contexts

23.3. Booleans

SELinux is based on the least level of access required for a service to run. Services can be run in a variety of ways; therefore, you need to specify how you run your services. Use the following Booleans to set up SELinux:

```
postfix_local_write_mail_spool
```

Having this Boolean enabled allows Postfix to write to the local mail spool on the system. Postfix requires this Boolean to be enabled for normal operation when local spools are used.



Note

Due to the continuous development of the SELinux policy, the list above might not contain all Booleans related to the service at all times. To list them, run the following command:

```
~]$ getsebool -a | grep service_name
```

Run the following command to view description of a particular Boolean:

```
~]$ sepolicy booleans -b boolean_name
```

Note that the additional *policycoreutils-devel* package providing the sepolicy utility is required.

23.4. Configuration Examples

23.4.1. SpamAssassin and Postfix

SpamAssasin is an open-source mail filter that provides a way to filter unsolicited email (spam messages) from incoming email.²

When using Fedora, the *spamassassin* package provides SpamAssassin. Run the following command to see if the *spamassassin* package is installed:

```
~]$ rpm -q spamassassin
package spamassassin is not installed
```

If it is not installed, use the DNF utility as root to install it:

```
~]# dnf install spamassassin
```

SpamAssassin operates in tandem with a mailer such as Postfix to provide spam-filtering capabilities. In order for SpamAssassin to effectively intercept, analyze and filter mail, it must listen on a network interface. The default port for SpamAssassin is TCP/783, however this can be changed. The following example provides a real-world demonstration of how SELinux complements SpamAssassin by only allowing it access to a certain port by default. This example will then demonstrate how to change the port and have SpamAssassin operate on a non-default port.

Note that this is an example only and demonstrates how SELinux can affect a simple configuration of SpamAssassin. Comprehensive documentation of SpamAssassin is beyond the scope of this document. See the official *SpamAssassin documentation*³ for further details. This example assumes the *spamassassin* is installed, that any firewall has been configured to allow access on the ports in use, that the SELinux targeted policy is used, and that SELinux is running in enforcing mode:

² Refer to the *SpamAssassin* [http://spamassassin.apache.org/] project page for more information.

³ http://spamassassin.apache.org/doc.html

Procedure 23.1. Running SpamAssassin on a non-default port

1. Use the semanage utility as root to show the port that SELinux allows the spamd daemon to listen on by default:

```
~]# semanage port -1 | grep spamd
spamd_port_t tcp 783
```

This output shows that TCP/783 is defined in spamd_port_t as the port for SpamAssassin to operate on.

2. Edit the /etc/sysconfig/spamassassin configuration file and modify it so that it will start SpamAssassin on the example port TCP/10000:

```
# Options to spamd
SPAMDOPTIONS="-d -p 10000 -c m5 -H"
```

This line now specifies that SpamAssassin will operate on port 10000. The rest of this example will show how to modify the SELinux policy to allow this socket to be opened.

3. Start SpamAssassin and an error message similar to the following will appear:

```
~]# systemctl start spamassassin.service
Job for spamassassin.service failed. See 'systemctl status spamassassin.service' and
'journalctl -xn' for details.
```

This output means that SELinux has blocked access to this port.

4. A denial message similar to the following will be logged by SELinux:

```
SELinux is preventing the spamd (spamd_t) from binding to port 10000.
```

5. As root, run semanage to modify the SELinux policy in order to allow SpamAssassin to operate on the example port (TCP/10000):

```
~]# semanage port -a -t spamd_port_t -p tcp 10000
```

6. Confirm that SpamAssassin will now start and is operating on TCP port 10000:

```
~]# systemctl start spamassassin.service

~]# netstat -lnp | grep 10000
tcp 0 0 127.0.0.1:10000 0.0.0.0:* LISTEN 2224/spamd.pid
```

7. At this point, spamd is properly operating on TCP port 10000 as it has been allowed access to that port by the SELinux policy.

DHCP

The dhcpd daemon is used in Fedora to dynamically deliver and configure Layer 3 TCP/IP details for clients.

The *dhcp* package provides the DHCP server and the dhcpd daemon. Run the following command to see if the *dhcp* package is installed:

```
~]# rpm -q dhcp
package dhcp is not installed
```

If it is not installed, use the DNF utility as root to install it:

```
~]# dnf install dhcp
```

24.1. DHCP and SELinux

When dhcpd is enabled, it runs confined by default. Confined processes run in their own domains, and are separated from other confined processes. If a confined process is compromised by an attacker, depending on SELinux policy configuration, an attacker's access to resources and the possible damage they can do is limited. The following example demonstrates dhcpd and related processes running in their own domain. This example assumes the *dhcp* package is installed and that the dhcpd service has been started:

1. Run the **getenforce** command to confirm SELinux is running in enforcing mode:

```
~]$ getenforce
Enforcing
```

The command returns **Enforcing** when SELinux is running in enforcing mode.

2. Run the following command as the root user to start dhcpd:

```
~]# systemctl start dhcpd.service
```

Confirm that the service is running. The output should include the information below (only the time stamp will differ):

```
~]# systemctl status dhcpd.service
dhcpd.service - DHCPv4 Server Daemon
Loaded: loaded (/usr/lib/systemd/system/dhcpd.service; disabled)
Active: active (running) since Mon 2013-08-05 11:49:07 CEST; 3h 20min ago
```

3. Run following command to view the dhcpd processes:

```
~]$ ps -eZ | grep dhcpd
system_u:system_r:dhcpd_t:s0 5483 ? 00:00:00 dhcpd
```

The SELinux context associated with the dhcpd process is $system_u:system_r:dhcpd_t:s0$.

24.2. Types

The following types are used with DHCP:

dhcp_etc_t

This type is mainly used for files in the /etc/ directory, including configuration files.

dhcpd_var_run_t

This type is used for the PID file for dhcpd, in the /var/run/ directory.

dhcpd_exec_t

This type is used for transition of DHCP executable files to the dhcpd_t domain.

dhcpd_initrc_exec_t

This type is used for transition of DHCP executable files to the dhcpd_initrc_t domain.



Note

To see the full list of files and their types for dhcpd, run the following command:

~]\$ grep dhcp /etc/selinux/targeted/contexts/files/file_contexts

References

The following references are pointers to additional information that is relevant to SELinux but beyond the scope of this guide. Note that due to the rapid development of SELinux, some of this material may only apply to specific releases of Fedora.

Books

SELinux by Example

Mayer, MacMillan, and Caplan

Prentice Hall, 2007

SELinux: NSA's Open Source Security Enhanced Linux

Bill McCarty

O'Reilly Media Inc., 2004

Tutorials and Help

Tutorials and talks from Russell Coker

http://www.coker.com.au/selinux/talks/ibmtu-2004/

Dan Walsh's Journal

http://danwalsh.livejournal.com/

Red Hat Knowledgebase

https://access.redhat.com/site/

General Information

NSA SELinux main website

http://www.nsa.gov/research/selinux/index.shtml

NSA SELinux FAQ

http://www.nsa.gov/research/selinux/faqs.shtml

Mailing Lists

NSA SELinux mailing list

http://www.nsa.gov/research/selinux/list.shtml

Fedora SELinux mailing list

http://www.redhat.com/mailman/listinfo/fedora-selinux-list

Community

SELinux Project Wiki

http://selinuxproject.org/page/Main_Page

SELinux community page

http://selinux.sourceforge.net/

IRC

irc.freenode.net, #selinux

Appendix A. Revision History

Revision 3.0-0 Mon July 11 2016 Fedora 24 release of the book.

Mirek Jahoda mirek. jahoda@redhat.com

Revision 2.0-0 Mon May 25 2015

Barbora Ančincová bancinco@redhat.com

Fedora 22 release of the book.

Barbora Ančincová bancinco@redhat.com

Revision 1.0-0 Mon Dec 8 2014 Fedora 21 release of the book.

Revision 0.1-1 Sun Oct 12 2014
Initial creation of the book for Fedora.

Barbora Ančincová bancinco@redhat.com