

# Linux “Top” Command Output Details

The top command allows users to monitor processes and system resource usage on Linux. It is one of the most useful tools in a sysadmin’s toolbox, and it comes pre-installed on every distribution. Unlike other commands such as `ps`, it is interactive, and you can browse through the list of processes, kill a process, and so on.

In this article, we are going to understand how to use the top command.

## Getting started

As you might have already guessed, you simply need to type this in to launch top:

`top`

This starts up an interactive command line application, similar to one in the screenshot below. The upper half of the output contains statistics on processes and resource usage, while the lower half contains a list of the currently running processes. You can use the arrow keys and Page Up/Down keys to browse through the list. If you want to quit, simply press “q”.

```
top - 15:39:37 up 90 days, 15:26, 2 users, load average: 0.00, 0.00, 0.00
Tasks: 27 total, 1 running, 26 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 524288 total, 22792 free, 119380 used, 382116 buff/cache
KiB Swap: 131072 total, 43716 free, 87356 used. 322002 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	37168	1192	680	S	0.0	0.2	2:21.51	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd/646
3	root	20	0	0	0	0	S	0.0	0.0	0:00.01	khelper/646
62	root	20	0	38896	1316	1188	S	0.0	0.3	1:08.85	systemd-journal
219	root	20	0	26012	328	200	S	0.0	0.1	0:11.09	cron
226	root	20	0	65464	924	220	S	0.0	0.2	0:13.11	sshd
229	syslog	20	0	184632	1524	464	S	0.0	0.3	0:28.85	rsyslogd
231	root	20	0	47572	504	40	S	0.0	0.1	0:07.80	rpcbind
274	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	agetty
275	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	agetty
293	root	20	0	308984	12800	2248	S	0.0	2.4	27:15.96	fail2ban-server
4452	root	20	0	92996	3124	3120	S	0.0	0.6	0:00.03	sshd
4461	supriyo	20	0	92996	1000	996	S	0.0	0.2	0:00.00	sshd
4462	supriyo	20	0	19472	1604	1600	S	0.0	0.3	0:00.05	bash
4696	root	20	0	92996	4036	3132	S	0.0	0.8	0:00.02	sshd
4705	supriyo	20	0	92996	1952	1008	S	0.0	0.4	0:00.02	sshd
4706	supriyo	20	0	19472	3364	1600	S	0.0	0.6	0:00.05	bash
4718	supriyo	20	0	36608	1784	1324	R	0.0	0.3	0:00.31	top
5830	root	20	0	41532	728	320	S	0.0	0.1	0:01.25	systemd-udev
13879	www-data	20	0	290032	2632	2612	S	0.0	0.5	0:01.18	php-fpm7.0
14031	cloud-t+	20	0	19788	9736	3276	S	0.0	1.9	10:11.27	cloud-torrent
14089	root	20	0	286060	560	452	S	0.0	0.1	1:11.67	php-fpm7.0
14091	www-data	20	0	289508	2168	2064	S	0.0	0.4	0:02.32	php-fpm7.0

There are a number of variants of top, but in the rest of this article, we will talk about the most common variant — the one that comes with the “procps-ng” package. You can verify this by running:

```
top -v
```

If you have this variant, this will show up in the output, like so:

```
procps-ng version 3.3.10
```

There is quite a bit going on in top’s interface, so we will break it down bit by bit in the next section.

## Understanding top’s interface: the summary area

As we have previously seen, top’s output is divided into two different sections. In this part of the article, we’re going to focus on the elements in half of the output. This region is also called the “summary area”.

### System time, uptime and user sessions

```
top - 15:39:37 up 90 days, 15:26, 2 users, load average: 0.00, 0.00, 0.00
Tasks: 27 total, 1 running, 26 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 524288 total, 22792 free, 119380 used, 382116 buff/cache
KiB Swap: 131072 total, 43716 free, 87356 used. 322002 avail Mem
```

At the very top left of the screen (as marked in the screenshot above), top displays the current time. This is followed by the system uptime, which tells us the time for which the system has been running. For instance, in our example, the current time is “15:39:37”, and the system has been running for 90 days, 15 hours and 26 minutes.

Next comes the number of active user sessions. In this example, there are two active user sessions. These sessions may be either made on a [TTY](#) (physically on the system, either through the command line or a desktop environment) or a [PTY](#) (such as a terminal emulator window or over SSH). In fact, if you log in to a Linux system through a desktop environment, and then start a terminal emulator, you will find there will be two active sessions.

If you want to get more details about the active user sessions, use the `who` command.

### Memory usage

```
top - 15:39:37 up 90 days, 15:26, 2 users, load average: 0.00, 0.00, 0.00
Tasks: 27 total, 1 running, 26 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 524288 total, 22792 free, 119380 used, 382116 buff/cache
KiB Swap: 131072 total, 43716 free, 87356 used. 322002 avail Mem
```

The “memory” section shows information regarding the memory usage of the system. The lines marked “Mem” and “Swap” show information about RAM and swap space respectively. Simply put, a swap space is a part of the hard disk that is used like RAM. When the RAM usage gets nearly full, infrequently used regions of the RAM are written into the swap space, ready to be retrieved later when needed. However, because accessing disks are slow, relying too much on swapping can harm system performance.

As you would naturally expect, the “total”, “free” and “used” values have their usual meanings. The “avail mem” value is the amount of memory that can be allocated to processes without causing more swapping.

The Linux kernel also tries to reduce disk access times in various ways. It maintains a “disk cache” in RAM, where frequently used regions of the disk are stored. In addition, disk writes are stored to a “disk buffer”, and the kernel eventually writes them out to the disk. The total memory consumed by them is the “buff/cache” value. It might sound like a bad thing, but it really isn’t — memory used by the cache will be allocated to processes if needed.

## Tasks

```
top - 15:39:37 up 90 days, 15:26, 2 users, load average: 0.00, 0.00, 0.00
Tasks: 27 total, 1 running, 26 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 524288 total, 22792 free, 119380 used, 382116 buff/cache
KiB Swap: 131072 total, 43716 free, 87356 used. 322002 avail Mem
```

The “Tasks” section shows statistics regarding the processes running on your system. The “total” value is simply the total number of processes. For example, in the above screenshot, there are 27 processes running. To understand the rest of the values, we need a little bit of background on how the Linux kernel handles processes.

Processes perform a mix of I/O-bound work (such as reading disks) and CPU-bound work (such as performing arithmetic operations). The CPU is idle when a process performs I/O, so OSes switch to executing other processes during this time. In addition, the OS allows a given process to execute for a very small amount of time, and then it switches over to another process. This is how OSes appear as if they were “multitasking”. Doing all this requires us to keep track of the “state” of a process. In Linux, a process may be in of these states:

- **Runnable (R):** A process in this state is either executing on the CPU, or it is present on the run queue, ready to be executed.
- **Interruptible sleep (S):** Processes in this state are waiting for an event to complete.
- **Uninterruptible sleep (D):** In this case, a process is waiting for an I/O operation to complete.

- Stopped (T): These processes have been stopped by a [job control signal](#) (such as by pressing Ctrl+Z) or because they are being traced.
- Zombie (Z): The kernel maintains various data structures in memory to keep track of processes. A process may create a number of child processes, and they may exit while the parent is still around. However, these data structures must be kept around until the parent obtains the status of the child processes. Such terminated processes whose data structures are still around are called zombies.

Processes in the D and S states are shown in “sleeping”, and those in the T state are shown in “stopped”. The number of zombies are shown as the “zombie” value.

## CPU usage

```
top - 15:39:37 up 90 days, 15:26, 2 users, load average: 0.00, 0.00, 0.00
Tasks: 27 total, 1 running, 26 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 524288 total, 22792 free, 119380 used, 382116 buff/cache
KiB Swap: 131072 total, 43716 free, 87356 used. 322002 avail Mem
```

The CPU usage section shows the percentage of CPU time spent on various tasks. The `us` value is the time the CPU spends executing processes in user space. Similarly, the `sy` value is the time spent on running kernel space processes.

Linux uses a “nice” value to determine the priority of a process. A process with a high “nice” value is “nicer” to other processes, and gets a low priority. Similarly, processes with a lower “nice” gets higher priority. As we shall see later, the default “nice” value can be changed. The time spent on executing processes with a manually set “nice” appear as the `ni` value.

This is followed by `id`, which is the time the CPU remains idle. Most operating systems put the CPU on a power saving mode when it is idle. Next comes the `wa` value, which is the time the CPU spends waiting for I/O to complete.

[Interrupts](#) are signals to the processor about an event that requires immediate attention. Hardware interrupts are typically used by peripherals to tell the system about events, such as a keypress on a keyboard. On the other hand, software interrupts are generated due to specific instructions executed on the processor. In either case, the OS handles them, and the time spent on handling hardware and software interrupts are given by `hi` and `si` respectively.

In a virtualized environment, a part of the CPU resources are given to each virtual machine (VM). The OS detects when it has work to do, but it cannot perform them because the CPU is busy on some other VM. The amount of time lost in this way is the “steal” time, shown as `st`.

## Load average

```
top - 15:39:37 up 90 days, 15:26, 2 users, load average: 0.00, 0.00, 0.00
Tasks: 27 total, 1 running, 26 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni, 100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 524288 total, 22792 free, 119380 used, 382116 buff/cache
KiB Swap: 131072 total, 43716 free, 87356 used. 322002 avail Mem
```

The load average section represents the average “load” over one, five and fifteen minutes. “Load” is a measure of the amount of computational work a system performs. On Linux, the load is the number of processes in the R and D states at any given moment. The “load average” value gives you a relative measure of how long you must wait for things to get done.

Let us consider a few examples to understand this concept. On a single core system, a load average of 0.4 means the system is doing only 40% of work it can do. A load average of 1 means that the system is exactly at capacity — the system will be overloaded by adding even a little bit of additional work. A system with a load average of 2.12 means that it is overloaded by 112% more work than it can handle.

On a multi-core system, you should first divide the load average with the number of CPU cores to get a similar measure.

In addition, “load average” isn’t actually the typical “average” most of us know. It is an “[exponential moving average](#)“, which means a small part of the previous load averages are factored into the current value. If you’re interested, [this](#) article covers all the technical details.

## Understanding top’s interface: the task area

The summary area is comparatively simpler, and it contains a list of processes. In this section, we will learn about the different columns shown in top’s default output.

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	37168	1192	680	S	0.0	0.2	2:21.51	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd/646
3	root	20	0	0	0	0	S	0.0	0.0	0:00.01	khelper/646
62	root	20	0	38896	1316	1188	S	0.0	0.3	1:08.85	systemd-journal
219	root	20	0	26012	328	200	S	0.0	0.1	0:11.09	cron
226	root	20	0	65464	924	220	S	0.0	0.2	0:13.11	sshd
229	syslog	20	0	184632	1524	464	S	0.0	0.3	0:28.85	rsyslogd
231	root	20	0	47572	504	40	S	0.0	0.1	0:07.80	rpcbind
274	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	agetty
275	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	agetty
293	root	20	0	308984	12800	2248	S	0.0	2.4	27:15.96	fail2ban-server
4452	root	20	0	92996	3124	3120	S	0.0	0.6	0:00.03	sshd
4461	supriyo	20	0	92996	1000	996	S	0.0	0.2	0:00.00	sshd
4462	supriyo	20	0	19472	1604	1600	S	0.0	0.3	0:00.05	bash
4696	root	20	0	92996	4036	3132	S	0.0	0.8	0:00.02	sshd
4705	supriyo	20	0	92996	1952	1008	S	0.0	0.4	0:00.02	sshd
4706	supriyo	20	0	19472	3364	1600	S	0.0	0.6	0:00.05	bash
4718	supriyo	20	0	36608	1784	1324	R	0.0	0.3	0:00.31	top
5830	root	20	0	41532	728	320	S	0.0	0.1	0:01.25	systemd-udevd
13879	www-data	20	0	290032	2632	2612	S	0.0	0.5	0:01.18	php-fpm7.0
14031	cloud-t+	20	0	19788	9736	3276	S	0.0	1.9	10:11.27	cloud-torrent
14089	root	20	0	286060	560	452	S	0.0	0.1	1:11.67	php-fpm7.0
14091	www-data	20	0	289508	2168	2064	S	0.0	0.4	0:02.32	php-fpm7.0

- **PID**

This is the process ID, a unique positive integer that identifies a process.

- **USER**

This is the “effective” username (which maps to a user ID) of the user who started the process. Linux assigns a real user ID and an effective user ID to processes; the latter allows a process to act on behalf of another user. (For example, a non-root user can elevate to root in order to install a package.)

- **PR and NI**

The “NI” field shows the “nice” value of a process. The “PR” field shows the scheduling priority of the process from the perspective of the kernel. The nice value affects the priority of a process.

- **VIRT, RES, SHR and %MEM**

These three fields are related with to memory consumption of the processes. “VIRT” is the total amount of memory consumed by a process. This includes the program’s code, the data stored by the process in memory, as well as any regions of memory that have been swapped to the disk. “RES” is the memory consumed by the process in RAM, and “%MEM” expresses this value as a percentage of the total RAM available. Finally, “SHR” is the amount of memory shared with other processes.

- **S**

As we have seen before, a process may be in various states. This field shows the process state in the single-letter form.

- **TIME+**

This is the total CPU time used by the process since it started, precise to the hundredths of a second.

- **COMMAND**

The COMMAND column shows the name of the processes.

## **Top command usage examples**

So far, we have discussed about top's interface. However, it can also manage processes, and you can control various aspects of top's output. In this section, we're going to take at a few examples.

In most of the examples below, you have to press a key while top is running. Keep in mind that these keypresses are case sensitive — so if you press “k” while Caps Lock is on, you have actually pressed a “K”, and the command won't work, or do something else entirely.

### **Killing processes**

If you want to kill a process, simply press ‘k’ when top is running. This will bring up a prompt, which will ask for the process ID of the process and press enter.



```
top - 19:05:02 up 2 days, 22:59, 2 users, load average: 0.00, 0.00, 0.00
Tasks: 27 total, 1 running, 26 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 524288 total, 336608 free, 104540 used, 83140 buff/cache
KiB Swap: 131072 total, 84420 free, 46652 used. 339668 avail Mem
PID to signal/kill [default pid = 69] 23706
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
69	root	20	0	38896	3068	2932	S	0.0	0.6	0:01.37	systemd-journal
72	root	20	0	41632	136	88	S	0.0	0.0	0:00.16	systemd-udevd
213	root	20	0	47572	104	40	S	0.0	0.0	0:00.24	rpcbind
219	root	20	0	26012	300	196	S	0.0	0.1	0:00.35	cron
230	root	20	0	65464	332	220	S	0.0	0.1	0:00.23	sshd
231	syslog	20	0	184632	884	584	S	0.0	0.2	0:00.61	rsyslogd
282	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	agetty
286	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	agetty
294	root	20	0	126072	304	200	S	0.0	0.1	0:00.00	nginx
295	www-data	20	0	126868	2412	1712	S	0.0	0.5	1:05.38	nginx
296	www-data	20	0	126916	2540	1812	S	0.0	0.5	1:20.61	nginx
299	root	20	0	286060	524	396	S	0.0	0.1	0:08.50	php-fpm7.0
301	www-data	20	0	286548	1580	988	S	0.0	0.3	0:00.06	php-fpm7.0
302	www-data	20	0	288728	4620	2336	S	0.0	0.9	0:00.29	php-fpm7.0
307	root	20	0	300484	10096	2480	S	0.0	1.9	1:00.34	fail2ban-server
23659	root	20	0	92996	4040	3132	S	0.0	0.8	0:00.02	sshd
23668	supriyo	20	0	92996	2032	1012	S	0.0	0.4	0:00.03	sshd
23669	supriyo	20	0	19472	3368	1600	S	0.0	0.6	0:00.04	bash
23690	supriyo	20	0	17868	1308	996	S	0.0	0.2	0:00.00	tmux
23692	supriyo	20	0	26448	1776	1264	S	0.0	0.3	0:00.14	tmux
23693	supriyo	20	0	19480	3368	1596	S	0.0	0.6	0:00.05	bash
23706	supriyo	20	0	4168	412	328	S	0.0	0.1	0:00.00	test

Next, enter the signal using which the process should be killed. If you leave this blank, top uses a SIGTERM, which allows processes to terminate gracefully. If you want to kill a process forcefully, you can type in SIGKILL here. You can also type in the signal number here. For example, the number for SIGTERM is 15 and SIGKILL is 9.

If you leave the process ID blank and hit enter directly, it will terminate the topmost process in the list. As we've mentioned previously, you can scroll using the arrow keys, and change the process you want to kill in this way.

## Sorting the process list

One of the most frequent reasons to use a tool like top is to find out which process is consuming the most resources. You can press the following keys to sort the list:

- 'M' to sort by memory usage
- 'P' to sort by CPU usage
- 'N' to sort by process ID
- 'T' to sort by the running time

By default, top displays all results in descending order. However, you can switch to ascending order by pressing 'R'.



You can also sort the list with the `-o` switch. For example, if you want to sort processes by CPU usage, you can do so with:

```
top -o %CPU
```

You can sort the list by any of the attributes in the summary area in the same way.

## Showing a list of threads instead of processes

We have previously touched upon how Linux switches between processes. Unfortunately, processes do not share memory or other resources, making such switches rather slow. Linux, like many other operating systems, supports a “lightweight” alternative, called a “thread”. They are part of a process and share certain regions of memory and other resources, but they can be run concurrently like processes.

By default, `top` shows a list of processes in its output. If you want to list the threads instead, press ‘H’ when `top` is running. Notice that the “Tasks” line says “Threads” instead, and shows the number of threads instead of processes.

```
top - 20:24:20 up 3 days, 18 min, 2 users, load average: 0.00, 0.00, 0.00
Threads: 59 total, 1 running, 58 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.1 us, 0.1 sy, 0.0 ni, 99.8 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 524288 total, 333112 free, 104884 used, 86292 buff/cache
KiB Swap: 131072 total, 84460 free, 46612 used. 339308 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	37372	2036	1232	S	0.0	0.4	0:04.58	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd/646
3	root	20	0	0	0	0	S	0.0	0.0	0:00.00	khelper/646
69	root	20	0	38896	3108	2972	S	0.0	0.6	0:01.40	systemd-journal
72	root	20	0	41632	136	88	S	0.0	0.0	0:00.17	systemd-udevd
213	root	20	0	47572	104	40	S	0.0	0.0	0:00.25	rpcbind
219	root	20	0	26012	300	196	S	0.0	0.1	0:00.36	cron
230	root	20	0	65464	332	220	S	0.0	0.1	0:00.23	sshd
231	syslog	20	0	184632	884	584	S	0.0	0.2	0:00.00	rsyslogd
272	syslog	20	0	184632	884	584	S	0.0	0.2	0:00.17	in:imuxsock
273	syslog	20	0	184632	884	584	S	0.0	0.2	0:00.42	rs:main Q:Reg
237	cloud-t+	20	0	135796	14360	4032	S	0.0	2.7	0:10.78	cloud-torrent
287	cloud-t+	20	0	135796	14360	4032	S	0.0	2.7	1:16.38	cloud-torrent
288	cloud-t+	20	0	135796	14360	4032	S	0.0	2.7	0:20.86	cloud-torrent
289	cloud-t+	20	0	135796	14360	4032	S	0.0	2.7	0:36.16	cloud-torrent
290	cloud-t+	20	0	135796	14360	4032	S	0.0	2.7	0:20.10	cloud-torrent
303	cloud-t+	20	0	135796	14360	4032	S	0.0	2.7	0:59.56	cloud-torrent
304	cloud-t+	20	0	135796	14360	4032	S	0.0	2.7	0:12.25	cloud-torrent
17904	cloud-t+	20	0	135796	14360	4032	S	0.0	2.7	0:05.94	cloud-torrent
17905	cloud-t+	20	0	135796	14360	4032	S	0.0	2.7	1:00.08	cloud-torrent
17906	cloud-t+	20	0	135796	14360	4032	S	0.0	2.7	0:29.77	cloud-torrent
17907	cloud-t+	20	0	135796	14360	4032	S	0.0	2.7	0:16.15	cloud-torrent

You may have noticed how none of the attributes in the process list changed. How is that possible, given that processes differ from threads? Inside the Linux kernel, threads and processes are handled using the same data structures. Thus, every thread has its own ID, state and so on.

If you want to switch back to the process view, press ‘H’ again. In addition, you can use the `-H` switch to display threads by default.

`top -H`

## Showing full paths

By default, top does not show the full path to the program, or make a distinction between kernelspace processes and userspace processes. If you need this information, press ‘c’ while top is running. Press ‘c’ again to go back to the default.

```
top - 21:45:07 up 3 days, 1:39, 2 users, load average: 0.00, 0.00, 0.00
Tasks: 26 total, 1 running, 25 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 524288 total, 321760 free, 112476 used, 90052 buff/cache
KiB Swap: 131072 total, 90360 free, 40712 used. 331772 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
23692	supriyo	20	0	27344	2616	1264	S	0.4	0.5	0:02.68	tmux
1	root	20	0	37372	2036	1232	S	0.0	0.4	0:04.66	init -z
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	[kthreadd/646]
3	root	20	0	0	0	0	S	0.0	0.0	0:00.00	[khelper/646]
69	root	20	0	38896	3160	3024	S	0.0	0.6	0:01.43	/lib/systemd/systemd-journald
72	root	20	0	41632	136	88	S	0.0	0.0	0:00.17	/lib/systemd/systemd-udev
213	root	20	0	47572	104	40	S	0.0	0.0	0:00.25	/sbin/rpcbind -f -w
219	root	20	0	26012	300	196	S	0.0	0.1	0:00.36	/usr/sbin/cron -f
230	root	20	0	65464	332	220	S	0.0	0.1	0:00.23	/usr/sbin/sshd -D
231	syslog	20	0	184632	888	584	S	0.0	0.2	0:00.63	/usr/sbin/rsyslogd -n
237	cloud-t+	20	0	135796	14368	4040	S	0.0	2.7	11:29.68	/home/cloud-torrent/cloud-torrent
282	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	/sbin/agetty --noclear --keep-baud console 115200 +
286	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	/sbin/agetty --noclear tty2 linux
294	root	20	0	126072	304	200	S	0.0	0.1	0:00.00	nginx: master process /usr/sbin/nginx -g daemon on +
295	www-data	20	0	127032	2664	1780	S	0.0	0.5	1:05.62	nginx: worker process
296	www-data	20	0	126916	2544	1816	S	0.0	0.5	1:21.01	nginx: worker process
299	root	20	0	286060	524	396	S	0.0	0.1	0:08.81	php-fpm: master process (/etc/php/7.0/fpm/php-fpm.+
301	www-data	20	0	288624	9924	6816	S	0.0	1.9	0:00.11	php-fpm: pool www
302	www-data	20	0	288792	9636	6576	S	0.0	1.8	0:00.34	php-fpm: pool www
307	root	20	0	300484	10096	2480	S	0.0	1.9	1:01.86	/usr/bin/python3 /usr/bin/fail2ban-server -s /var/+
23659	root	20	0	92996	4040	3132	S	0.0	0.8	0:00.02	sshd: supriyo [priv]
23668	supriyo	20	0	93140	2184	1016	S	0.0	0.4	0:00.37	sshd: supriyo@pts/0

Kernelspace processes are marked with square brackets around them. As an example, in the above screenshot there are two kernel processes, `kthreadd` and `khelper`. On most Linux installations, there will usually be a few more of them.

Alternatively, you can also start top with the `-c` argument:

`top -c`

## Forest view

Sometimes, you may want to see the child-parent hierarchy of processes. You can see this with the forest view, by pressing ‘v’/’V’ while top is running.

```
top - 23:07:22 up 3 days, 3:01, 2 users, load average: 0.00, 0.00, 0.00
Tasks: 26 total, 1 running, 25 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 524288 total, 319940 free, 113816 used, 90532 buff/cache
KiB Swap: 131072 total, 90688 free, 40384 used, 330428 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	37372	2036	1232	S	0.0	0.4	0:04.74	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	`- kthreadd/646
3	root	20	0	0	0	0	S	0.0	0.0	0:00.00	`- khelper/646
69	root	20	0	38896	3188	3052	S	0.0	0.6	0:01.44	`- systemd-journal
72	root	20	0	41632	136	88	S	0.0	0.0	0:00.17	`- systemd-udev
213	root	20	0	47572	104	40	S	0.0	0.0	0:00.26	`- rpcbind
219	root	20	0	26012	300	196	S	0.0	0.1	0:00.37	`- cron
230	root	20	0	65464	332	220	S	0.0	0.1	0:00.24	`- sshd
24366	root	20	0	92996	4040	3132	S	0.0	0.8	0:00.03	`- sshd
24375	supriyo	20	0	92996	1956	1008	S	0.0	0.4	0:00.04	`- sshd
24376	supriyo	20	0	19472	3364	1600	S	0.0	0.6	0:00.04	`- bash
24390	supriyo	20	0	17868	1300	988	S	0.0	0.2	0:00.00	`- tmux
231	syslog	20	0	184632	888	584	S	0.0	0.2	0:00.63	`- rsyslogd
237	cloud-t+	20	0	135796	14336	4040	S	0.0	2.7	11:34.33	`- cloud-torrent
282	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	`-agetty
286	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	`-agetty
294	root	20	0	126072	304	200	S	0.0	0.1	0:00.00	`-nginx
295	www-data	20	0	127032	2664	1780	S	0.0	0.5	1:05.90	`-nginx
296	www-data	20	0	126732	2548	1872	S	0.0	0.5	1:21.10	`-nginx
299	root	20	0	286060	524	396	S	0.0	0.1	0:08.98	`-php-fpm7.0
301	www-data	20	0	288688	10320	7172	S	0.0	2.0	0:00.12	`-php-fpm7.0
302	www-data	20	0	288792	11304	8044	S	0.0	2.2	0:00.37	`-php-fpm7.0

As you can tell from the screenshot above, the `systemd` process was the first one to start up on the system. It has started processes such as `sshd`, which in turn has created other `sshd` processes, and so on.

## Listing processes from a user

To list processes from a certain user, press ‘u’ when top is running. Then, type in the username, or leave it blank to display processes for all users.

```
top - 20:13:10 up 3 days, 7 min, 2 users, load average: 0.00, 0.00, 0.00
Tasks: 26 total, 1 running, 25 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni,100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 524288 total, 334016 free, 104724 used, 85548 buff/cache
KiB Swap: 131072 total, 84460 free, 46612 used, 339556 avail Mem
Which user (blank for all)
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	37372	2036	1232	S	0.0	0.4	0:04.56	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd/646
3	root	20	0	0	0	0	S	0.0	0.0	0:00.00	khelper/646
69	root	20	0	38896	3108	2972	S	0.0	0.6	0:01.40	systemd-journal
72	root	20	0	41632	136	88	S	0.0	0.0	0:00.17	systemd-udev
213	root	20	0	47572	104	40	S	0.0	0.0	0:00.25	rpcbind
219	root	20	0	26012	300	196	S	0.0	0.1	0:00.36	cron
230	root	20	0	65464	332	220	S	0.0	0.1	0:00.23	sshd
282	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	agetty
286	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	agetty
294	root	20	0	126072	304	200	S	0.0	0.1	0:00.00	nginx
299	root	20	0	286060	524	396	S	0.0	0.1	0:08.64	php-fpm7.0
307	root	20	0	300484	10096	2480	S	0.0	1.9	1:00.96	fail2ban-server
23659	root	20	0	92996	4040	3132	S	0.0	0.8	0:00.02	sshd

Alternatively, you can run the top command with the `-u` switch. In this example, we have listed all processes from the root user.

```
top -u root
```

## Filtering through processes

If you have a lot of processes to work with, a simple sort won't work well enough. In such a situation, you can use top's filtering to focus on a few processes. To activate this mode, press 'o'/'O'. A prompt appears inside top, and you can type a filter expression here.

```
top - 23:47:23 up 3 days, 3:41, 2 users, load average: 0.00, 0.00, 0.00
Tasks: 26 total, 1 running, 25 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.1 us, 0.1 sy, 0.0 ni, 99.8 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 524288 total, 321356 free, 112216 used, 90716 buff/cache
KiB Swap: 131072 total, 90712 free, 40360 used, 332048 avail Mem

add filter #1 (ignoring case) as: [!]FLD?VAL !COMMAND=getty

PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ COMMAND
1 root 20 0 37372 2036 1232 S 0.0 0.4 0:04.78 init -z
2 root 20 0 0 0 0 S 0.0 0.0 0:00.00 [kthreadd/646]
3 root 20 0 0 0 0 S 0.0 0.0 0:00.00 [khelper/646]
69 root 20 0 38896 3212 3076 S 0.0 0.6 0:01.45 /lib/systemd/systemd-journald
72 root 20 0 41632 136 88 S 0.0 0.0 0:00.17 /lib/systemd/systemd-udevd
213 root 20 0 47572 104 40 S 0.0 0.0 0:00.26 /sbin/rpcbind -f -w
219 root 20 0 26012 300 196 S 0.0 0.1 0:00.37 /usr/sbin/cron -f
230 root 20 0 65464 332 220 S 0.0 0.1 0:00.24 /usr/sbin/sshd -D
231 syslog 20 0 184632 888 584 S 0.0 0.2 0:00.63 /usr/sbin/rsyslogd -n
237 cloud-t+ 20 0 135796 14336 4040 S 0.0 2.7 11:36.61 /home/cloud-torrent/cloud-torrent
282 root 20 0 12788 8 4 S 0.0 0.0 0:00.00 /sbin/agetty --noclear --keep-baud console 115200 +
286 root 20 0 12788 8 4 S 0.0 0.0 0:00.00 /sbin/agetty --noclear tty2 linux
294 root 20 0 126072 304 200 S 0.0 0.1 0:00.00 nginx: master process /usr/sbin/nginx -g daemon on+
295 www-data 20 0 127032 2664 1780 S 0.0 0.5 1:06.05 nginx: worker process
296 www-data 20 0 126732 2548 1872 S 0.0 0.5 1:21.11 nginx: worker process
299 root 20 0 286060 524 396 S 0.0 0.1 0:09.06 php-fpm: master process (/etc/php/7.0/fpm/php-fpm.+
301 www-data 20 0 288688 10320 7172 S 0.0 2.0 0:00.12 php-fpm: pool www
302 www-data 20 0 288792 11304 8044 S 0.0 2.2 0:00.37 php-fpm: pool www
307 root 20 0 300484 10096 2480 S 0.0 1.9 1:02.91 /usr/bin/python3 /usr/bin/fail2ban-server -s /var/+
24519 root 20 0 92996 4040 3132 S 0.0 0.8 0:00.02 sshd: supriyo [priv]
24528 supriyo 20 0 92996 1948 1008 S 0.0 0.4 0:00.04 sshd: supriyo@pts/0
24529 supriyo 20 0 19472 3360 1596 S 0.0 0.6 0:00.05 -bash
```

A filter expression is a statement that specifies a relation between an attribute and a value. Some examples of filters are:

- `COMMAND=getty`: Filter processes which contain “getty” in the COMMAND attribute.
- `!COMMAND=getty`: Filter processes which do not have “getty” in the COMMAND attribute.
- `%CPU>3.0`: Filter processes which have a CPU utilization of more than 3%.

Once you've added a filter, you can further prune down things by adding more filters. To clear any filters you have added, press '='.

## Changing the default look of the CPU and memory statistics

If you are mostly at home in a GUI environment, you might not like top's default way of showing CPU and memory statistics. You can press 't' and 'm' to change the style of the CPU and memory statistics. Here is a screenshot of top, where we have pressed 't' and 'm' once.

```

top - 23:17:09 up 3 days, 3:11, 2 users, load average: 0.00, 0.00, 0.00
Tasks: 26 total, 1 running, 25 sleeping, 0 stopped, 0 zombie
%Cpu(s):  0.2/0.0  0[
KiB Mem : 37.0/524288  [|||||]
KiB Swap: 30.8/131072  [|||||]

```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	37372	2036	1232	S	0.0	0.4	0:04.74	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd/646
3	root	20	0	0	0	0	S	0.0	0.0	0:00.00	khelper/646
69	root	20	0	38896	3196	3060	S	0.0	0.6	0:01.44	systemd-journal
72	root	20	0	41632	136	88	S	0.0	0.0	0:00.17	systemd-udev
213	root	20	0	47572	104	40	S	0.0	0.0	0:00.26	rpcbind
219	root	20	0	26012	300	196	S	0.0	0.1	0:00.37	cron
230	root	20	0	65464	332	220	S	0.0	0.1	0:00.24	sshd
231	syslog	20	0	184632	888	584	S	0.0	0.2	0:00.63	rsyslogd
237	cloud-t+	20	0	135796	14336	4040	S	0.0	2.7	11:34.87	cloud-torrent
282	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	agetty
286	root	20	0	12788	8	4	S	0.0	0.0	0:00.00	agetty
294	root	20	0	126072	304	200	S	0.0	0.1	0:00.00	nginx
295	www-data	20	0	127032	2664	1780	S	0.0	0.5	1:05.94	nginx
296	www-data	20	0	126732	2548	1872	S	0.0	0.5	1:21.10	nginx
299	root	20	0	286060	524	396	S	0.0	0.1	0:09.00	php-fpm7.0
301	www-data	20	0	288688	10320	7172	S	0.0	2.0	0:00.12	php-fpm7.0
302	www-data	20	0	288792	11304	8044	S	0.0	2.2	0:00.37	php-fpm7.0
307	root	20	0	300484	10096	2480	S	0.0	1.9	1:02.64	fail2ban-server
23692	supriyo	20	0	28268	3568	1276	S	0.0	0.7	0:04.38	tmux
23693	supriyo	20	0	19480	3404	1628	S	0.0	0.6	0:00.09	bash
24366	root	20	0	92996	4040	3132	S	0.0	0.8	0:00.03	sshd

If you press ‘t’ or ‘m’ repeatedly, it cycles through four different views. In the first two presses, it cycles through two different types of progress bars. If you press the key for a third time, the progress bar is hidden. If you press the key again, it brings back the default, text based counters.

## Saving your settings

If you’ve made any changes to top’s output, you can save them for later use by pressing ‘W’. top writes its configuration to the `.toprc` file in your home directory.

## Conclusion

The top command is extremely helpful for monitoring and managing processes on a Linux system. This article only scratches the surface, and there is quite a bit we haven’t covered. For example, there are a lot more columns you can add to top. For all this stuff, be sure to check out the man page by running `man top` on your system.